



Nebula: Performance and Energy Efficiency in Serverless Computing

A Comparative Study of WebAssembly and Docker

Marius Nilsen Kluften

Informatics: Programming and System Architecture
60 ECTS Study Points

Department of Informatics
Faculty of Mathematics and Natural Sciences

Spring 2024



Nebula

Performance and Energy Efficiency in
Serverless Computing

A Comparative Study of WebAssembly and Docker

Marius Nilsen Kluften

© 2024 Marius Nilsen Kluften

Nebula

<https://duo.uio.no/>

Printed: Reprosentralen, University of Oslo

Abstract

The ever increasing demand for cloud services has resulted in the expansion of energy-intensive data centers, the ICT industry accounts for about 1 % of global electricity use, highlighting a need for sustainable options in cloud computing architectures.

This thesis investigates WebAssembly, a technology originally intended for running in the browser, as a potential contender in the space of technologies to consider in cloud native applications. Leveraging the inherent efficiency, portability and lower startup times of WebAssembly modules, this thesis presents an approach that aligns with green energy principles, while maintaining performance and scalability, essential for cloud services.

Preliminary findings suggest that programs compiled to WebAssembly modules have reduced startup and runtimes, which hopefully leads to lower energy consumption, offering a viable pathway towards more sustainable cloud computing.

Acknowledgments

The inspiration for this thesis originated from an episode of the podcast "Rustacean Station," where Matt Butcher, the CEO of Fermyon, discussed the capabilities of WebAssembly running on the server with the aid of WebAssembly System Interface. This piqued my interest and set in motion the events that culminated in this thesis. I would like to express my gratitude to Matt Butcher and the team at Fermyon for inadvertently inspiring my topic.

I also extend my sincere thanks to my two supervisors, Joachim Tilsted Kristensen and Michael Kirkedal Thomsen, whom I somehow managed to convince to guide me through such a cutting-edge topic. Their guidance and insight have been invaluable throughout the past semesters.

The completion of this thesis alongside a full-time job was made possible, in part, by the support and patience of my girlfriend, Ingvild Stølen, whose feedback throughout the project proved invaluable.

I would also like to express my gratitude to my peer, Håkon Thorkildsen Smørvik, for his valuable feedback and assistance in reviewing and testing the prototype during its development, often breaking it in unexpected ways.

Finally, I would like to thank Syrus Akbary, founder of Wasmmer. If I hadn't met him at the WasmIO 2024 conference in Barcelona, my startup times would still be 100 times slower.

Contents

Abstract	i
Acknowledgements	ii
Contents	iii
List of Figures	vii
List of Tables	ix
I Overview	1
1 Introduction	2
1.1 Motivation	3
1.2 The Project	3
1.3 Problem Statement	3
1.4 Outline	4
2 Three Waves of Cloud Compute	5
2.1 Ashore: Before the Waves	5
2.2 The First Wave: Virtual machines	7
2.3 The Second Wave: Containers	8
2.4 The Third Wave: WebAssembly Modules	9
3 Background	11
3.1 Cloud Computing Overview	11
3.2 Energy Consumption and Sustainability in Cloud Computing	13
3.3 Virtualization and Virtual Machines	15
3.4 Containers and Container Orchestration	16
3.5 Serverless Computing	18
3.5.1 Functions-as-a-Service	19

3.6	WebAssembly and WASI	20
3.6.1	asm.js	21
3.6.2	WebAssembly	21
3.6.3	WebAssembly System Interface	23
3.6.4	WebAssembly Runtimes	23
3.7	Energy Monitoring	25
3.7.1	MQTT	25
3.7.2	Z-Wave	25
3.7.3	Aeotec Smart Switch	25
3.7.4	Modbus TCP	26
3.7.5	Gude Expert Power Control 1105	27
II	Project	28
4	Methodology	29
4.1	Experimental Framework	29
4.1.1	Prototyping	30
4.1.2	Controlled Experimentation	31
4.1.3	Benchmarking	32
4.1.4	Benchmark Functions	32
4.1.5	Measurement and Data Collection	33
4.1.6	Data Analysis	36
4.1.7	Comparative Analysis	36
4.1.8	Reliability and Validity	36
5	Designing Nebula	38
5.1	Prototype Scope and Objectives	38
5.2	System Architecture	39
5.3	Nebula	40
5.3.1	Function Packaging and Deployment	41
5.3.2	Function Execution	42
5.3.3	Interface	42
5.4	Infrastructure	42
5.5	Energy Measurements	43
5.6	Benchmarking utility	43
5.7	Experiment Design	44
6	Implementing Nebula	45
6.1	Choosing a Tech Stack	46

6.1.1	Rust and WebAssembly	46
6.1.2	Wasmtime	46
6.1.3	Docker Environment	46
6.2	Nebula Core	47
6.2.1	Building the Web Server	47
6.3	Calling Functions	48
6.3.1	Executing WebAssembly (Wasm) Modules	50
6.3.2	Executing Docker Containers	52
6.3.3	Challenges and Insights	54
6.4	Function Development and Deployment	57
6.4.1	Development Environment	57
6.4.2	Shared Library	57
6.4.3	Implementing a Benchmark Function	61
6.5	Deployment Setup	62
6.5.1	Raspberry Pi Setup	62
6.5.2	Virtual Machine Setup (Norwegian Research and Education Cloud (NREC))	62
6.5.3	Software Setup	63
6.5.4	Tying It All Together	63
6.6	Benchmarking	64
6.6.1	Stress-Testing Nebula	64
6.6.2	Measuring Power	67
6.6.3	Reading from Modbus TCP	68
6.6.4	Attaching Power Measurement to Function Invocations	69
6.6.5	Energy Measurement Challenges	72
6.7	Testing - Validity and Rigidity	73
6.7.1	Rust's Compiler Guarantees	73
6.7.2	Rigorous Testing and Benchmarking	73
III	Analysis	75
7	Results	76
7.1	Methodology Recap	76
7.2	Startup and Runtimes	77
7.3	Power Load and Energy Consumption Results	82
8	Discussion	87
8.1	Startup Latency and Runtime Performance	87
8.2	Power Consumption and Energy Efficiency	88

8.3 Function storage	89
8.4 Related Work	89
9 Conclusion	92
10 Future work	94
10.1 Investigating Docker Performance Discrepancies	94
10.2 Other Programming Languages as Source	95
10.3 Distributed Deployment and Power Measurement	95
10.4 Accurate Power Consumption Measurements	95

List of Figures

2.1	Example of a company that host their own infrastructure.	6
2.2	Example of “Headbook“ building their services on EC2 and using S3 for storage.	7
2.3	DevOps engineer deploying services as containers on Google Cloud Platform (GCP).	8
2.4	An example polyglot app built with different languages that all compile to Wasm and can interface together.	10
3.1	Projected energy consumption in 2026: Skien data center and Norway.	13
3.2	Virtual machines running on top of Hypervisor on a computer.	16
3.3	Containers running on the Docker Engine.	17
3.4	Source code in C/C++ compiled to asm.js and run in browser.	21
3.5	Source code compiled to WebAssembly and embedded in browser.	22
3.6	Source code compiled to wasm_wasi32 and deployed on platforms that support running the binaries.	23
3.7	Source code compiled to wasm_wasi32 that can run anywhere a WebAssembly runtime can be installed.	24
3.8	PubSub model of the Message Queuing Telemetry Transport (MQTT) protocol.	25
3.9	Smart Switch 6 communicating with zwave-js-ui through Aeotec Smart Stick.	26
3.10	TCP/IP package exchange between a client and a server over Modbus TCP.	27
3.11	Client communicating with Gude over Modbus TCP.	27
4.1	Reading energy consumption of our Raspberry Pi.	33
5.1	Overall architecture of our benchmarking setup.	40
5.2	User interacting with Nebula.	40
5.3	Flow of a user invoking a function on Nebula.	41
6.1	Streamlined flow of building and deploying functions.	60

6.2	Our virtual machine configured on NREC.	63
6.3	A Modbus TCP packet.	67
6.4	Power measurements that overlap with function invocations.	70
6.5	The physical setup, including the Raspberry Pi connected to power through our Gude power supply. The Raspberry Pi and Gude controller is connected to the benchmark computer on a local network through ethernet.	72
7.1	Mean startup and runtimes on both our setups.	77
7.2	Wasm runtime measurements on VM during benchmarks.	80
7.3	Startup and Runtime measurements on Raspberry Pi 4B during benchmarks.	81
7.4	Average energy consumption for our functions.	82
7.5	Power load on Raspberry Pi during Fibonacci benchmark.	83
7.6	Energy consumption on Raspberry Pi during Fibonacci benchmark.	83
7.7	Power load on Raspberry Pi during Factorial benchmark.	84
7.8	Energy consumption on Raspberry Pi during Factorial benchmark.	84
7.9	Power load on Raspberry Pi during Exponential benchmark.	85
7.10	Energy consumption on Raspberry Pi during Factorial benchmark.	85
7.11	Power load on Raspberry Pi during Prime Number benchmark.	86
7.12	Energy consumption on Raspberry Pi during Prime Number benchmark.	86
8.1	Machine learning training with SGD with Faaslets (FAASM) and containers (Knative). Reprinted with permission.	90
8.2	Different Kubernetes operators memory usage. Comparing Rust to WASM vs Rust container vs Golang container. © 2022 IEEE	91

List of Tables

3.1	WebAssembly features	22
4.1	Nebula requirements	31
4.2	Gude Expert Power Control 1105 sensors.	34
7.1	Comparison of Wasm vs. Docker cold starts on VM.	79
7.2	Comparison of Wasm vs. Docker cold starts on Raspberry Pi.	79
8.1	Comparison of storage size of functions.	89
8.2	Comparison of Faaslets vs. container cold starts. Reprinted with permission.	90

Part I

Overview

Chapter 1

Introduction

If WASM+WASI existed in 2008, we wouldn't have needed to created Docker. That's how important it is. Webassembly on the server is the future of computing. A standardized system interface was the missing link. Let's hope WASI is up to the task!

—Solomon Hykes, *Founder of Docker*

In the digital age, cloud computing has emerged as a foundational technology in the technological landscape, driving innovation and increased efficiency across various sectors. Its growth over the past decade has not only transformed how consumers store, process, and access data, but it has also raised environmental concerns as more and more data centers are built around the globe to accommodate the traffic, consuming vast amounts of power. The Information and Communication Technology (ICT) industry, with cloud computing at its core, accounts for an estimated 2.1% to 3.9% of global greenhouse gas emissions. Data centers, the backbone of cloud computing infrastructures, are responsible for about 200 TWh/yr, or about 1% of the global electricity consumption, a figure projected to escalate, potentially reaching 15% to 30% of electricity consumption in some countries by 2030 (Freitag et al., 2021).

The sustainability of cloud computing is thus under scrutiny, and while some vendors strive to achieve a net-zero carbon footprint for their cloud computing services, many data centers still rely on electricity generated by fossil fuels, a leading contributor to climate change (Mytton, 2020). This reality emphasizes an urgent need to explore alternative technologies that promise enhanced energy efficiency while meeting customers demands. In this vein, serverless computing has emerged as a compelling paradigm, offering scalability and flexibility by enabling functions

to execute in response to requests, rather than having a server running all the time. However, the inherent startup latency associated with containerized serverless functions pose a challenge, particularly for on-demand applications. To mitigate this, vendors often opt for keeping the underlying servers *warm* to keep the startup latency as low as possible for serving functions. Reducing the startup time for serving a function should reduce the need for keeping servers warm and therefore reduce the standby power consumption of serverless architectures.

1.1 Motivation

The environmental footprint of cloud computing, particularly the energy demands of data centers, is a pressing issue. As the digital landscape continues to evolve, the quest for sustainable solutions has never been more critical. This thesis is motivated by the need to reconcile the growing demand for cloud services with the pressing need for environmental sustainability. Through the lens of WebAssembly and WebAssembly System Interface (WASI), this thesis aims to investigate innovative deployment methods that promise to reduce energy consumption without sacrificing performance, thereby contributing to the development of a more sustainable cloud computing ecosystem.

1.2 The Project

This thesis explores Wasm with WASI as an innovative choice for deploying functions to the cloud, through developing a prototype Functions-as-a-Service (FaaS) platform named Nebula. This platform will run functions compiled to Wasm, originally designed for high-performance tasks in web browsers, which coupled with WASI, allows us to give WebAssembly programs access to the underlying system. This holds potential for a more efficient way to package and deploy functions, potentially reducing the startup latency and the overhead associated with traditional serverless platforms. Wasm and WASI offers a pathway, where the demands of today is met, while reducing the carbon footprint for cloud applications.

1.3 Problem Statement

In this thesis, we are interested in exploring the potential of Wasm and WASI as an energy-efficient alternative for deploying functions in a serverless computing environment. The primary focus is to investigate whether Wasm-based functions can offer improved performance and reduced energy consumption compared to

traditional containerized functions.

Based on this, we can define these goals that we want to achieve with this thesis:

- Develop a prototype cloud computing platform for the FaaS paradigm.
- Use this prototype to compare the performance of Wasm-based functions to containerized functions in terms of startup latency and execution time
- Find out if Wasm-based functions consume less energy than containerized functions when executing equivalent tasks

By achieving these goals this thesis seeks to shed light on the feasibility and implications of adopting Wasm and WASI for a greener cloud.

1.4 Outline

The thesis is structured to guide the reader through the background, design, implementation, and analysis of the Nebula project. The key aspects covered include:

- The evolution of cloud computing and the emergence of WebAssembly as a promising technology for serverless computing.
- The design and implementation of the Nebula prototype, highlighting the integration of Wasm and WASI, using the Rust programming language.
- The experimental setup and methodology employed to compare the performance and energy consumption of Wasm-based functions against containerized functions.
- An in-depth analysis and discussion of the experimental results, addressing the implications for energy efficiency in serverless computing.
- Concluding remarks and suggestions for future work, exploring the potential for further research and development in this area.

Chapter 2

Three Waves of Cloud Compute

If VMs were the first wave and containers were the second, WebAssembly is the third wave of cloud compute.

—Matt Butcher, CEO of Fermyon

The evolution of cloud computing represents a transformative adventure, driven by the pursuit for efficiency, scalability and reliability, yet it also poses challenges, notably its environmental impact. This chapter steps introduces the concept of the “Three waves of cloud computing”, coined by the WebAssembly community (Butcher & Dodds, 2024; Leonard, 2024). Where the two first waves of cloud compute represent the shift from virtual machines to containerization, the third wave encompasses utilizing Wasm and WASI to build the next era of cloud compute with the potential to significantly reduce the carbon footprint.

2.1 Ashore: Before the Waves

Before delving into the waves themselves, it is useful to understand the landscape that preceded cloud computing. Prior to the cloud era, companies were required to building and mantaining the physical infrastructure for their digital services in-house. This required companies to invest heavily into both expensive hardware and expensive engineers to buy, upkeep and oversee their own physical servers and network hardware. (See Figure 2.1 for an example)

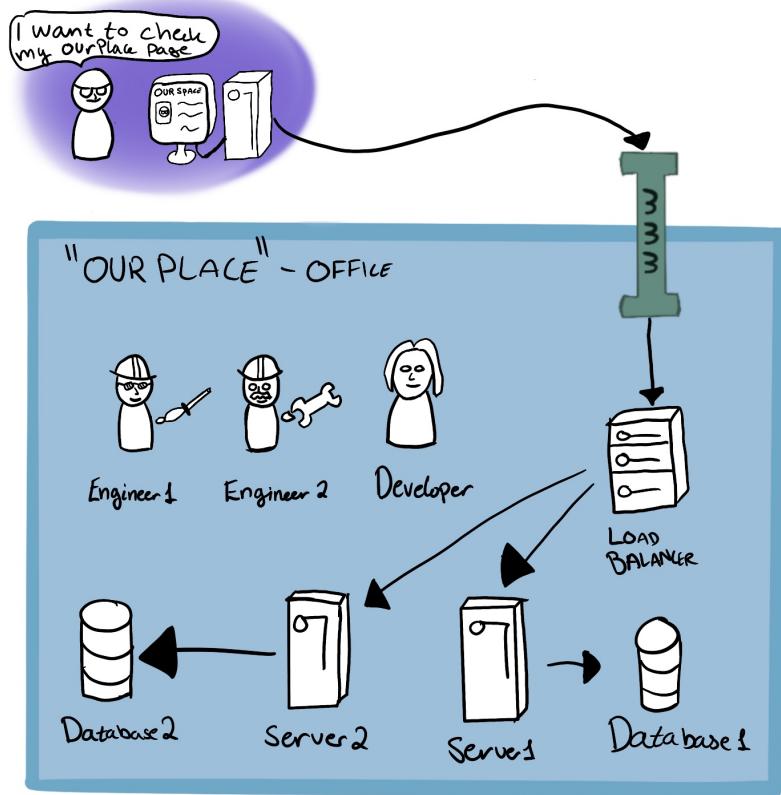


Figure 2.1: Example of a company that host their own infrastructure.

Setting up such an infrastructure comes with a significant cost, both upfront for purchasing such a setup and for employing engineers who can set up and maintain this. This puts a considerable financial strain on organizations, and kept smaller companies without the financial backing to invest in this at a disadvantage (Etro, 2009). Having to maintain infrastructure also poses a challenge when the services start to gain traction. To meet the increased traffic on the services, a company will have to scale up and extend the infrastructure, which also can be expensive and difficult (Armbrust et al., 2010).

As a response to these challenges, some companies found a market for taking on the responsibility of managing infrastructure, and offer Infrastructure-as-a-Service (IaaS) to an evolving market that relies more and more on digital solutions. IaaS provides consumers with the ability to provision computing resources where they can deploy and run software, including operating systems and applications (Mell & Grance, 2011). On these managed infrastructures companies could deploy their services on top of virtual machines that allowed more flexibility, and lowered the bar to new companies.

2.2 The First Wave: Virtual machines

The start of cloud computing can be traced back to the emergence of virtualization, more specifically virtual machines, a response to the costly and complex nature of managing traditional, on-premise data centers. During the mid-2000s, Amazon launched its subsidiary, Amazon Web Services (AWS), who in turn launched Amazon S3 in March 2006, followed by Elastic Compute Cloud (EC2) in August the same year (Barr, 2006). With these services, AWS positioned itself as a pioneer in this space, marking a major turning point in application development and deployment, and popularized cloud computing. EC2, as an IaaS platform, empowered developers to run virtual machines remotely. (See Figure 2.2 for an example of this kind of architecture)

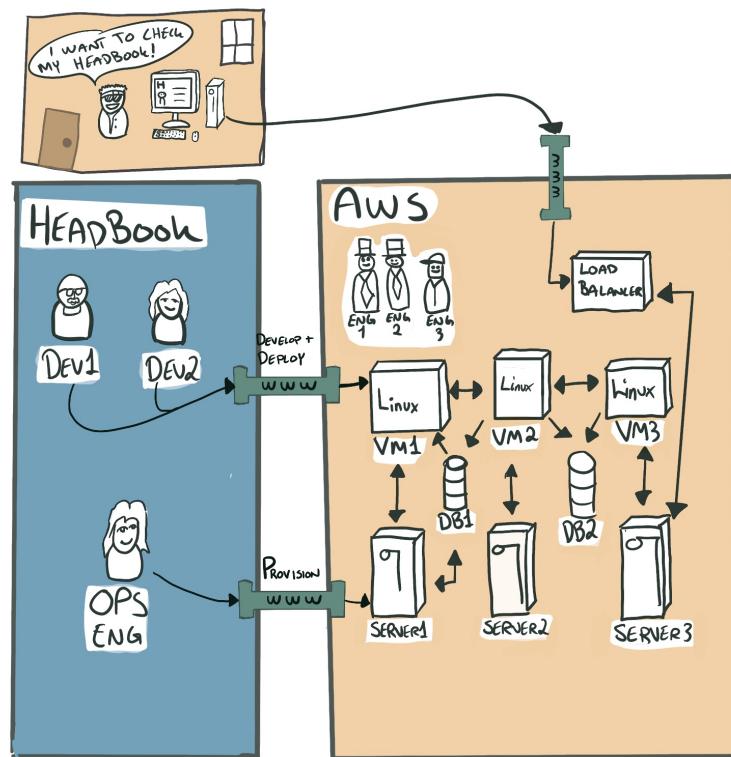


Figure 2.2: Example of “Headbook“ building their services on EC2 and using S3 for storage.

While similar services existed before 2006, with Amazon’s existing large customer base helped them gain significant traction, and ushered in a the first era, or wave, of *cloud computing*.

2.3 The Second Wave: Containers

As we entered the 2010s, the focus shifted from virtual machines to containers, largely due to the limitations of VMs in efficiency, resource utilization, and application deployment speed. Containers, being a lightweight alternative to VMs, designed to overcome these hurdles (Bao et al., 2016).

In contrast to VMs, which require installation of resource-intensive operating systems and minutes to start up, containers along with their required OS components, could start up in seconds. Typically managed by orchestration tools like Kubernetes¹, containers enabled applications to package alongside their required OS components, facilitating scalability in response to varying service loads. Consequently, an increasing number of companies have since established platform teams to build orchestrated developer platforms, thereby simplifying application development in Kubernetes clusters. (See Figure 2.3 for an example where a fictive company WacDonalds and their container workflow)

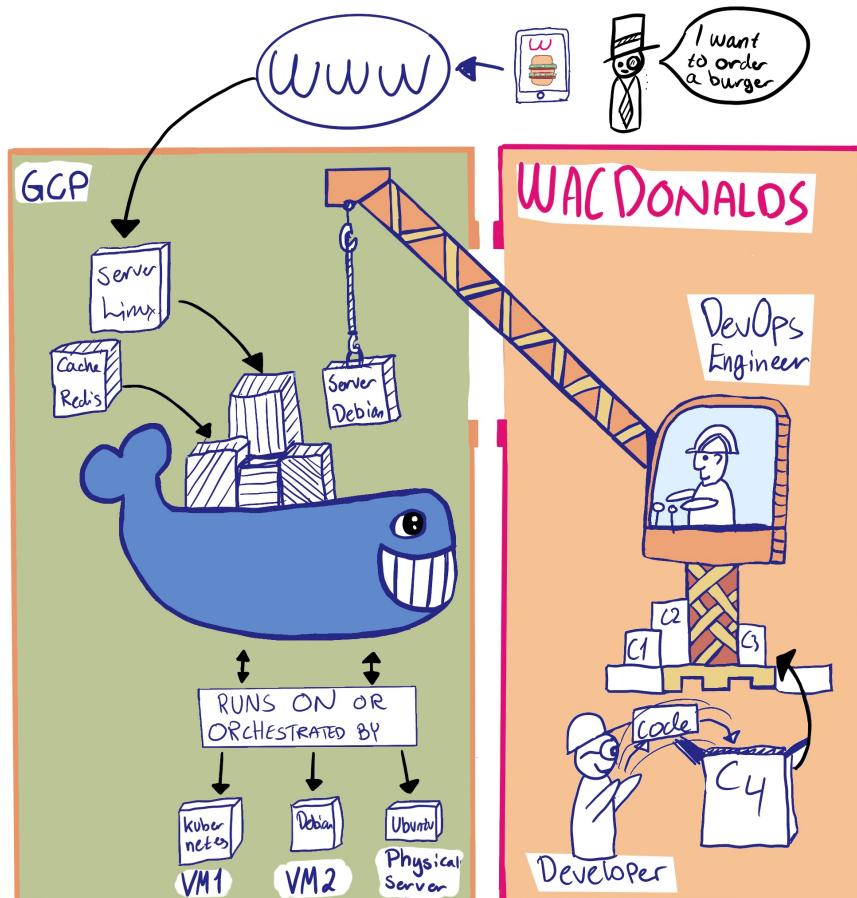


Figure 2.3: DevOps engineer deploying services as containers on GCP.

¹<https://kubernetes.io>

Containers are not a perfect solution however, and while they simplify the means of developing and deploying applications, docker images can easily reach Gigabytes in image size (Durieux, 2024), can take a long time to start up, and building applications that target multiple platforms can be difficult.

These solutions are more efficient than manually installing an operating system on a machine, but they still have leave a large footprint. Is there a more efficient way to package and deploy our programs? Wasm and WASI, as mentioned in epigraph of chapter 1, has positioned itself as a potential contender for how applications are built, packaged and deployed to the cloud.

2.4 The Third Wave: WebAssembly Modules

WebAssembly has had a surge of popularity the past three to four years when developers discovered that what it was designed for - to truly run safely inside the browser - translated well into a cloud native environment as well. Containers have, with the benefits mentioned in section 2.3, had a positive impact on the cloud native landscape. However, with the limitations - like large image sizes, slow startups and complexity of cross-platform - there is space for exploring alternative technologies for building our cloud-native applications.

WebAssembly is a compilation target with many languages adopting support, and by itself, it is sandboxed to run in a WebAssembly virtual machine (VM) without access to the outside world, meaning that it cannot access the underlying system. This means that a “vanilla” WebAssembly module cannot write to the file system, update a Redis cache or transmit a POST request to another service.

To make this possible, the WebAssembly System Interface project was created. This project allows developers to write code that compiles to WebAssembly that can access the underlying system. This is the key project that turned many developers onto the path of exploring WebAssembly as a potential contender for building cloud applications. With WebAssembly, developers can write programs in a programming language that supports it as a compilation target, such as Rust, C, C++, Go, and build tiny modules that can run on a WebAssembly runtime. These WebAssembly runtimes can run on pretty much any architecture with ease, the resulting binary size are quite small, and the performance is near-native. These perks combined with the potential for reduced overhead, smaller image sizes, and faster startup times make WebAssembly and WASI a promising candidate for the third wave of cloud compute with a lower impact on the environment.

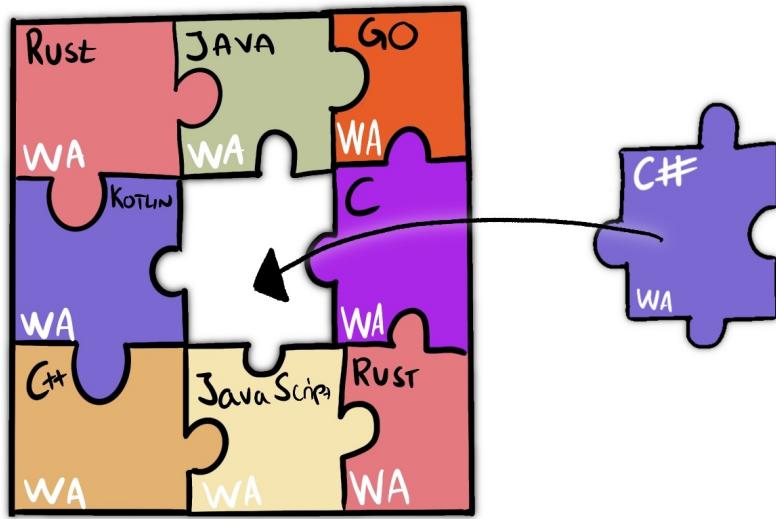


Figure 2.4: An example polyglot app built with different languages that all compile to Wasm and can interface together.

In summary, the three waves of cloud computing - virtual machines, containers, and now Wasm with WASI - represent the industry's pursuit of more efficient, scalable and reliable solutions for building cloud applications. While each wave has attempted to tackle pressing challenges of its time, it is exciting to see how Wasm and WASI can be leveraged in this third wave and see if its promise of more efficient applications can lead to reducing the environmental impact of ICT.

Chapter 3

Background

We were interested in applications that needed to start, run to completion, and shut down nearly instantly. Examples of this are AWS Lambda and Azure Functions. But those were built on VMs, which are slow to start. So they had to pre-warm huge numbers of VMs, which sat idle until they received a request. We were looking for a more efficient runtime that could cold-start in under one millisecond. And that's what brought us to WebAssembly.

—Matt Butcher, CEO of Fermyon

3.1 Cloud Computing Overview

Cloud computing, commonly referred to as “*the cloud*”, refers to the delivery of computing resources served over the internet, as opposed to traditional on-premise hardware setups. The National Institute of Standards and Technology (NIST) defines cloud computing like so:

NIST definition of Cloud Computing

Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction.

(Mell & Grance, 2011)

Cloud computing traces its root back to the 1960s, with the Compatible Time-Sharing System (CTSS) project at MIT, which demonstrated the potential for multiple users accessing and sharing computing resources simultaneously (Crisman, 1963). While CTSS was a localized system, it paved the way for the concept of shared computing resources, a fundamental principle of cloud computing.

Over the following decades, advancements in networking, virtualization and the ubiquity of the internet led to the development of todays sophisticated cloud services. The term “cloud computing” was first coined in the year 1996 by Compaq, (Favaloro & O’Sullivan, 1996), but it was not until Amazon launched its subsidiary AWS in the 2006 that the adoption became wide spread.

The launch of AWS Amazon S3 cloud service in March 2006, followed by Elastic Compute Cloud (EC2) in August the same year (Barr, 2006), marked a major turning point in application development and deployment, and popularized cloud computing. EC2, as an Infrastructure-as-a-Service platform, empowered developers to run virtual machines remotely. By providing these services over the internet on a pay-as-you-go basis, AWS drastically lowered the bar for accessing computing resources, making it easier and more cost-effective for businesses and developers to build and deploy applications without the need for considerable upfront investment in hardware and infrastructure.

With the success of AWS, other major technology companies saw their fit to enter the cloud computing market. In 2008, Google launched the Google App Engine (McDonald, 2008), a platform for building and hosting web applications in Google’s data centers. Microsoft followed with the launch of Azure in 2010, its cloud computing platform that offers a range of services comparable to AWS.

The rapid growth of cloud computing also fueled the rise of DevOps practices and containerization technologies like Docker, which facilitate the development, deployment and management of applications on the cloud. Orchestration tools like Docker Swarm and Kubernetes further simplify the process of managing and scaling containerized applications across cloud environments (Bernstein, 2014).

Today, cloud computing has become an essential part of modern IT infrastructure, where major cloud providers, like AWS, Microsoft and Google, continue to innovate and expand their offerings. Cloud computing has also enable new paradigms like serverless computing and edge computing, allowing for even more efficient and distributed computing models. (Baldini et al., 2017)

3.2 Energy Consumption and Sustainability in Cloud Computing

A major drawback of cloud computing is the increasing energy consumption required to power data centers. With the increased demand for energy consumption, comes an increased impact on the environment. As mentioned in chapter 1, the ICT industry accounts for an estimated 2.1% to 3.9% of global greenhouse emissions (Freitag et al., 2021). According to the International Energy Agency (IEA), data centers across the globe consumed between from 240 to 340 TWh, accounting for 1 to 1.3% of the global electricity use.

In Norway, for example, Google is constructing a data center in Skien, expected to be fully operational by 2026. As of April 2024, they have been granted a capacity of 240 Megawatts, but they have applied for a total capacity of 860 Megawatts (Rivrud, 2024). At full capacity at 860 MW, Google's data center is aiming to consume 7.5 TWh each year, and according to Google's most recent sustainability report, they consumed a total of 22.29 TWh globally in 2022 ("Google 2023 Environmental Report", 2023). In other words, in 2026 the data center in Skien alone is projected to consume ~33% of the energy Google consumed globally in 2022. The energy consumption in Norway is projected to reach 150-158 TWh in 2026 (Gunnerød, 2022), meaning that the data center in Skien could account for 5% of the energy consumption in the country. Figure 3.1 below illustrates the amount of energy the new data center will consume compared to the entire energy consumption of Norway.

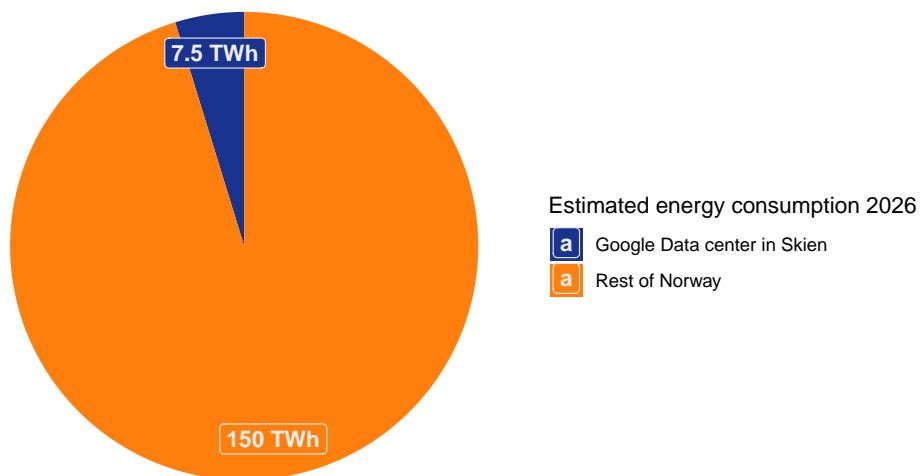


Figure 3.1: Projected energy consumption in 2026: Skien data center and Norway.

The flipside of this, is that Google is committed to reach a net-zero carbon footprint by 2030, and the data center in Skien is built to reflect this. Google has been carbon

neutral since 2007 and has matched 100% of its annual electricity consumption with renewable energy since 2017 (Google, n.d.). Norway's abundant hydropower, rising wind power production, investment into solar energy and other renewable energy sources make it an ideal location for building data centers aiming to be powered by renewable energy (Norwegian-Energy, 2023).

It is important to note that while Google is committed to using renewable energy, their presence in the area does not necessarily lead to an increase in renewable energy infrastructure. If Google purchases all the available green energy in the region, it may result in other consumers relying more on non-renewable energy sources (Sundt & Rehdanz, 2015).

Like Google, other major cloud providers have set ambitious targets for renewable energy adoption and have invested in large-scale renewable projects. Microsoft has committed to shifting to 100% renewable energy by 2025 for all its data centers, buildings and campuses, and to be carbon *negative* by 2030 (Smith, 2020), while Amazon has pledged to transition to 100% renewable energy by 2030 for its cloud subsidiary and to have a net-zero carbon footprint by 2040 (Amazon, 2019). These efforts have contributed to reducing greenhouse gas emissions in the cloud computing industry, but this commitment is not universally adopted, and many data centers still rely on electricity generated by fossil fuels, a leading contributor to climate change (Mytton, 2020).

It is crucial to acknowledge that the increasing demand for data centers also leads to a higher overall power demand. This increased demand for energy can result in more natural habitats being destroyed to accommodate the necessary power infrastructure. The construction of new data centers and the associated energy infrastructure can have a significant impact on the environment and local ecosystems (Bologna & Aquino, 2020).

Several factors make up the energy consumption required to service a data center. One of these factors is cooling down the servers while running, and a study from 2017 discovered that cooling accounted for about 38% of total energy consumption in data centers, ranging from 21% to 61% depending the effectiveness of the facility's heating, ventilation, and air conditioning (HVAC) system (Ni & Bai, 2017).

One innovative example of attempting to mitigate the environmental impact of cooling data centers can be found by data center providers like DeepGreen¹, who submerge their servers into dielectric fluid, which gets warmed up by the excess heat of the computers. This heat is then transferred to a host's hot water system via a heat exchange and used for heating up swimming pools in London. Another

¹<https://deepgreen.energy/faqs/>

broad strategy cloud providers opt for is the implementation of power management techniques, such as dynamic voltage and frequency scaling (DVFS), which adjusts the power consumption of servers based on workload demands (Beloglazov et al., 2012).

Virtualization and resource pooling, two key components of cloud computing, also contribute to energy efficiency. By consolidating virtual machines onto shared physical servers, cloud providers can improve resource utilization and reduce the energy consumption of their data centers. (Beloglazov et al., 2012)

To summarize, in the context of the growing energy consumption of data centers and its environmental impact, major cloud providers such as Google, Microsoft, and Amazon have made significant commitments to renewable energy adoption and sustainability. However, it is important to acknowledge that reducing energy consumption is the most sustainable approach. Researchers and industry experts have emphasized the importance of prioritizing energy efficiency, optimizing resource utilization, and adopting innovative solutions like submerged servers and power management techniques to minimize the environmental footprint of cloud computing {shehabiUnitedStatesData2016,myton2020,danilakCouncilPost-Why2017}. The European Climate, Infrastructure and Environment Executive Agency underscores this point, stating, “The greenest and most ethical energy is the energy we don’t consume” (Cappelletti, 2022).

3.3 Virtualization and Virtual Machines

Virtualization is the process of creating a virtual version of a physical resource such as an operating system, a server, a storage device, or a network resource (Chiueh & Nanda, 2005). Virtualization allows multiple virtual instances to share the underlying physical hardware, enabling more efficient resource utilization and consolidation.

One of the most common forms of virtualization is the creation of virtual machines (VMs). Barham et al. (2003) described that a virtual machine is a software-based emulation of a physical computer system, including its processor, memory, storage, and network interfaces. Furthermore they describe that VMs run on top of a hypervisor, a software layer that manages and allocates the physical hardware resources to the virtual machines.

Figure 3.2 below illustrates virtual machines running in such an environment.

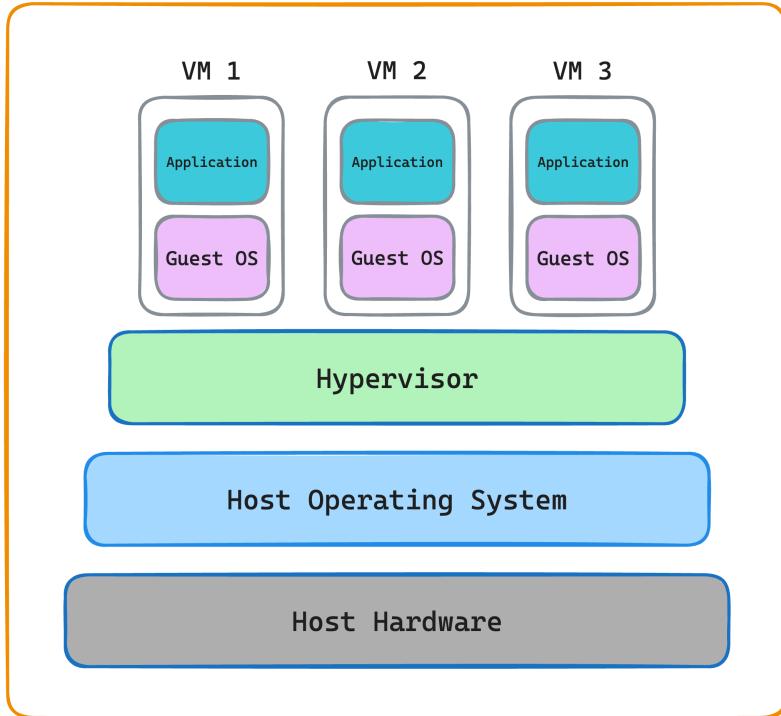


Figure 3.2: Virtual machines running on top of Hypervisor on a computer.

By running several virtual machines on the same physical server, data centers can drastically decrease the number of physical machines physical machines needed to operate, leading to a reduction in energy consumption (Kaplan et al., 2008).

3.4 Containers and Container Orchestration

While virtual machines provide isolation at the hardware level, containers offer a more lightweight form of virtualization by isolating applications at the operating system level (Merkel, 2014). Containers share the host operating system's kernel, enabling them to be more lightweight and efficient compared to traditional virtual machines. They can run physical servers, as well as on VMs (Bernstein, 2014).

Merkel (2014) also explains that containers package an application and its dependencies into a single unit. This includes libraries, configuration files and other necessary files. This allows applications to be deployed consistently across different environments, ensuring predictable behavior and reducing compatibility issues (Sergeev et al., 2022)

By enabling more granular and efficient use of system resources, containers contribute to lower energy usage compared to virtual machines due to their lightweight nature and faster startup times (Shirinbab et al., 2020). The adoption of containerization technologies have been shown to optimize energy costs, as containers allow

for a higher density of applications running on the same physical hardware without the overhead associated with full virtualization (Cuadrado-Cordero et al., 2018).

Docker is one of the most widely adopted container platforms, providing tools for building, shipping, and running applications in containers (Merkel, 2014). Docker containers are based on open standards and can run on various operating systems and cloud platforms (Sergeev et al., 2022). Figure 3.3 below illustrates how containers run ontop of a Docker Engine on a virtual or physical machine.

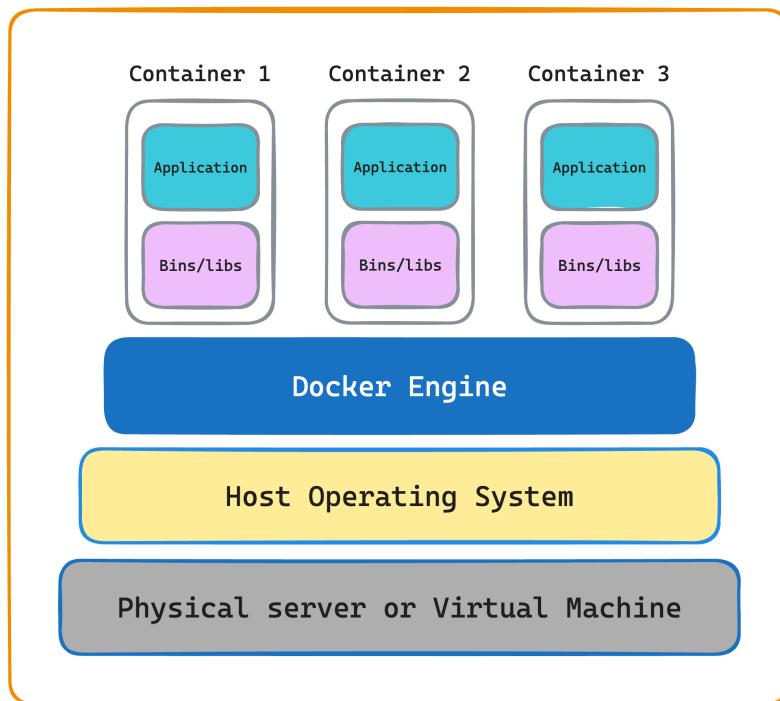


Figure 3.3: Containers running on the Docker Engine.

As the number of containers in an environment grows, managing and orchestrating them becomes increasingly complex. Container orchestration tools, such as Kubernetes and Docker Swarm, help automate the deployment, scaling, and management of containerized applications across multiple hosts (Burns et al., 2016).

These tools provide features like:

1. **Automated deployment and scaling:** Containers can be automatically provisioned, scaled up or down based on demand, and load-balanced across multiple hosts (Burns et al., 2016).
2. **Self-healing and monitoring:** Orchestration tools can monitor the health of containers and automatically restart or reschedule them in case of failures².

²<https://kubernetes.io/docs/concepts/overview/>

3. Service discovery and load balancing: Applications running in containers can be easily discovered and accessed by other services, enabling microservice architectures³.

Container orchestration has become an essential component of modern cloud-native architectures, enabling efficient management and scaling of containerized applications in dynamic environments.

3.5 Serverless Computing

Serverless computing has emerged as an alternative to traditional infrastructure management approaches, such as managing physical servers or building developer platforms as described in Section 3.4 (Baldini et al., 2017). In serverless computing, the underlying infrastructure is abstracted away, allowing developers to focus on writing code and deploying programs without the need to provision or manage servers (Baldini et al., 2017; Roberts, 2018). Castro et al. (2019) defines serverless as such:

Serverless definition

Serverless computing is a platform that hides server usage from developers and runs code on-demand automatically scale and billed only for the code running.

(Castro et al., 2019)

According to Baldini et al. (2017), in a serverless model, the cloud provider is responsible for managing the infrastructure, including resource allocation, scaling, and maintenance. This approach enables more efficient resource utilization, as resources are dynamically allocated based on demand. Furthermore, developers can deploy containers or functions independently, promoting modularity and scalability.

On top of this, severless computing offers several more benefits, including:

1. Reduced operational overhead: Developers do not need to manage servers or infrastructure, freeing up time and resources for application development (Baldini et al., 2017).

³<https://kubernetes.io/docs/concepts/services-networking/>

2. Faster time-to-market: With serverless, developers can quickly deploy and iterate on functions without the need for time consuming infrastructure setup (Adzic & Chatley, 2017).
3. Cost efficiency: Serverless platforms typically employ a pay-per-use pricing model, where users are charged based on the actual execution time and resources consumed by their applications (Eismann et al., 2021).

However, serverless architectures also introduce certain challenges, such as:

1. Potential vendor lock-in: Serverless platforms may have provider-specific APIs and services, which can make it challenging to change providers or migrate applications (Gottlieb, 2018).
2. Cold start latencies: Containers or functions that have not been invoked recently may experience longer startup times, cold starts, which can impact application performance (Golec et al., 2023).
3. Resource overhead and efficiency: Serverless platforms typically rely on containerization technologies, such as Docker, to encapsulate and isolate function executions. The use of containers can introduce resource overhead and impact the efficiency of serverless applications (Akkus et al., 2018).

Examples of widely used platforms that build on the serverless model can be found at the major cloud providers, like Google Cloud Run⁴, AWS Fargate⁵ and Microsoft's Azure Container Instances (ACI)⁶. These three example platforms allow developers to deploy containers onto the cloud without worrying about orchestration behind the scenes.

3.5.1 Functions-as-a-Service

FaaS is a cloud computing model, derived from serverless, that allows developers to execute individual functions in response to events or triggers without needing to manage the underlying infrastructure (Sewak & Singh, 2018).

Examples of these FaaS platforms among the big three cloud providers are; AWS Lambda⁷, Google Cloud Functions⁸, and Microsoft's Azure Functions⁹. On FaaS platforms like these, developers write and deploy small, self-contained functions that perform specific tasks. These functions are typically stateless and can be

⁴<https://cloud.google.com/run>

⁵<https://aws.amazon.com/fargate/>

⁶<https://azure.microsoft.com/en-us/products/container-instances>

⁷<https://aws.amazon.com/lambda/>

⁸<https://cloud.google.com/functions>

⁹<https://azure.microsoft.com/en-us/products/functions>

written in various programming languages supported by the FaaS provider (Baldini et al., 2017).

When a function is triggered, the FaaS platform automatically allocates the necessary resources to execute the function, such as CPU, memory, and network bandwidth. The platform also handles the scaling of the function based on the incoming requests, ensuring that the function can handle varying workloads by itself (McGrath & Brenner, 2017).

FaaS platforms often utilize container technology, like Docker, to provide an isolated environment for each function execution. Containers offer several benefits, such as fast startup times, efficient resource utilization, and the ability to package functions with their dependencies (Van Eyk et al., 2018). However, the use of containers in FaaS also introduce some challenges including, cold start latency and the overhead associated with container initialization and management (Wang et al., 2018).

Cold start latency refers to the time it takes for a FaaS platform to provision a new container instance when a function is invoked after a period of inactivity. This latency can be significant, especially for applications with strict performance requirements (Wang et al., 2018). The limitations of containers in FaaS environments have led to the exploration of alternative approaches, such as using Wasm and {WASI}. As Solomon Hykes, the creator of Docker, stated:

“If WASM+WASI existed in 2008, we wouldn’t have needed to create Docker. That’s how important it is. WebAssembly on the server is the future of cloud computing. A standardized system interface was the missing link. Let’s hope WASI is up to the task.” (Solomon Hykes [@solomonstrel], 2019)

This statement highlights the potential of Wasm and WASI to face the challenges with containers in FaaS and to provide a more efficient and platform-agnostic approach to serverless computing, a potential explored and supported by Kjorveziroski et al. (2022).

3.6 WebAssembly and WASI

WebAssembly, commonly referred to as Wasm, is a modern binary instruction format that has risen to prominence as a versatile technology across a diverse amount of computing environments, originating in the web browser. This section introduces the project that WebAssembly evolved from - *asm.js* - and illustrate

how WebAssembly lets developers write programs in a high-level language and run them across a multitude of platforms.

3.6.1 asm.js

Mozilla released the first version of asm.js in 2013, and it was designed be a subset of JavaScript, allowing web applications written in other languages than JavaScript, such as C or C++, to run in the browser. The intention of asm.js was to enable web applications to run at performance closer to native code than applications written in standard JavaScript. A simplified flow for how source code written in C/C++ is compiled to bytecode that can be executed in the browser can be found in figure 3.4 below.

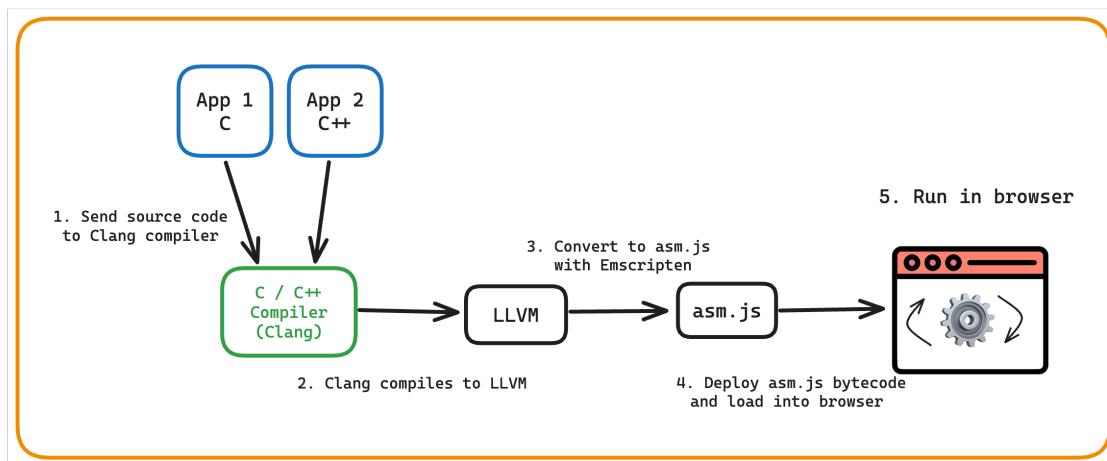


Figure 3.4: Source code in C/C++ compiled to asm.js and run in browser.

While asm.js was a great leap forward, being a subset of JavaScript limited its scope, leading to its deprecation in 2017 and the development of a more efficient and portable format (WebAssembly.org, n.d.).

3.6.2 WebAssembly

The team at Mozilla built upon the lessons learned from asm.js and went on to develop WebAssembly, launching the first public version in 2017. WebAssembly is a low-level code format designed to serve as a compilation target for high-level programming languages (Haas et al., 2017). It is a binary format that gets executed by a stack-based virtual machine, comparable to how Java bytecode runs on the Java Virtual Machine (VM) (Haas et al., 2017). In Figure 3.5 below, a simplified flow for compiling web applications written in different languages can be compiled and run inside a web browser.

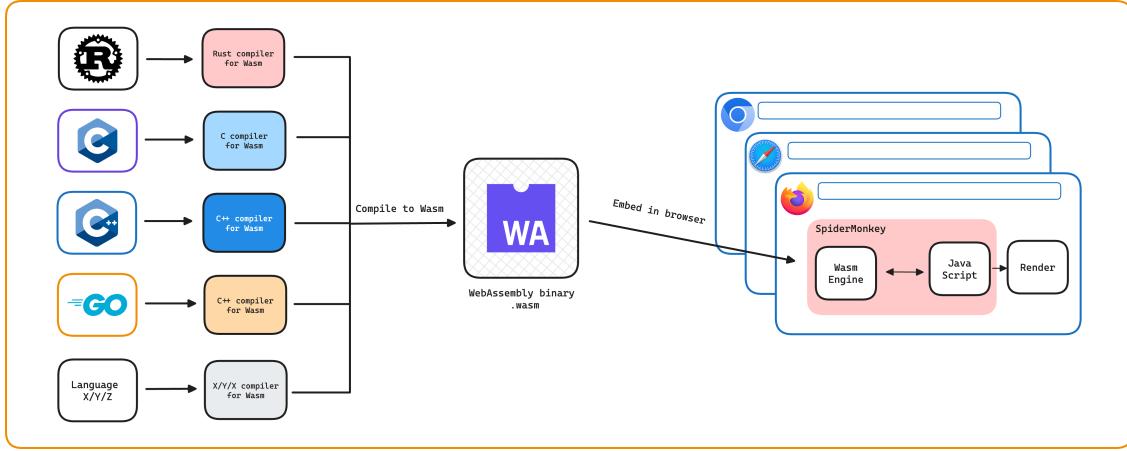


Figure 3.5: Source code compiled to WebAssembly and embedded in browser.

According to Haas et al. (2017), some key features that the team behind Wasm sought out to implement were:

Table 3.1: WebAssembly features

Feature	Description
Portability	Wasm can run on different hardware and operating systems due to its hardware-independent design and sandboxed execution environment.
High Performance	Wasm aims to achieve near-native performance by leveraging hardware capabilities and ahead-of-time compilation.
Security	Wasm modules run in a sandboxed environment, isolated from the host system, with strict memory and control flow limits, mitigating security vulnerabilities.
Modular design	As an open standard with a modular design, Wasm can be extended and integrated with various programming languages and tools, enabling broad applications beyond web browsers.
Efficient compression	Wasm's binary format is designed for efficient compression, reducing code size and improving download times, especially on resource-constrained devices.

By addressing performance, security and portability concerns, Wasm offers an alternative to traditional approaches for running untrusted code on the web and in

other computing environments, such as cloud and edge computing (Haas et al., 2017).

3.6.3 WebAssembly System Interface

While Wasm was initially designed to run in web browsers, its potential for use in other environments, such as cloud and edge computing, led to the development of the WebAssembly System Interface¹⁰. WASI is a modular system interface that provides a standardized set of functions for interacting with the host operating system, enabling Wasm modules to run outside of web browsers.

WASI defines a set of APIs for performing tasks like file system operations, networking, and other system-level operations, allowing Wasm modules to be portable across different platforms and environments. Its first iteration aimed to be implement as many POXIS-like features as possible, but has since extended beyond this and in their latest Preview 2, they have implemented support for websockets and HTTP interfaces (“WASI/Preview2/README.Md at Main · WebAssembly/WASI”, n.d.).

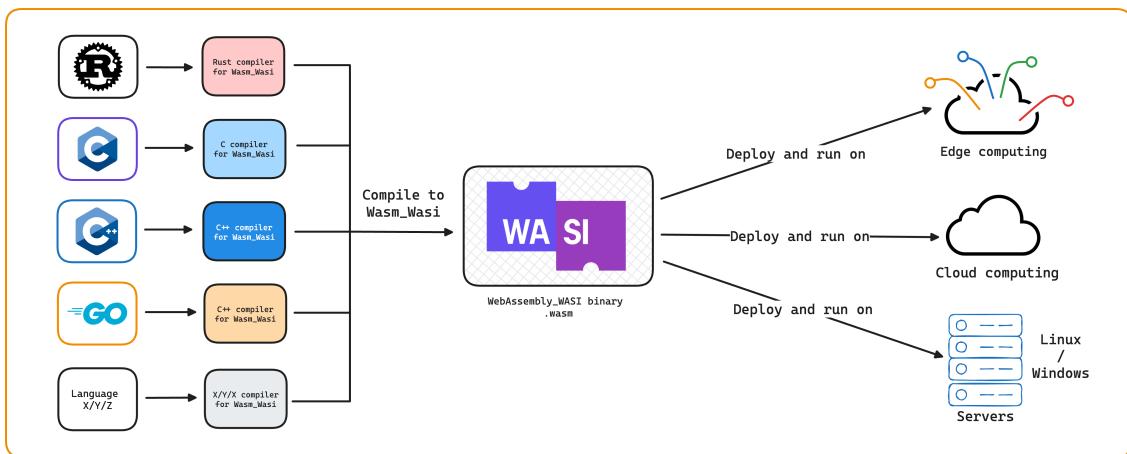


Figure 3.6: Source code compiled to `wasm_wasi32` and deployed on platforms that support running the binaries.

3.6.4 WebAssembly Runtimes

WebAssembly modules cannot run independently; they require a runtime environment to interpret and execute them. Several WebAssembly runtimes have been developed to support the execution of WebAssembly modules in different environments, such as Wasmtime, Wasmer, and WasmEdge (Zhang et al., 2024).

¹⁰<https://wasi.dev/>

Zhang et al. (2024) explains that these runtimes provide a secure and efficient execution environment for WebAssembly modules, enabling them to run on a wide range of platforms, including cloud servers, edge devices, and even embedded systems. They explain that some runtimes offer features like ahead-of-time (AOT) compilation, which can further improve the performance of WebAssembly modules. Figure 3.7 below illustrates how code written in a programming languages with support for compiling to Wasm can run cross-platform.

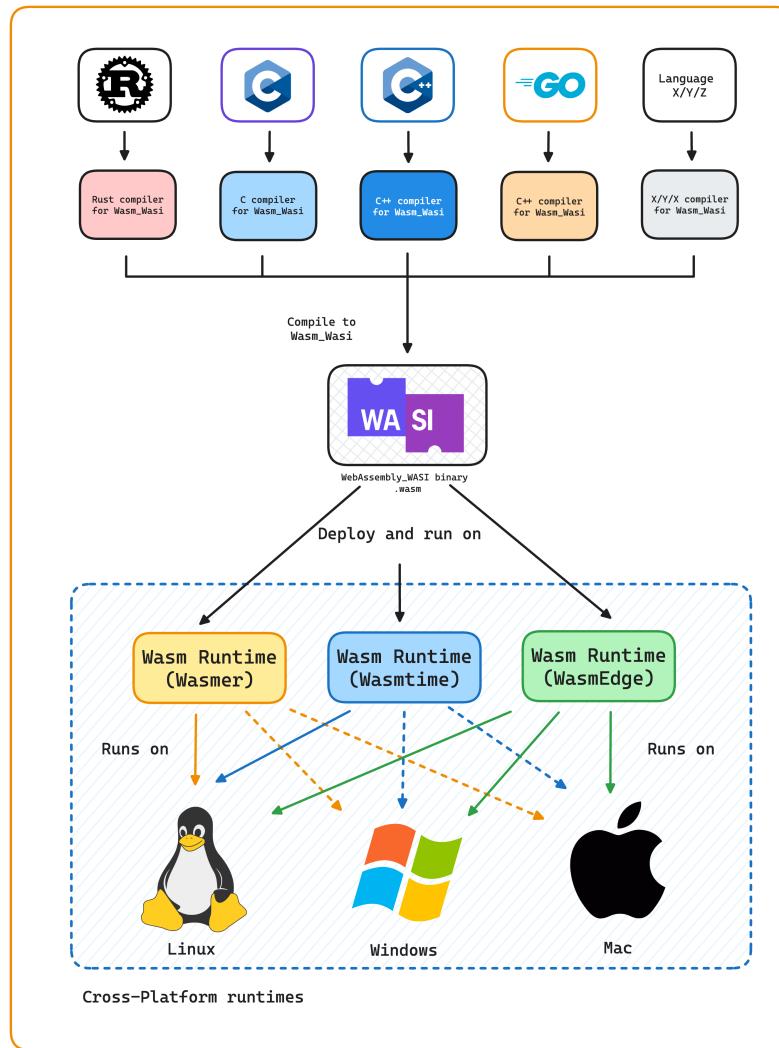


Figure 3.7: Source code compiled to `wasm_wasi32` that can run anywhere a WebAssembly runtime can be installed.

The combination of WebAssembly, WASI, and efficient runtimes has sparked interest in using WebAssembly as an alternative to traditional containerization technologies, such as Docker, in cloud-native and serverless environments (Sebrechts et al., 2022; Shillaker & Pietzuch, 2020).

3.7 Energy Monitoring

Monitoring and measuring energy consumption can be useful for understanding and optimizing the energy efficiency of cloud computing environments. Energy measurements provide insights into the impact of different technologies, architectures, and workloads on overall energy usage (Al-Fuqaha et al., 2015; Shehabi et al., 2016).

Several protocols and techniques have been explored to collect and analyze energy consumption data, and this thesis will explore some of these. These protocols enable the transmission of energy data, which can be used for optimizing energy efficiency (Al-Fuqaha et al., 2015).

3.7.1 MQTT

The MQTT protocol is a lightweight and efficient protocol that has been used for energy monitoring. In an MQTT setup, a broker server acts as an intermediary, facilitating communication between devices and servers. This enables efficient data exchange between publishers and subscribers (Al-Fuqaha et al., 2015). (See Figure 3.8).

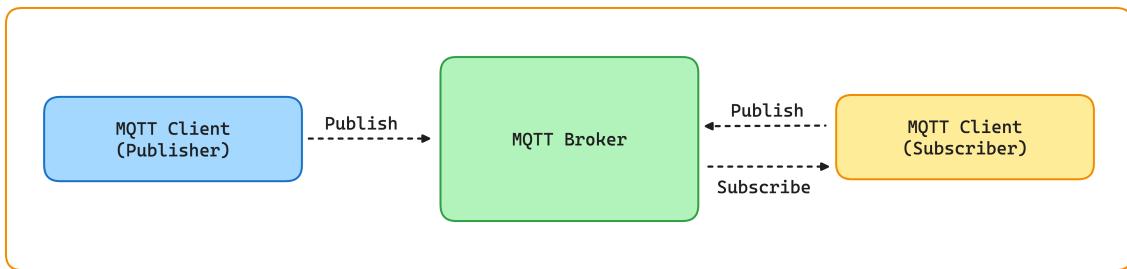


Figure 3.8: PubSub model of the MQTT protocol.

3.7.2 Z-Wave

Z-Wave is a wireless communication protocol designed for home automation and energy management. It has been used in research studies to develop energy monitoring systems for residential buildings, enabling monitoring and control of energy consumption (Al-Fuqaha et al., 2015).

3.7.3 Aeotec Smart Switch

The Aeotec Smart Switch 6 is a power switch that is primarily used for home automation. It allows its users to integrate with it through a Aeotec Z-Stick what

communicates with the switch through Z-wave. A supporting library; `zwave-js-ui`¹¹, which can be setup as a Docker container on a computer, can read energy measurements from the Smart Switch at a given interval and publish the values on a MQTT-topic. This is in turn picked up by the MQTT broker and published to its subscribers. This switch supports reporting measurements at an interval of 1 second (Aeotec, n.d.).

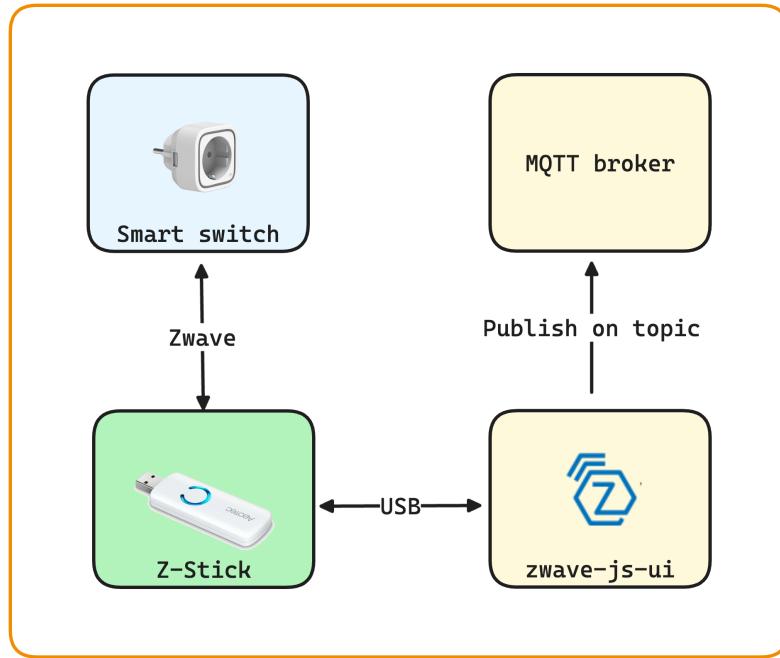


Figure 3.9: Smart Switch 6 communicating with `zwave-js-ui` through Aeotec Smart Stick.

3.7.4 Modbus TCP

An alternative to reading energy measurements through a pub-sub model with MQTT, is to utilize the Modbus TCP protocol. Modbus TCP is a widely adopted industrial communication protocol for exchanging data between devices and control systems. It has been used to develop energy monitoring and auditing systems for industrial applications, facilitating energy audits and energy-saving strategies (Tong et al., 2015). It communicates with TCP/IP packets, and an illustration on how a client-server request and response could look like can be found in Figure 3.10.

¹¹<https://github.com/zwave-js/zwave-js-ui>

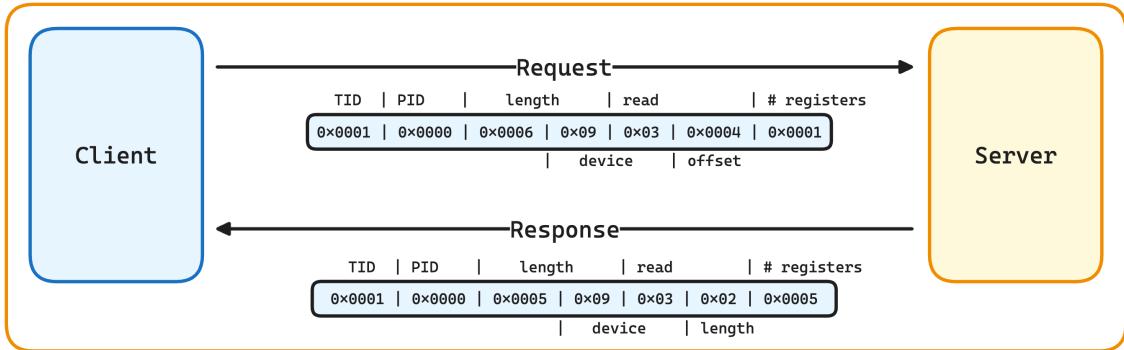


Figure 3.10: TCP/IP package exchange between a client and a server over Modbus TCP.

3.7.5 Gude Expert Power Control 1105

The Gude Expert Power Control 1105 is a switched PDU with an integrated current metering and monitoring that supports communicating over TCP/IP. It has a server that can measure energy, current, power factor, phase angle, voltage, and active / apparent / reactive power. To read these measurements, the PDU supports a handful of interfaces, such as REST API, HTTPS, SNMP, Telnet, MQTT and Modbus TCP (GmbH, 2023). In Figure 3.11 below, a simplified setup with the Gude acting as the server for communicating over Modbus TCP is shown.

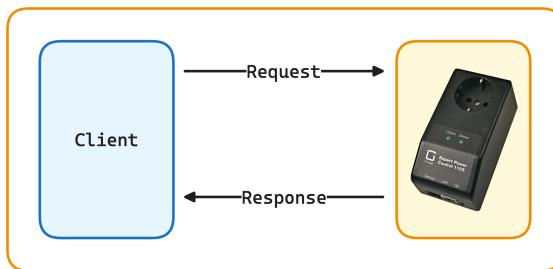


Figure 3.11: Client communicating with Gude over Modbus TCP.

This device can measure a lot of different sensor, including current, voltage, power draw and power factor. Voltage is the amount of volts the system has access to, typically 240V in Norway on AC. The current is the amount of electricity measured in ampere that the system is consuming. Power draw is a product of current and voltage, $P = U \times I$, while power factor measures between 0 to 1 and describes the amount of power consumed spent on actual usage. For example, a power factor of 0.5 means that 50% of the power consumed in a circuit goes to waste. Most data centers aim to get as close to a power factor 1 as possible, to make most of the energy they pay for not go to waste (Rasmussen, 2018).

Part II

Project

Chapter 4

Methodology

The only true wisdom is in knowing you know nothing.

—Socrates

To reiterate, the goal of this thesis is to investigate the problem statements presented in Section 1.3. We are going to develop a FaaS prototype, named Nebula, which we are going to use for conducting experiments to see if the claim that Wasm modules offer a more efficient way to develop and deploy our cloud native applications holds true. On this platform we are going to deploy functions written in Rust that are either compiled to Wasm modules or compiled to Rust binaries that get packaged into a Docker image.

This chapter will present the intended methodology on how we will conduct our experimental research method. We will employ an experimental research method, where we will build a prototype, perform experiments on it in a controlled environment, collect measurements and compare the results.

4.1 Experimental Framework

The focus of this thesis is to build a prototype and perform experiments on it, to measure the capabilities of Wasm modules as a model for deploying and running cloud applications. An experimental approach was chosen to test the problem statements, as we can have controlled manipulation of variables, measure the outcomes, and compare directly between two deployment environments; Wasm modules and Docker containers.

We will perform a series of controlled experiments designed to measure startup latencies (cold start), total runtime, and energy consumption for each invocation

of a function. This allows for a precise evaluation of how Wasm modules compare against traditional container-based deployments in the context of cloud-native applications.

To capture energy measurements, we will need to use hardware locally that we can control. For this we will look at the protocols mentioned in Section 3.7, and use a Raspberry Pi as our local deployment environment. We will also capture startup and runtime data on these invocations, but not many cloud applications today are deployed on Raspberry Pi's, so we will also need to deploy Nebula on a more typical setup.

4.1.1 Prototyping

The prototype we are going to build for this project will be named Nebula, named after the celestial phenomenon that can be observed in the form of a giant cloud out in space. On Nebula we will develop a way to deploy and invoke functions in the form of Wasm modules or as spun up Docker containers. To perform the desired experiments, the prototype will need to be able to deploy to both a Raspberry Pi 4b, and on a typical machine hosted on the cloud. For this, we will be using NREC¹, an IaaS platform available to students, where we can spin up a Debian virtual machine and set up a service on a public IP address.

4.1.1.1 Hardware specifications

Experiments will be performed on two types of hardware:

1. Raspberry Pi: A small Single board computer (SBC)s with an ARM64 architecture that can run flavors of Linux. This device will be used to measure energy consumption relative to function invocations, which it is well suited for, as it consumes small amounts of energy by itself, making it ideal for comparing power consumption under load (Bekaroo & Santokhee, 2016).
2. Virtual Machine: A Debian-based VM running on an IaaS platform (e.g., NREC), to measure performance relative to how a typical cloud native application would perform.

¹<https://nrec.no/>

4.1.1.2 Nebula Requirements

Based on this, we can summarize what we need Nebula to be able to meet these requirements.

Table 4.1: Nebula requirements

Requirement	Description
Function deployment	We should be able to deploy functions to Nebula.
Function execution	The prototype should execute functions in response to user invocations, either as Wasm modules or as Docker containers.
Measurement capabilities	The prototype should be able to collect data for startup latencies, total runtime, and energy consumption for each function invocation.
Portability	The prototype should be deployable on both a Raspberry Pi for local testing and measurement, and on a Debian virtual machine hosted on an IaaS platform (e.g., NREC) for a more typical cloud setup.
Performance	The prototype should be built in such a manner that it presents little overhead when running our functions, isolating the resulting efficiency and energy readings to each function invocation.

4.1.2 Controlled Experimentation

Controlled experimentation will be an important aspect of this research, as it let us isolate and study the effects of the deployment environment on performance and energy consumption. Various factors, such as input parameters for the benchmark functions and hardware configurations, will be controlled or kept constant across experiments to minimize the influence of external factors and increase the validity and reliability of the results.

To experiment on Nebula, we will craft a set of benchmark functions. As the input value to these functions scale up, the computational load on the server will increase, letting us measure the execution time across different workloads.

4.1.3 Benchmarking

Benchmarking will be employed as a research method to evaluate the performance of different configurations. Benchmark functions representing various computational workloads will be implemented and executed on both environments to assess the suitability of the FaaS prototype for diverse application scenarios.

4.1.4 Benchmark Functions

A total of four benchmarking functions will be selected for testing, representing various computational workloads. These functions will be written in Rust and compiled to both Wasm modules and Rust binaries packaged in Docker images. The selection of these functions is based on their ability to stress different aspects of Nebula, including function invocation mechanisms, memory management, computational intensity, and handling of large datasets.

The functions we will implement are:

1. *Fibonacci*: This function calculates the n th Fibonacci number in the sequence, following the formula: $F_n = F_{n-1} + F_{n-2}$. The recursive nature of the Fibonacci function involves a significant number of function calls and memory accesses, making it a relevant benchmark for evaluating the overhead and efficiency of function invocations and memory utilization in the FaaS prototype.
2. *Exponential*: This function calculates the value of Euler's number raised to the n th power, following the formula: $F_n = e^n$. This benchmark tests the performance of the prototype in handling intensive floating-point calculations, which are common in scientific computing and data analysis applications.
3. *Factorial*: This function calculates the sum of the factorial of n , following the formula: $n! = \prod_{k=1}^n k$. Similar to the Fibonacci function, the factorial function involves recursive computations and can stress the function invocation and memory management aspects of the FaaS platform.
4. *Prime number*: This function finds the n th prime number using a brute-force approach, looping through a vast range of numbers (0 through $2^{64} - 1$ to check for primality. This benchmark stresses the CPU and memory capabilities of Nebula by simulating computationally intensive workloads involving large datasets and long-running calculations, commonly found in areas like cryptography.

Choosing what functions to implement is not an easy task. While these functions are relatively simple in nature, and software applications are commonly more

complex, these four functions still provide a well-rounded evaluation by stressing various aspects of Nebula. With these functions we can test computational intensity, memory management, and handling of large datasets. The selection balances comprehensiveness with practicality, letting us perform a thorough evaluation within a reasonable scope.

4.1.5 Measurement and Data Collection

To quantify the efficiency and energy consumption associated with function invocations on each environment, measurement techniques will be employed. We will deploy Nebula, along with the Wasm modules and Docker images, to our target computers. For measuring startup and runtimes of each function like a typical setup, we will deploy Nebula and the functions to a virtual machine running Debian on NREC. While a Raspberry Pi 4 will serve as our deployment target for measuring startup latencies and runtimes of a device more akin to embedded systems, but more importantly, allowing us to measure energy consumption on a controlled device.

The Raspberry Pi will be connected to a power supply that can report energy readings, enabling the collection of energy consumption data during function execution. See Figure 4.1 below for a rough sketch on how we will set up benchmarking and energy measurements for our experiments.

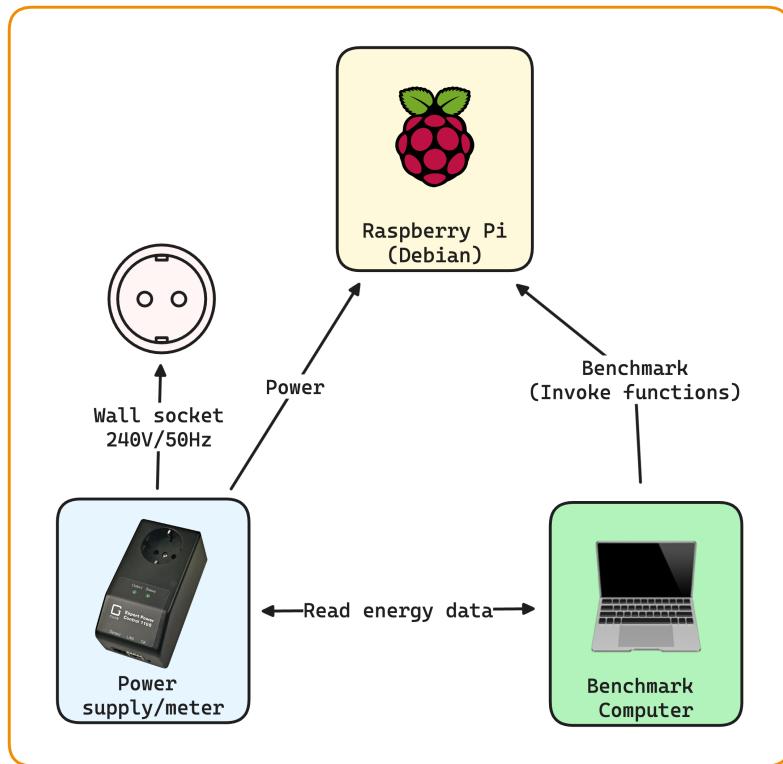


Figure 4.1: Reading energy consumption of our Raspberry Pi.

As we will discuss in Section 6.6.5, we will opt for using the Gude Expert Power Control 1105 for measuring our energy readings. Table 4.2 below shows a sample of sensors we can read from this device. To read from each sensor we need its address represented in hex values as we will discuss in more detail Section 6.6.2.

Table 4.2: Gude Expert Power Control 1105 sensors.

Offset	Sensor.Field	Address	Unit	Inaccuracy
0	Absolute Active Energy	0x400	Wh	< 1.5%
1	Power Active	0x402	W	< 1.5%
2	Voltage	0x404	V	< 1%
3	Current	0x406	mA	< 1.5%
4	Frequency	0x408	0.01hz	< 0.03%
5	Power Factor	0x40a	0.001	< 3%

As we can see from Table 4.2, we can read *Absolute Active Energy*, a value that is the cumulative sum of all energy consumed on our device, measured in Wh. *Power Active* reports the current W being drawn at the moment of reading the sensor, the product of *Voltage* and *Current*. At first glance, these two sensors seem ideal for our benchmarkings, but as we will reveal in Chapter 7, the energy it takes to invoke a function as either a Wasm module or Docker container, will be measured in μ Wh. Measurements in whole Watts will not provide accurate enough data for our experiments. The three most important sensors for our experiments will be *Voltage*, *Current* and *Power Factor*.

Throughout the experiments, we will collect various data points to evaluate the performance and energy efficiency of both Wasm modules and Docker containers. As we can see in Table 4.2 there will be some measure of inaccuracy involved with our experiments, with the most impactful one being the power factor at < 3%. We will collect multiple measurements for each permutation of our function invocations which should help account for this inaccuracy, but it is still important to keep in mind.

To isolate how much power our functions invoke on top of the idle load of the underlying system that hosts Nebula, we will also benchmark a baseline power level, to measure how much power Nebula consumes while idle. Any power consumed above this should then reflect how much power each function invokes when called.

The following measurements will be recorded on both hardware configurations:

- Startup latency / Cold start: For each function invocation, the time taken from the initial request is received until the function starts executing will be measured. This startup latency, known as cold start, is a critical metric for evaluating the responsiveness and scalability of serverless function deployments.
- Total runtime: The total execution time of each function will be measured, from the same time we start measuring cold start until the function returns the result will be measured. This metric will let us compare the computational efficiency of each environment under different workloads.

Additionally, the following measurements will be gathered on the Raspberry Pi system, as shown in Figure 4.1:

- Power consumption: During the function execution on the Raspberry Pi, we will measure the power the server consumes at regular intervals. To get a more accurate reading, we will measure the current (I) in mA and the power factor (PF), while assuming the voltage (U) is a constant 240V. Using these numbers we will calculate the power with the following equation:

$$P = U \times I \times PF$$

Two sets of power consumption data will be collected.

- Average Power: The average power consumption of the Raspberry Pi device, including any background processes, including Nebula, or idle draw, while invoking a function. This will be an estimation based on energy readings measured during the lifetime of each function, and calculating the average of these. We will turn off wifi, bluetooth and hdmi on our device to minimize external factors that typically draw extra power.
- Energy Consumption: Based on the power consumption measurements and the total runtime durations, the energy consumption for each function invocation will be calculated. We will calculate the energy consumption with the following formula:

$$E_{Wh} = P_W \times t_h$$

Two sets of energy consumption values will be derived:

- Energy consumption: The total energy in Wh consumed of the entire Raspberry Pi device during function execution.

This data will then be saved in a JSON format and used for analysis.

4.1.6 Data Analysis

The collected data on startup latency, total runtime, power load and energy consumption will be analyzed using statistical techniques to identify patterns, trends, and potential correlations. Visualization techniques, including line graphs, scatter plots, and bar charts, will be used to present the data effectively and make it easy to interpret the data. These visual representations will make it easier for us to identify relationships between the metrics, such as impact of input size on runtime and energy consumption.

4.1.7 Comparative Analysis

A comparative analysis will be performed to evaluate the differences in performance and energy efficiency between Wasm modules and Docker containers. The following comparisons will be made:

- Startup latency: The startup times (cold starts) of Wasm modules and Docker containers will be compared across various input sizes to assess the responsiveness and scalability of each deployment method.
- Total runtime: The total execution times of functions deployed as Wasm modules and Docker containers will be compared to evaluate their computational efficiency under different workloads.
- Power load: The average power loads during function execution will be compared between Wasm modules and Docker containers to identify potential differences in power requirements.
- Energy consumption: The estimated energy consumption for each function invocation will be compared across Wasm modules and Docker containers to evaluate their relative energy efficiency.

4.1.8 Reliability and Validity

To ensure the reliability and validity of the results we will take some measures. Controlled experimentation will be employed, where our experiments will be conducted in a controlled environment with variables such as hardware configurations and input parameters kept constant across runs. This minimizes the influence of external factors on the results.

We will replicate each experiment multiple times to ensure the consistency and reproducibility of the findings. The data collection process will be automated

using a benchmarking utility to minimize human error and ensure consistent measurement techniques across all experiments.

The results obtained from this research will be compared with findings from existing studies in the field to help validate the observations made in this study.

By implementing controlled experimentation, replication, automation, and comparison with existing research, the reliability and validity of the experimental results will be enhanced. Implementing these measures will increase our credibility and validity of our findings from this study.

Chapter 5

Designing Nebula

The most dangerous phrase in the language is “We have always done it this way”

—Grace Hopper

Based on the requirements outlined in Section 4.1.1.2, we will design a prototype named Nebula, which is a serverless FaaS platform capable of executing deployed functions as both Wasm modules and Docker containers. The primary goal of Nebula is to serve as an experimental testbed for evaluating the performance and energy efficiency of Wasm modules to traditional container-based deployments in the context of cloud-native applications. On top of this, we will need supporting infrastructure and utilities for benchmarking our prototype.

5.1 Prototype Scope and Objectives

While popular serverless platforms offered by major cloud providers are typically closed-source, meaning how the work behind the scenes is hard to evaluate from the outside, this prototype focuses on the aspect of function deployment and execution of FaaS. This controlled approach lets us concentrate on an evaluation of deploying and running functions as Wasm modules versus Docker containers.

As a student thesis project, the prototype’s scope is inherently limited. We are not going to implement features such as automatic scaling, distributed function deployments, authentication and authorization, integrations with file systems and databases, and more. However, this focused design is intentional, as we want to minimize other factors, and rather hone in on a controlled comparison between the two deployment methods.

5.2 System Architecture

We will design Nebula to be a FaaS platform that lets us execute functions deployed as both Wasm modules and Docker containers. The system architecture will consist of several key components that work together to facilitate the development, deployment, and execution of functions, as well as the benchmarking and evaluation of their performance and energy efficiency.

1. *Nebula core*: The central component of the prototype, responsible for handling incoming requests from users, spinning up deployed functions, and executing them. Nebula will support both Wasm modules and Docker containers, providing an interface for interacting with the deployed functions.
2. *Function development and deployment*: A separate environment where example functions can be developed, either compiled to Wasm modules or packaged into Docker images, and then uploaded to Nebula.
3. *Infrastructure*: Two environments will be used for performing the experiments: a Raspberry Pi, a low-power ARM-based SBC serving as the deployment target for measuring energy, and a Debian-based virtual machine hosted on an IaaS platform to simulate a typical cloud-native application deployment environment.
4. *Energy measurement setup*: A power supply capable of reporting energy readings will be used to supply the Raspberry Pi with power, allowing the collection of energy consumption data during function execution.
5. *Benchmarking utility*: A separate benchmarking application for performing our experiments. We will develop this to let us create a benchmark suite that can be reused and test a variety of variables. This utility will be responsible for reading energy measurements from our power supply, and consolidate energy readings with function executions.

Figure 5.1 illustrates the overall architecture of our Nebula prototype system, including benchmarking and energy measurements for our experiments.

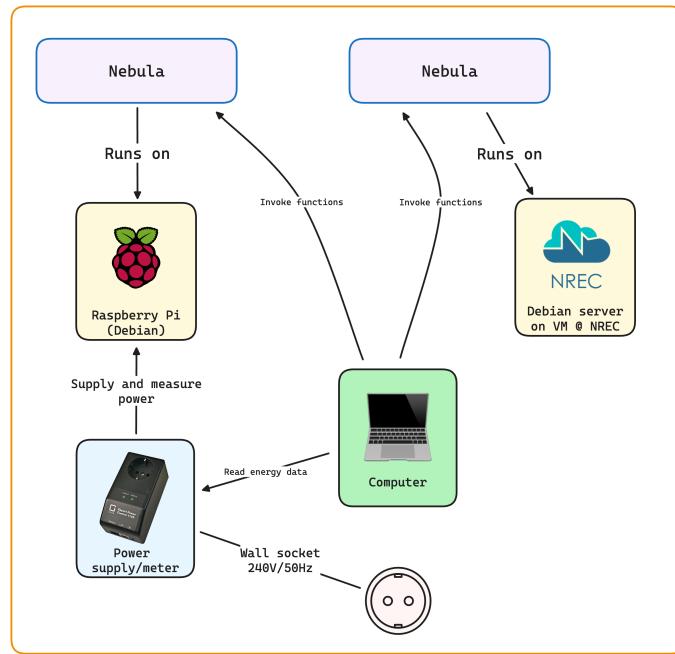


Figure 5.1: Overall architecture of our benchmarking setup.

5.3 Nebula

Nebula will serve as the core component of the prototype. It will handle incoming requests, manage the lifecycle of deployed functions, execute them, and measure startup and runtimes of each execution. We will develop an interface for users to interact with the deployed functions, either as Wasm modules or Docker containers, like illustrated in Figure 5.2.

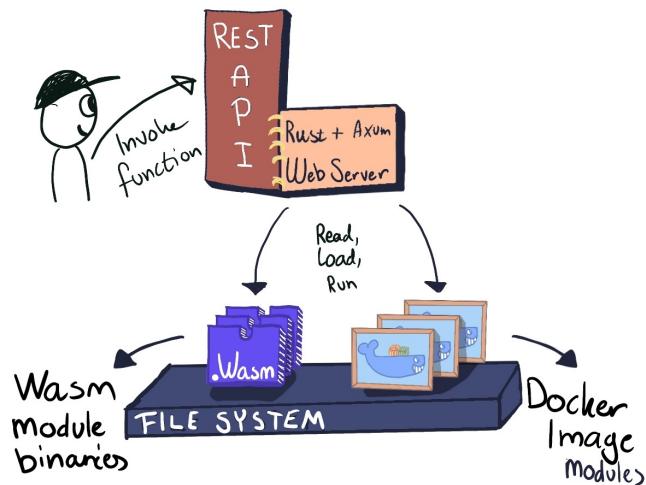


Figure 5.2: User interacting with Nebula.

Nebula's function execution process will involve several steps, as illustrated in

Figure 5.3. When a user invokes a function through the interface, Nebula receives the request and routes it to the corresponding function based on the specified deployment method. Nebula will then start the function, execute it with the provided input data, and return the result back to the user.

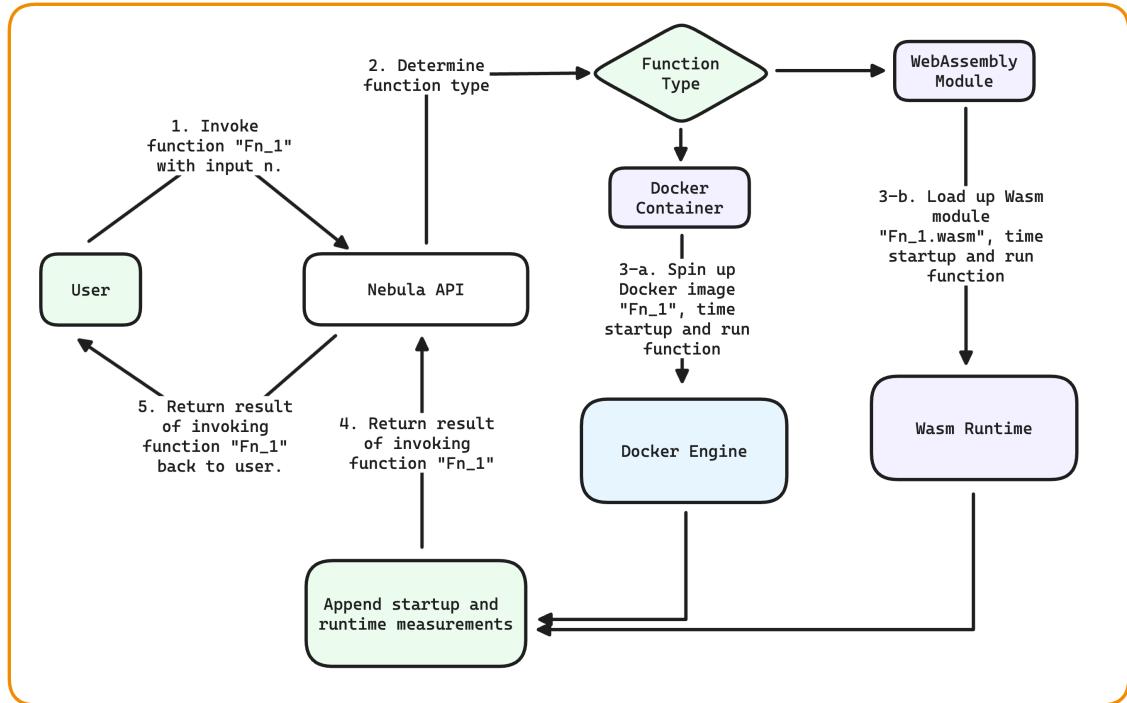


Figure 5.3: Flow of a user invoking a function on Nebula.

This function execution process will also include collecting metrics for each function invocation. We will start a timer at the start of our function execution that first times when our function is loaded into memory and starts executing, and then takes the time again at the end of the function. These metrics will be added to our function response and provide accurate measurements on how fast our functions start up and how much time is attributed to the function running itself.

5.3.1 Function Packaging and Deployment

Our development environment will need to support packaging and deploying functions as WebAssembly (Wasm) modules or Docker images. For Wasm, the compiled binary will be transferred to the server through SCP (Secure Copy Protocol). For Docker, a Dockerfile specifying the image build instructions will describe how Docker should package our function, before we gzip the image and upload it to the server through SCP.

For Wasm modules no further steps are required, as Nebula will handle loading the binary into the Wasm runtime environment. For the Docker images we will need

to inform Docker on the server that we want to load our image into its repository, making it available to run as a container. This process will ensure that the functions are ready to be invoked and executed when needed.

5.3.2 Function Execution

To run functions we will need to provide the proper environments they can run inside of. For Wasm modules, we will need a way to include a Wasm runtime, as mentioned in Section 3.6.4, which will let us run our functions in a sandboxed environment. The runtime will handle the loading and instantiation of the modules.

For Docker containers, we will use the containerization capabilities of Docker. Each function will run in its own isolated container, with the Docker Engine runtime managing its startup and shutdown.

For both methods, we will need a way to pass the users input to the function as a parameter and receive the result of each function invocation. The chosen way to pass input and receive output should be the same for each environment, to ensure controlled experimentation.

5.3.3 Interface

To serve as an interface for the user to invoke our functions, we will also develop a web server that will expose end points that our users can communicate with over HTTP in a RESTful manner. In other word, we are going to develop a web server that exposes a REST API, where users can send requests that result in function invocation based on the request body.

5.4 Infrastructure

For our experiments, we are interested in finding out two things: how efficient is Wasm modules compared to Docker containers, and how much energy can we potentially save by shifting to Wasm. To measure a real case scenario for how much more efficient our function executions can be, we will deploy Nebula where the digital services normally are: on the cloud. This will let us measure realistic startup and runtimes for our function invocations. However, we do not have access to the power socket that the data center we will use for hosting Nebula on, so we can not measure the power consumed by our functions. Therefore we will target two environments for our experiments on Nebula:

- Virtual Machine: A Debian-based virtual machine hosted on an IaaS platform

will be used to simulate a typical cloud-native application deployment environment. Nebula will be deployed on this virtual machine, and the functions will be executed in this environment to assess their performance in a cloud-like setting.

- Raspberry Pi: A low-power, ARM-based SBC will serve as the deployment target for measuring energy consumption during function execution. The Raspberry Pi will run Nebula and the deployed functions, and it will be connected to the energy measurement setup to collect power consumption data.

5.5 Energy Measurements

To measure the energy consumption of our functions executed on the Raspberry Pi, a power supply capable of reporting energy readings will be used. The Raspberry Pi will be connected to the power supply and provide energy readings continually over the lifetime of our benchmarks. In Section 3.7 a couple of devices and protocols were presented that will be relevant for our system. We will attempt to set up energy measurements with each configuration, and see which best matches our desired outcome.

The collected energy measurements data will be stored in a fitting data format and parsed for analysis and comparison between Wasm modules and Docker containers.

5.6 Benchmarking utility

We want to minimize the amount of computing power attributed to our prototype as much as possible. With the exception of invoking the functions from a client, we want to isolate the computing power used to run our functions as much as possible. Because of this, we will use a separate computer for performing benchmarking. While there exists tools out there that could help us achieve parts of this, a custom application will be built that can both interface with the server through the REST API, and interface with the power supply over an appropriate protocol as we saw in Figure 5.3.

The utility will be responsible for invoking the functions deployed on Nebula, both as Wasm modules and Docker containers, and collect our metrics. The collected data will be stored and organized in a JSON format, and then analysed using the statistical techniques mentioned in Section 4.1.6.

5.7 Experiment Design

To evaluate the performance and energy efficiency of Wasm modules compared to Docker containers, we will perform a series of experiments on our prototype. The experiments, as described in {sect:experiments}, will focus on measuring startup times, runtimes, power measurements and calculate the resulting energy consumption. Each experiment will involve deploying the same function as both a Wasm module and as a Docker container on Nebula. The functions will be invoked using the benchmarking utility, which will collect the relevant metrics for analysis.

The experiments will be performed on both the Raspberry Pi and the virtual machine environments to assess the performance in different deployment scenarios. Both environments will provide time usage measurements in the form of startup and runtimes. The Raspberry Pi setup will additionally give us insights into the energy consumption of our functions, while the virtual machine will represent a cloud-native deployment environment, with a more powerful CPU and more memory, like we would expect of our cloud servers.

The data we end up with from the experiments will be analyzed and compared to evaluate the performance and energy efficiency of Wasm modules versus Docker containers. We will perform multiple invocations of the same set of independent variables; function, module type and input, and derive median, mean and average of each invocation to strengthen our findings.

Finally, the results of our experiments will be visualized and discussed, relating the findings to the problem statements posed for this thesis. Here we will highlight the potential benefits and trade-offs of using Wasm modules for serverless function deployment compared to traditional Docker containers, if the hypothesis we set out to explore ends up holding true.

Chapter 6

Implementing Nebula

Talk is cheap. Show me the code.

—Linus Torvalds, creator of Linux

As mentioned in the acknowledgments at the start of this thesis, the idea for this project came from a podcast episode where the CEO of Fermyon, Matt Butcher, told the journey of how they ended up developing their cloud platform to offer a FaaS powered by Wasm modules. They have built a cloud application platform for Wasm-based serverless functions named Fermyon Cloud, which is supported by their developer framework, Spin. Spin is a framework for building Wasm-based functions that can be deployed on their cloud, and this setup acts as an inspiration for the prototype we are going to build. Both projects are open-sourced and available on Github, so we can take a look at how they have implemented their services to run Wasm modules in a serverless environment.

In this chapter, the code snippets are presented in a condensed form, focusing on the parts relevant to our experiments. There is also an interactive GUI built with HTMX and Askama HTML templating which the user can use to manually interact with each deployed function, but this aspect is more a novelty for manual testing, rather than a feature that supports this thesis. A lot of the source code will not be shown, we are going to focus on the code that lets us run functions as either Wasm modules or Docker containers. The souce code can be found on Github at the public repository for this project¹, for closer inspection.

¹<https://github.com/brehen/nebula/tree/v1.0>

6.1 Choosing a Tech Stack

The programming language Rust was chosen to build our prototype. This choice was driven by Rust's strong relationship with Wasm and its suitability for building efficient and scalable applications. Developed by Mozilla Research, the same company that developed Wasm, Rust emphasizes memory safety, concurrency, and parallelism while combining high performance with strong safety guarantees. Its syntax is somewhat similar to TypeScript, allowing developers write low level programs in a high-level manner

6.1.1 Rust and WebAssembly

Rust's close relationship with Wasm stems from their shared origin at Mozilla Research, where both technologies stems from. Rust's ecosystem offers a wide range of crates (Rust's naming convention for libraries) for various Wasm runtimes, enabling us to ship the entire Wasm runtime required for running our Wasm modules as part of the web server itself. This approach eliminates the need to install a separate runtime on our server, simplifying the prototype development.

6.1.2 Wasmtime

For our Wasm runtime, we opted for Wasmtime². Wasmtime offers seamless integration with Rust through the official wasmtime crate and provides portability across platforms, including ARM for our Raspberry Pi deployment. The choice for Wasmtime was inspired by Fermyon's choice to build their platform on it. A viable alternative is Wasmer³, which can also be installed as a crate and shipped with a Rust application.

6.1.3 Docker Environment

For our Docker environment, we installed Docker Engine on both our development computer and each deployment target. Docker Engine supports Debian on ARM64 for our Raspberry Pi deployment. It also supports Debian on amd for our virtual machine. We implemented a Rust function to spin up a command for running our Docker images during execution.

²<https://wasmtime.dev/>

³<https://wasmer.io/>

6.2 Nebula Core

In this section we will describe how we built the core of Nebula, with the following components:

1. **Web Server:** A RESTful API server written in Rust and utilities used by both the web server and function runners. built using the Actix Web framework, responsible for receiving function invocation requests and managing the execution of functions, serving as the interface for our users.
2. **Function Runners:** Separate modules for executing functions deployed as Wasm modules or Docker containers, encapsulating the logic for invoking and managing the respective execution environments.
3. **Shared Library:** A common library containing shared types the function runners.

6.2.1 Building the Web Server

The web server serves as the entry point for Nebula, exposing a RESTful API for function invocation requests. We built it using the Actix Web framework, which leverages the asynchronous Rust runtime Tokio for handling concurrent requests.

```
// Include crates
use axum::{routing::post, Router};
use std::net::{IpAddr, Ipv4Addr, SocketAddr};
use tokio::net::TcpListener;
// Supporting library for our function execution
use nebula_server::api::call_function::call_function;
// Entry point, Tokio main function
#[tokio::main]
async fn main() {
    let router =
        Router::new().route("/invoke_function", post(call_function));
    // Precompile and serialize our Wasm modules, more on that later.
    serialize_modules();
    let localhost = IpAddr::V4(Ipv4Addr::new(127, 0, 0, 1));
    let listener = TcpListener::bind(&SocketAddr::new(localhost, 3090))
        .await
        .unwrap();
```

```
        axum::serve(listener, router.into_make_service()).await?;  
    }  
  
The main entry point defines a router with a single POST endpoint /invoke_function,  
which is bound to the call_function handler. This handler function determines  
the requested function type (Wasm or Docker) and routes the request to the  
corresponding function runner module.
```

6.3 Calling Functions

To handle function invocation requests, Nebula defines a set of types to represent the request parameters, function results, and associated metrics. This is stored in the shared library and lets us reuse the same types for our runner functions and the web server implementation.

```
mod serialize_wasm;  
  
enum FunctionType {  
    Docker,  
    Wasm,  
}  
struct FunctionRequest {  
    function_name: String,  
    input: String,  
    module_type: FunctionType,  
}  
struct Metrics {  
    startup_time: u64,  
    total_runtime: u64,  
    start_since_epoch: u64,  
    end_since_epoch: u64,  
}  
struct FunctionResult {  
    func_type: FunctionType,  
    func_name: String,  
    input: String,  
    result: String,
```

```
    metrics: Metrics,  
}
```

The `call_function` handler function is responsible for dispatching the request to the appropriate function runner based on the requested function type.

```
use axum::{http::StatusCode, response::IntoResponse, Form};  
use nebula_lib::{docker_runner::run_docker, wasm_runner::run_wasm};  
use serde::Serialize;  
  
// Define response body, that can be serialized to JSON  
#[derive(Serialize)]  
struct Response {  
    result: FunctionResult,  
}  
  
// Call function that takes the request as a form and invoke the correct type  
async fn call_function(  
    Form(request): Form<FunctionRequest>,  
) -> impl IntoResponse {  
    let result: FunctionResult = match request.function_type {  
        FunctionType::Docker => {  
            let docker_function = format!(  
                "nebula-function-{}-{}",  
                request.function_name, request.base_image  
            );  
            run_docker(  
                &docker_function,  
                &request.input,  
                request.function_name,  
                request.base_image,  
            )  
            .unwrap()  
        }  
        FunctionType::Wasm => {  
            let function_path = get_file_path(&request.function_name);  
            run_wasm(&request.input, function_path, &request.function_name)  
            .unwrap()  
        }  
    };
```

```

    let body = serde_json::to_string(&Response { result }).unwrap();

    return (StatusCode::OK, body).into_response();
}

```

The function runners for Wasm modules and Docker containers are implemented in separate modules within the shared library. This separation of concerns makes it easier to maintain and possibly extend the codebase later, in case anyone wants to build upon this project in the future.

6.3.1 Executing Wasm Modules

The `run_wasm` function is responsible for executing a Wasm module based on the provided input and module path. We are using the `Wasmtime` crate, the official Rust library for executing Wasm modules, and the `wasmtime-wasi` crate for providing a WebAssembly System Interface (WASI) implementation, which lets us pass input as `stdin` to our functions.

```

use std::{path::PathBuf, time::Instant};

use wasi_common::pipe::{ReadPipe, WritePipe};

use wasmtime::*;

use wasmtime_wasi::sync::WasiCtxBuilder;

/* Load Wasm binary that matches module path, execute its main
 * function and return the result. Take timestamps for defining
 * startup and total runtime for each invocation. */
fn run_wasm(
    input: &str,
    wasm_module_path: PathBuf,
    func_name: &str,
) -> Result<FunctionResult, anyhow::Error> {
    let start_since_epoch = current_millis()?;
    let start = Instant::now();
    // 1.1.
    let engine = Engine::default();
    let mut linker = Linker::new(&engine);
    // 1.2.
    wasmtime_wasi::add_to_linker(&mut linker, |s| s)?;
    let stdin = ReadPipe::from(input);

```

```

let stdout = WritePipe::new_in_memory();
let wasi = WasiCtxBuilder::new();
let mut store = Store::new(&engine, wasi)
    .stdin(Box::new(stdin.clone()))
    .stdout(Box::new(stdout.clone()))
    .build();
// 1.3.
let module = load_module(&engine, wasi_module_path)?;
let startup_time = start.clone().elapsed().as_micros();
linker.module(&mut store, "", &module)?;
// 1.4.
linker
    .get_default(&mut store, "")?
    .typed::<(), ()>(&store)?
    .call(&mut store, ())?;
drop(store);
// 1.5.
let contents: Vec<u8> = stdout
    .try_into_inner()
    .map_err(|_err| anyhow::Error::msg("sole remaining reference"))?
    .into_inner();
let result = String::from_utf8(contents)? .trim() .to_string();
Ok(FunctionResult {
    result,
    metrics: Metrics {
        startup_time,
        start_since_epoch,
        total_runtime: start.elapsed().as_micros(),
        end_since_epoch: start_since_epoch + total_runtime,
    },
    func_type: ModuleType::Wasm,
    func_name: func_name.to_string(),
    input: input.to_string(),
})
}

```

The key steps involved in executing a Wasm module are:

- 1.1. Initialize the Wasmtime engine and linker.

- 1.2. Set up the WASI context, including stdin and stdout pipes for passing input and receiving output.
 - 1.3. Load the correct Wasm module from the provided path. This path is determined in the *call_function* and lets our runner know where our .wasm binary files are stored on our server. In our case it is in `/root/modules/wasm/``.
 - 1.4. Instantiate the module and execute the module's default function, which corresponds to each function's main function, which we will show later.
 - 1.5. Capture the output from stdout and return it as the function result.

Throughout the execution process, we collect startup and runtime metrics which are included in the returned `FunctionResult` struct.

6.3.2 Executing Docker Containers

The `run_docker` function is responsible for executing a function deployed as a Docker container. It spawns a new Docker container based on the provided image and function name, passing the input data as `stdin`.

```
use std::{

    io::{Error, Result, Write},
    process::{Command, Stdio},
    time::Instant,
};

/* Spin up Docker container matching image name for
function and take timestamps for startup which
is used to determine actual startup, and total runtime */

pub fn run_docker_image(
    image_name: &str,
    input: &str,
    func_name: String,
) -> Result<FunctionResult> {
    let start_since_epoch = current_millis()?;
    let start = Instant::now();
    // 2.1.

    let mut child = Command::new("docker")
        .args(["run", "--rm", "-i", image_name])
        .stdin(Stdio::piped())
        .stdout(Stdio::piped())
        .stderr(Stdio::piped());
    let mut output = String::new();
    let mut error = String::new();

    if let Err(error) = child.stdout.read_to_string(&mut output) {
        return Err(error);
    }

    if let Err(error) = child.stderr.read_to_string(&mut error) {
        return Err(error);
    }

    let end = Instant::now();
    let duration = end.duration_since(start).as_nanos() as f64 / 1_000_000.0;
    let millis = duration as f64 * 1000.0;

    Ok(FunctionResult { image_name, func_name, start, end, duration, millis, output, error })
}
```

```

.spawn()?;
// 2.2.

let cmd_start = current_micros()?;
// 2.3.

child.stdin.as_mut().unwrap().write_all(input.as_bytes())?;
// 2.4.

let output = child.wait_with_output()?;
let stdout = String::from_utf8_lossy(&output.stdout);
// 2.5.

let (result, actual_startup) = parse_output(&stdout, cmd_start)?;
Ok(FunctionResult {
    result,
    metrics: Metrics {
        startup_time: actual_startup,
        total_runtime: start.elapsed().as_micros(),
        start_since_epoch,
        end_since_epoch: start_since_epoch + total_runtime,
    },
    func_type: ModuleType::Docker,
    func_name,
    input: input.to_string(),
})
}

```

The key steps involved in executing a Docker container are:

- 2.1. Start up a Command process, from Rust's standard library, which spawns a 'docker run' command that we provide with image name and append stdin and stdout processes.
- 2.2. Time the moment the container spawns, this will be compared against the timestamp timed inside container, to get an accurate cold start time.
- 2.3. Write input to the stdin, which the function that is packaged inside our container is waiting on.
- 2.4. Wait for function inside the Docker container to execute and write out the result as stdout.
- 2.5. Parse output from the Docker container with a helper function, which comes in the form of a tuple of (result, micro_seconds_since_epoch). We will use

the second item in this tuple pair for determining the actual startup time of our function, which when the Docker container starts executing.

Similar to the Wasm module execution, we collect relevant metrics for startup and runtime of our functions invoked as Docker containers.

6.3.3 Challenges and Insights

During the implementation of Nebula's core functionality, several challenges were encountered that required creative problem-solving and valuable insights from experts from the Wasm community.

Initially, passing input data to Wasm modules proved challenging due to the need for shared memory and exported functions. This issue was resolved by leaning into the the POSIX-like behavior of WASI and provide input via stdin and output to be captured from stdout, like traditional command-line programs. As Malmgren (2022) suggested in their helpful article, using this feature of WASI provided a simpler approach for handling input and output in our Wasm module execution model. This approach was also implemented for passing input and output to and from our Docker containers, and while using stdin and stdout should incur some cost in terms of our metrics, this cost applies to both deployment methods, ensuring a fair assessment.

Another challenge arose when timing the startup of Docker containers. The first iteration involved defining a startup_time variable immediately after spawning the Command process. However, this naive implementation resulted in misleading “cold start” times of practically zero seconds, which would have invalidated our hypothesis. Fortunately, a former colleague, Roger Slaaen, suggested timing the startup at the moment the Docker images start running. This proved to be the key for solving this challenge as this accurately captures the cold start of our Docker-based functions when the container are loaded and start running.

Finally, the first implementation of the Wasm module execution had longer cold start times than expected, despite maintaining a fast runtime. While attending the WasmIO conference in 2024, I met Syrus Akbary, the founder of Wasmer, whom I showed my prototype to and mentioned my cold start issue.

Akbary, who works on the core of Wasmer runtime, shared key insights into how a runtime prepares a Wasm module for execution. The long startup latency turned out to be caused by a naive understanding of how the modules should be loaded, and with the first iteration, Wasmtime had to load “raw” Wasm binary files. To execute our modules on the runtime, we have to compile raw binary files into a

binary file that Wasmtime can understand. This step was crucial in unlocking the performance we will explore in Chapter 7. This method was backed by the employees at Fermyon, who were also attending the same conference, and shared that their framework Spin does the same thing while building Wasm modules for their platform. The code that implements this feature is as follows:

```
use std::fs::File, io::Read, io::Write, path::PathBuf;
use wasmtime::{Engine, Module};

// Serialize all modules into bytecode and save as files
fn serialize_wasm_modules(
    module_dir: PathBuf,
    serialized_dir: PathBuf,
) {
    // 3.1.
    let modules = list_files(module_dir.to_str().unwrap())?;
    // 3.2.
    let engine = Engine::default();
    // 3.3.
    for module in modules {
        let module_name = module.file_name().unwrap();
        // 3.4.
        let module = Module::from_file(&engine, &module)?;
        let module_bytes = module.serialize()?;

        let target_module = serialized_dir.join(module_name);

        // 3.5.
        let mut file = std::fs::File::create(target_module)?;
        file.write_all(&module_bytes)?;
    }
}
```

The steps involved in serializing our modules are:

- 3.1. Get a list of all the Wasm modules present in our deployment folder.
- 3.2. Instantiate the same kind of Wasmtime engine that we use in our function execution function.

- 3.3. Loop through all the deployed modules in our folder.
- 3.4. Load the raw Wasm binary and compile into a ‘Module’ which Wasmtime can run on its engine. Then serialize the resulting compiled binary represented by bytecode.
- 3.5. Write the resulting bytecode into a file that we store in a separate folder from our deployment storage.

This sequence precompiles our raw Wasm binaries into bytecode we can load and execute upon request. The code for loading the binaries can be seen here:

```
use std::fs::File, io::Read, path::PathBuf;
use wasmtime::{Engine, Module};

// Read modules from file and deserialize as bytes
fn load_module(
    engine: &Engine,
    wasi_module_path: PathBuf,
) -> anyhow::Result<Module, anyhow::Error> {
    let mut module_file = File::open(wasi_module_path)?;
    let mut module_bytes = Vec::new();
    module_file.read_to_end(&mut module_bytes)?;
    let module = unsafe { Module::deserialize(engine, &module_bytes)? };
    Ok(module)
}
```

Here we load the requested precompiled and serialized Wasm binary from our file system and deserialize it into a Wasm module we can run on Wasmtime. This function is called on each invocation of our Wasm modules. It is important to note the presence of an unsafe code block surrounding the `Module::deserialize` function. As Rust is a statically typed program, we must provide the types for our code ahead of time, before compilation, and even if we are “pretty sure” that we know what the contents of our `module_bytes` are, we can not be entirely sure, so we have to opt out of Rust’s safety in this case. This is not a crucial downside for our experiments, but hopefully something that can be changed in a later version of Wasmtime, in order to minimize the risk of this insecurity.

The resulting application built from this source is now ready to execute functions as either Wasm modules or Docker containers. However, to achieve this, we need to implement our environment for developing and deploying these functions.

6.4 Function Development and Deployment

Now that we have a web server that we can start on our deployment targets, we are going to need some functions that we can deploy and run on Nebula. For that we are going to develop the functions described in Section 4.1.4 in Rust and set up a way to package our functions in a way that we can spin them up and conduct the experiments we have planned.

6.4.1 Development Environment

For our development environment, we are going to set up some parts:

- A shared library: We are relying on stdin and stdout to pass input and output from our functions, and ideally we will not have to implement that individually for each of our functions. Therefore we are going to build a shared library that exposes functions for reading from stdin and printing to stdout.
- Timing docker cold start: In this library we will also implement a helper function for getting the actual time of cold start for our Docker containers.
- A Make file and Dockerfile: for streamlining the build and deployment of our functions to our target.

6.4.2 Shared Library

We have three goals with our shared library code. First we need a single way to run our functions that wraps our example functions with reading from stdin and writing to stdout. Secondly, we want a way to time the actual startup times of our docker containers. Finally we will write a Dockerfile that we will use to build your Rust functions into Docker images and a Makefile we can use to run to streamline the build and deploy of each function with ease.

The pipeline we want to support is:

```
cargo new example_function
# Implement example function in ./example_function/src/main.rs
make build_wasm build_docker
make deploy_wasm deploy_docker
```

To achieve this pipeline we will implement the following:

```
fn get_stdin<T: FromStr>() -> Result<T> {
    let mut input = String::new();
```

```

// Lock process and wait for stdin to be passed in
stdin()
    .lock()
    .read_line(&mut input)
    .context("Failed to read line")?;
// Parse input to a String and return the value
input.trim().parse::<T>().map_err(|_| {
    Error::msg(format!(
        "Failed to parse input as {}",
        type_name::<T>()
    ))
})
}

// If timestamp was recorded for function, print it to stdout
fn print_stdout<T: Display>(output: T, timestamp: Option<String>) {
    match timestamp {
        Some(timestamp) => println!("{}|{}", output, timestamp),
        None => println!("{}", output),
    }
}

pub fn run_function<T, F, R>(func: F, func_type: FunctionType)
where
    T: FromStr,
    F: Fn(T) -> R,
    R: std::fmt::Display,
{
    // If function is Docker, get timestamp, else if Wasm it is of type "None"
    let timestamp = match func_type {
        FunctionType::Docker => docker::get_epoch_timestamp(),
        FunctionType::Wasm => wasm::get_epoch_timestamp(),
    };
    // Read from input
    let input: T = get_stdin().expect("To parse correctly");
    // Call function with input and receive result
    let result = func(input);
    // Print input to stdout, with timestamp if func_type was Docker
    print_stdout(result, timestamp)
}

```

In this library we have exposed a `run_function` wrapper that we will import into the code for each of our benchmark functions. For each of our functions we want to start with timing the start of our function, if it is Docker. If it is not, the timestamp will be empty and not returned to Nebula. In hindsight, we could have measured the start time in the same way for our Wasm modules, but some manual testing revealed a negligible difference, which should not have any major impact on our results.

After timing the start of the function, we read from the `stdin`, run the function with the input value of type `T` (generic type, which can change depending on our functions). Then we print the result to `stdout`, either as a pair of `result | timestamp`, or just the result.

Then we had to specify a Dockerfile that our Docker build process will read from that describes how our functions are to be packaged into a Docker image. When we build with this Dockerfile, we also provide the target architecture we want our Docker container to run on. For our NREC deployment, we will target `linux/amd64` and for the Raspberry Pi, we will target `linux/arm64`.

```
FROM rust:1.70 as build
WORKDIR /usr/src/nebula
# Copy shared library into layer
COPY shared /usr/src/shared
# Expose function name (folder) as an argument for Docker
ARG PROGRAM_NAME
COPY $PROGRAM_NAME .
# Build release binary with timestamp
RUN cargo build --release --features docker
# Target smaller release docker image
FROM debian:bullseye-slim
# Install required libraries/dependencies and update
RUN apt-get update && \
    apt-get install -y libc6 && \
    rm -rf /var/lib/apt/lists/* ;\
# Copy over rust binary, and a custom run.sh script
COPY --from=build /usr/src/nebula/target/release/$PROGRAM_NAME /usr/local/bin/
COPY --from=build /usr/src/nebula/run.sh /usr/local/bin/
# When container is invoked, run this shell command
CMD ["run.sh"]
```

In this Dockerfile, we use the official Docker image for Rust as the build layer.

Then we copy over the required images into this layer, build our functions with the release flag (ensuring best performance), and flag that we want to print the timestamp for these Docker functions through a feature flag.

Then we use a smaller debian:bullseye-slim image as a target for our final Docker image as we do not need the entire Rust toolchain installed in our final image to run our functions as Rust binaries.

Finally, we add a simple run script that the image invokes when we run our Docker containers as described in Section 6.3.2, invoking the main function of our benchmark functions.

While explaining the code for this file is not crucial for our research, Figure 6.1 illustrates the flow it follows.

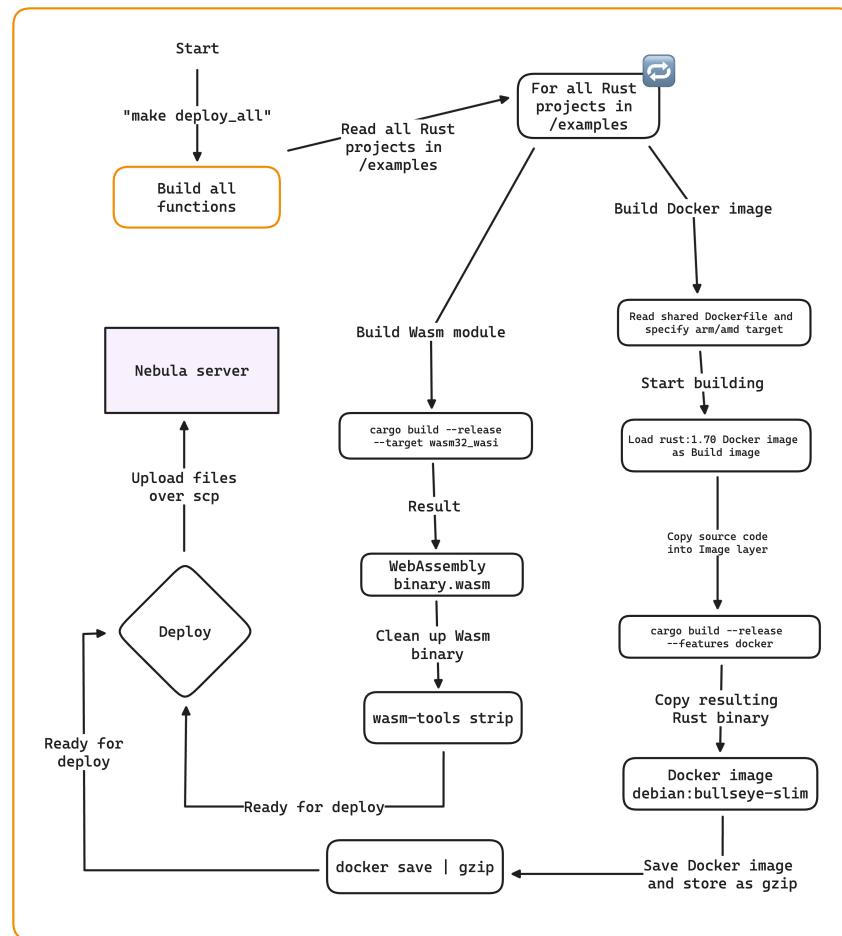


Figure 6.1: Streamlined flow of building and deploying functions.

One thing to note in this flow, is the step *wasm-tools strip*. *wasm-tools* is a CLI tool maintained by the Bytecode Alliance that includes a series of tools we can use to manipulate our Wasm modules. Compiling a Rust program to Wasm includes some unnecessary sections that we can strip away. The Wasm binaries we built for

our experiments resulted in ~2Mb in size. Using this tool lets us cut this down to ~200Kb instead, while still maintaining its function. More on that in Section 8.3.

6.4.3 Implementing a Benchmark Function

Now that we have set up an environment for building and deploying functions to Nebula, we can look at how we developed the benchmark functions themselves. We will implement the recursive Fibonacci function in this section.

```
use shared::{run_function, FunctionType};  
/* Reads std in as input, retrieves the fibonacci sequence  
* and returns the last number of the fibonacci sequence of  
* the provided size */  
fn main() {  
    // Determine if function type is docker or wasm  
    let func_type = if cfg!(feature = "docker") {  
        FunctionType::Docker  
    } else {  
        FunctionType::Wasm  
    };  
    run_function(fib, func_type)  
}  
// Recursive function that calculates the nth fibonacci number  
fn fib(size: i64) -> i64 {  
    match size {  
        0 => 0,  
        1 => 1,  
        n => fib(n - 1) + fib(n - 2),  
    }  
}
```

Our shared library handles the core functionality. As described in Section 6.3.3, we treat our functions as normal programs, which we do by running them by invoking their main function. In this main function we determine if the function will be compiled to a Wasm module or a Docker container based on the --features docker flag we pass in our build command. (See Figure 6.1)

Then we implement the `fib` function which takes in what number in the Fibonacci sequence we are interested in, which we have implemented recursively, and pass this into our `run_function` wrapper from our shared library.

With that in place, we just need to set up our deployment targets - the Raspberry Pi and the Debian-based virtual machine on NREC.

6.5 Deployment Setup

For our research we have decided to set up Nebula both on a Raspberry Pi running Debian on an ARM-based CPU, and on a virtual machine running on NREC, which gives us an AMD-based virtual CPU.

6.5.1 Raspberry Pi Setup

The Raspberry Pi, a low-cost and energy-efficient SBC, was chosen as one of our deployment targets to measure power consumption of our function executions. The model we chose was the Raspberry Pi 4 Model B, with 4GB of RAM, running the lite version of the Raspberry Pi OS, which is built on top of Debian.

6.5.1.1 Hardware Setup

1. Use the Raspberry Pi image tool for installing Raspberry Pi OS (lite, without desktop). Configure the OS with:
2. Wifi credentials, if not using ethernet.
3. SSH key, created with ssh-keygen, used to authenticate when connecting to device through SSH. This is also useful for uploading our files through scp.
4. Insert the microSD into the device and connect it to power, which boots it up.
5. If device connects to internet, it will be discoverable as `raspberrypi.local`, which can be used instead of its IP-address.

6.5.2 Virtual Machine Setup (NREC)

For our experiments, we chose NREC as our IaaS platform for hosting our virtual machine and benchmark a more typical deployment environment. To use this service, we had to contact the administration and verify that we were a student at the University of Oslo.

Once inside the NREC platform, we can launch an instance of our virtual machine, following their guide for doing so. For our experiments we chose the pre-configured flavor of `m1.medium`, which gives us 1 virtual CPU, 4GB of RAM and 20GB storage space. The virtual machine was built with the base image of Debian 12. After this is set up, we can ssh into the server with a ssh-key we provided during setup.

Instance Name	Image Name	IP Address	Flavor
nebula	Debian 12	158.39.200.15	m1.medium

Flavor Details: m1.medium

ID	c76cbcb9-df2d-4b8c-9587-b9b9bc232685
VCPUs	1
RAM	4GB
Size	20GB

Figure 6.2: Our virtual machine configured on NREC.

6.5.3 Software Setup

As both our setups are running a flavor of Debian, setting up the software required to run our experiments are identical. First we must update the packages on our Debian server, and then install Docker Engine, following their official guide in their manual⁴.

At this point, we are ready to deploy Nebula to each server and deploy our functions.

6.5.4 Tying It All Together

Once we have install the required software on our servers, we can start to deploy our Nebula web server and upload our Wasm modules and Docker images.

In our main project folder, we build Nebula using the cargo build command. We add the flag `--release`, which optimizes the binary for production, and the flag `--target <target-arch>`, which cross-compiles Nebula to the target architecture. Our prototype was developed on a M2-based Macbook pro, so the default build target is not supported on our target platforms. Cross-compiling works fine, but the easiest and fastest way to build Nebula is to do it on the target architecture itself.

For the Raspberry Pi, this was done by cloning the Github repository with the source code onto the device, and building it from there. For the NREC implementation, a Github action was implemented for building the application on a Linux VM as part of their CI/CD service. This is not a crucial detail for our research, but valuable for anyone who would want to recreate the experiment and expand upon it.

Finally, we build and deploy our functions as described in Section 6.4.2, and upload our Wasm binaries and our gzipped Docker images to each server. Part of

⁴<https://docs.docker.com/engine/install/debian/>

deploying the Docker images is loading them into the local Docker image list, done with the command `docker load -i /path/to/image.tar.gz`. When we start our server we call the `serialize_wasm()` function and serialize our Wasm modules into `byte_code` that we can load and run on Wasmtime. With all this in place, we can now perform our experiments.

6.6 Benchmarking

Now that we have our Nebula FaaS prototype up and running on both our Raspberry Pi and virtual machine on NREC, we can start implementing the benchmarking utility that will be used to conduct the experiments and collect our data. While there exist tools out there for performing experiments like this, the setup we have implemented would be difficult to cover with an already existing tool. For our experiments we will need to hammer Nebula with a large amount of requests, which we in theory could do with something like Postman. But, we also need to measure power data while hammering our Raspberry Pi over Modbus TCP and then bind these power readings to our function invocations. Because of this, we will develop a custom benchmark utility in Rust, which will let us collect the data we are interested in.

For our experiments we will need to implement two core functionalities. First, we need to implement a function that loops through a pre-defined set of input values for each function and perform POST requests to Nebula, storing the function results, which includes our efficiency metrics, on our computer. We also need a function for measuring power measurements while performing our POST requests to Nebula on our Raspberry Pi.

6.6.1 Stress-Testing Nebula

To benchmark Nebula's cold starts and runtime efficiencies, we will need to perform a range of POST requests with the intent of gathering a wide range of data that we can compare later. For this we rely on the widely used Rust crate, `reqwest`, a library for performing http requests in Rust. The code for performing the request is omitted from this example, but the suite of functions and their corresponding inputs we are testing with can be seen here:

```
async fn stress_nebula(
    client: Client,
    url: &str,
) -> anyhow::Result<Vec<FunctionResult>> {
```

```

// 4.1. function name, how many requests, steps
let functions: Vec<(&str, u32, u32)> = vec![
    ("exponential", 353, 2),
    ("factorial", 130, 1),
    ("fibonacci-recursive", 40, 1),
    ("prime-number", 600, 10),
];
let mut function_results: Vec<FunctionResult> = Vec::new();
// 4.2.
for function in functions {
    // 4.3.
    for input_value in 0..=function.1 {
        // 4.4.
        let input_value = input_value * function.2;
        // 4.5.
        for _ in 0..6 {
            // 4.6.
            let wasm_results = make_request(
                &client, url, function.0,
                "Wasm", &input_value.to_string()
            ).await?;
            function_results.extend(wasm_results);
            let docker_results = make_request(
                &client, url, function.0,
                "Docker", &input_value.to_string(),
            ).await?;
            function_results.extend(docker_results);
        }
    }
}
// 4.7.
Ok(function_results)
}

```

Through the development and manual testing of Nebula, we identified a range of benchmark values that provide a clear picture of our prototype's performance when running functions as WebAssembly (Wasm) modules or Docker containers.

The steps we take in our stress-test function are:

4.1. Define a list of modules that we are going to benchmark. For each function we have a tuple that consist of three values.

- The name of the function.
- The amount of input values we want to test for each function.
- The factor we want to multiply our input values with, to cover a wider range of input values for our experiments.

4.2. Loop through our list of functions.

4.3. For each function, loop through a range of <0, steps>.

4.4. Determine input value by multiplying current step with factor.

4.5. Perform each set of function and input 6 times. This is lets us gather a mean of each metric for our results later, to account for variance in our results.

4.6. Perform a function invocation for each function as both Wasm and Docker sequentially, then append result to the final result.

4.7. Return the list of results.

The result of this, is a long list of values that we store as JSON, for further analysis.

An example of a function result from benchmarking on our Raspberry Pi looks like this:

```
{  
  "func_name": "fibonacci-recursive",  
  "func_type": "Wasm",  
  "input": "2",  
  "result": "1",  
  "metrics": {  
    "startup_time": 3778,  
    "total_runtime": 8197,  
    "start_since_epoch": 1714755235731991,  
    "end_since_epoch": 1714755235740188  
  }  
}
```

Part of our metrics is the timings for microseconds since epoch for the start and end of each function invocation, this is important for our power measurements.

6.6.2 Measuring Power

For communicating with our Gude expert control over Modbus TCP, we are going to use the `tokio-modbus` crate. This library will let us write an interface against the device over our network.

As was presented in Section 3.7.4, we saw how a client and server communicates through Modbus TCP packets. In figure Figure 6.3 below, we zoom into this packet, as they are important for our implementation for communicating with our Gude power supply.

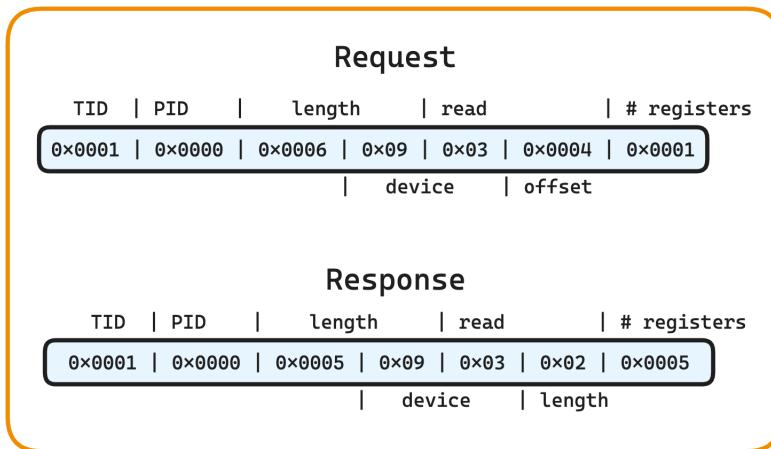


Figure 6.3: A Modbus TCP packet.

To read from the packets we receive from our Gude power supply, we need to know which register each measurement it can report to are registered. From the manual of our Gude Expert Power Control 1105, we find this table of sensor that it can report on. In Table 4.2, we extracted the most relevant sensors for our measurements, its register address, and how accurate its reported units are.

To re-iterate, the formula for calculating power is $P = U \times I$, where U is voltage and I is current. For example, with a constant voltage of 240V and a current of 24mA, we will get $240V \times 24mA = 5,76mW$. Another important sensor to read from is the Power factor, which is a factor that indicates how much power the connected device is actually consuming, compared to the total draw of current through the power cable. For our system it varied between 0.50 and 0.55, so with an example of $PF = 0.55$ for our example above, we get $240V \times 24mA \times 0.55 = 3,168mW$. This product is what we are going to collect data on.

In Section 6.6.5 we will discuss some challenges regarding this setup, and most importantly that we are going to assume a constant voltage of 240V for calculating our power measurements. This is because for every sensor we want to read from

our Gude device at a given measurement adds up to the total amount of time it takes to read this measure. When we limit our readings to focus on current and power factor, our energy measurements come in fast enough to accurately time them against our function executions.

6.6.3 Reading from Modbus TCP

For our sensor data, we are going to define a set of types and configure a mapping of our sensor to register address data. For our measurements, we will measure the current, power factor and time the start and end of our readings. These timings will help us map power readings to the function execution.

```
pub struct SensorData {  
    pub power: f32,  
    pub start_read: u128,  
    pub end_read: u128,  
}  
  
pub enum LineInEnergySensor {  
    Current = 0x406,  
    PowerFactor = 0x40a,  
}
```

With this in place, and we can define our function for reading from modbus. The goal of this function is to run in a parallel thread while we conduct our stress-testing of Nebula on our Raspberry Pi. This function will be run continually until the parent thread breaks, so we will read measurements from Gude, described in code:

```
use serde_derive::{Deserialize, Serialize};  
use tokio_modbus::prelude::*;

// Start tcp client that connects to Gude controller and  

// reads sensor values over register addresses in tcp packets
pub async fn read_modbus_data(  
    addr: &str,  
) -> Result<SensorData, anyhow::Error> {  
    // 5.1.  
    let mut client = tcp::connect("192.168.68.70:502".parse()?).await?;  
    let start_read = current_millis();
```

```

// 5.2.

let current_register = client
    .read_input_registers(LineInEnergySensor::Current as u16, 2)
    .await?;

let raw_current =
    (i32::from(current_register[0]) << 16) | i32::from(current_register[1]);
let current = (raw_current as f32) / 1000.0;
// 5.3.

let power_factor_register = client
    .read_input_registers(LineInEnergySensor::PowerFactor as u16, 2)
    .await?;

let raw_power_factor =
    (i32::from(power_factor_register[0]) << 16) |
    i32::from(power_factor_register[1]);
let power_factor = (raw_power_factor as f32) / 1000.0;
// 5.4.

let power = 240.0 * current * power_factor;
let end_read = current_micros()?;
Ok(SensorData { power, start_read, end_read })
}

```

The steps we go through here are:

- 5.1. Define our tcp client and connect it to our Gude device on the modbus tcp port (e.g., 192.168.68.71:502), and create an instance of our SensorData struct, where we set values to 0, and time the start of our measurement in microseconds since epoch.
- 5.2. Read current from its register, cast it as a floating point number and divide by 1000, so we get the value in *A*, instead of *mA*.
- 5.3. Read power factor from the register, which is multiplied by 1000 in the measurements, so we cast it to f32 and divide it by 1000 to get the power factor.
- 5.4. Calculate power at the time of measurement to $240V \times I \times PF$ and return the power product and start and end times for the measurement.

6.6.4 Attaching Power Measurement to Function Invocations

When we benchmark Nebula on our Raspberry Pi, we spawn two threads that run in parallel, using the Tokio async runtime. One of these threads performs out stress-

test function from Section 6.6.1, while the other reads from our `read_modbus_data` function from Section 6.6.3. We continue to read sensor data until the stress-test function is completed, and save both lists as a variable. Then we pass both the resulting function invocation results and the power readings to a function that finds power measurements that coincide with a function invocation.

Figure 6.4 below illustrates how power measurements are paired with coinciding function invocations to estimate the power required to run our functions.

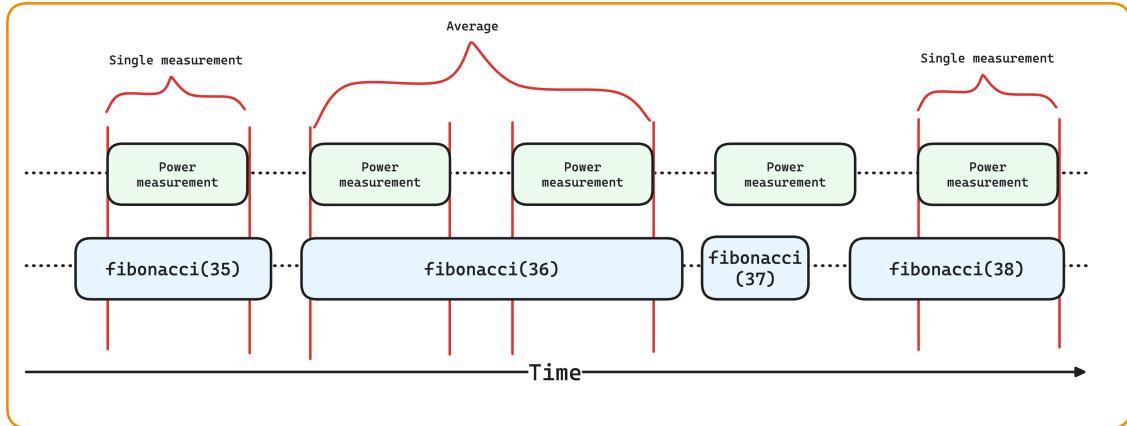


Figure 6.4: Power measurements that overlap with function invocations.

In some of our benchmarks, there were instances of function invocations being too fast to be able to pair with a power measurement. In these cases, we ran the functions with its input multiple times, to ensure that every input value for each function has the same amount of invocations to improve the validity of our data. In other cases, we saw functions that were so slow that multiple power measurements were performed over the functions lifetime, and in this case we sum all the corresponding measurements and get the average power consumption. For example, a slow recursive fibonacci function might have 30 power readings associated with it.

The code below shows how this is done. It loops through each function result, finds the energy measurements that begin and end within the start and end of each function invocation. Then it updates each function result with the average power that was consumed by the system during the functions execution. We also calculate the estimated μWh in energy consumption, with t in microseconds, with the formula:

$$P \times \frac{t}{3600}$$

```
// Loop through all energy measurements and mutate
// function results with energy measurements that
```

```

// were measured during their lifetime
pub fn associate_power_measurements(
    mut function_results: Vec<FunctionResult>,
    sensor_data: &Vec<SensorData>,
) -> Vec<FunctionResult> {
    let mut processed_results = Vec::new();
    let microseconds_to_hour = 1000.0 * 1000.0 * 60.0 * 60.0;

    for mut function_result in function_results {
        let function_start_time =
            function_result.metrics.start_since_epoch;
        let function_end_time = function_result.metrics.end_since_epoch;
        let mut total_power = 0.0;
        let mut num_readings = 0;
        for data in sensor_data {
            if data.start_read >= function_start_time
                && data.end_read < function_end_time
            {
                total_power += data.power;
                num_readings += 1;
            }
        }
        if num_readings > 0 {
            let avg_power = total_power / num_readings as f32;
            let duration = function_end_time - function_start_time;
            let energy_consumption_wh =
                (avg_power * duration as f32) / microseconds_to_hour;

            function_result.metrics.average_power = avg_power;
            function_result.metrics.energy_consumption_wh =
                energy_consumption_wh;
            processed_results.push(function_result);
        }
    }

    processed_results
}

```

With this in place, we can now perform our experiments on both our virtual ma-

chine and on the Raspberry Pi. Figure 6.5 below shows the physical setup for the experiments on Nebula on our Raspberry Pi.

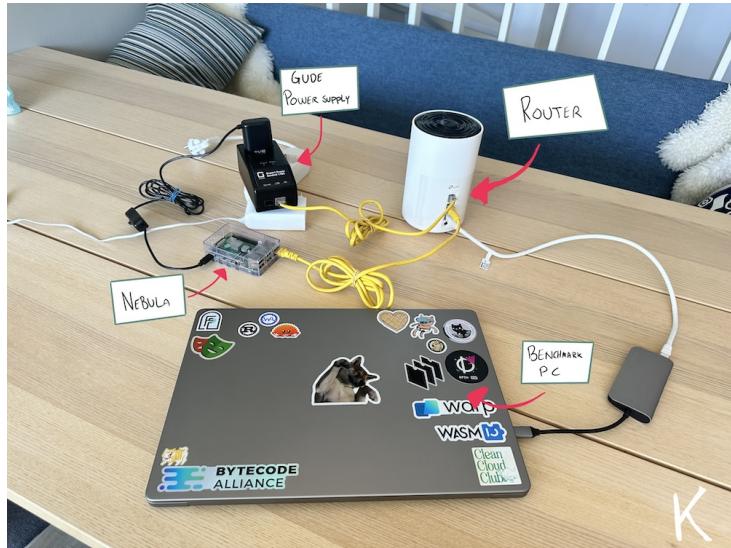


Figure 6.5: The physical setup, including the Raspberry Pi connected to power through our Gude power supply. The Raspberry Pi and Gude controller is connected to the benchmark computer on a local network through ethernet.

6.6.5 Energy Measurement Challenges

In Section 3.7, we saw several protocols and measurement devices, and how they work, described. The first iteration of measuring power was inspired by Qian (2022), and consisted of an Aotec Smart Switch 7, Aotec Z-stick 7 and relied on the zwave-js-ui application to act as a publisher for reporting energy measurements over MQTT. This setup had one major flaw; energy reporting through MQTT could only be reported once every 30 seconds at a minimum. This turned out to be too slow for attempting to measure accurate readings for our experiments.

This limitation was unexpected, but Qian had used an Aotec Smart Switch 6, the previous generation, which turned out to support reporting once every second. This device was luckily found on Finn.no, and allowed us to upgrade the accuracy down to measuring every second instead.

However, as we will present in Chapter 7, our function executions on our Raspberry Pi can come down to single-digit milliseconds, meaning that once every second is still not accurate enough to pinpoint the current draw on our system for each function invocation.

The Gude Expert Power Control 1105 turned out to be the most suited device for our experiments. This device supports many protocols, one of which is the Modbus TCP protocol. By switching to this power supply, connecting our Raspberry Pi

and benchmarking computer to the same network over ethernet, and changing the communication method from MQTT to Modbus TCP, we were able to perform energy readings in microseconds.

6.7 Testing - Validity and Rigidity

While formal unit testing was not exercised during the development of Nebula, the choice of the Rust programming language provided inherent benefits in terms of code reliability and correctness, ensuring the validity and rigidity of the prototype.

6.7.1 Rust's Compiler Guarantees

As outlined in Section 6.1, choosing Rust as the programming language was driven by its strong relationship with Wasm and its suitability for building efficient and scalable applications. Significant factors to this are Rust's strict compiler checks and emphasis on memory safety, concurrency, and parallelism.

The Rust compiler enforces a set of rules that prevent common programming errors, such as null pointer dereferences, data races, and memory leaks. By catching these issues at compile-time, Rust eliminates entire classes of runtime errors, increasing the overall stability and predictability of the codebase.

Furthermore, Rust's ownership model and borrowing rules ensure safe and efficient memory management, without the need for manual intervention or a garbage collector. This approach reduces the risk of memory-related bugs, which are very difficult to detect and resolve. We can not guarantee that there are not any present in Nebula, as testing will never be able to guarantee the absence of bugs. However, we saw that during extensive stress-testing, our prototype consistently ran for several hours straight during our benchmarks.

One exception to these guarantees is the deserialization of our Wasm modules, as discussed in Section 6.3.3. In this case, we opt out of Rust's compiler guarantees by telling the compiler that we have a better understanding of the situation and can safely read our Wasm binary files as bytecode. While this represents a potential risk if malicious binaries were to be introduced to the web server, it did not pose an issue for our research.

6.7.2 Rigorous Testing and Benchmarking

Although formal unit testing was not a primary focus, the code underwent rigorous manual testing and code review processes. As detailed in Section 6.6, the compo-

nents of Nebula were thoroughly tested with various input scenarios, edge cases, and stress tests to ensure correctness and robustness.

Moreover, the benchmarking suite itself, described in Section 6.6, served as a form of integration testing, exercising the entire system end-to-end and validating the correctness of the results.

The combination of Rust’s language features, idiomatic coding practices, and extensive manual testing through the benchmarking suite provided a strong foundation for the validity and rigidity of the Nebula prototype, ensuring reliable and robust performance during the experiments and evaluation.

Part III

Analysis

Chapter 7

Results

One accurate measurement is worth a thousand expert opinions.

—Grace Hopper

With the first goal of developing a prototype cloud computing platform for the FaaS paradigm accomplished, we can turn our attention to the remaining objectives of this thesis. Leveraging the prototype implemented in Chapter 6 and the supporting benchmark utility, experiments were conducted on Nebula to evaluate its performance. This chapter presents the results, while we will discuss these results further in Chapter 8.

7.1 Methodology Recap

The experiments involved running a series of benchmark functions on both a Raspberry Pi model 4B and a virtual machine on NREC, representing different operational environments. Each function was tested as a Wasm module and a Docker container to capture a comprehensive dataset under controlled conditions. The key metrics measured were:

- Startup latency (Cold start): The time elapsed between invoking the function and its execution start.
- Total runtime: The complete lifecycle duration from start to result return.
- Power consumption: The power draw (in watts) on the Raspberry Pi during function execution loads, calculated as the product of measured current, power factor, and an assumed 240V voltage.

Additionally, by combining total runtime and power consumption data, an estimation of the energy consumed by function executions throughout their lifetime could be derived. Due to the fast execution times of our functions, energy consumption (E) is expressed in microwatt-hours (μWh), calculated as $E = P \times t$, where P is the power in watts and t is the time in microseconds.

To ensure accuracy and repeatability, data collection was automated using the benchmarking utility developed in Section 6.6, proving invaluable for this report.

Of particular note is the extreme performance disparity between Wasm modules and Docker containers, making it difficult to visualize their data on the same graphs effectively. Consequently, where applicable, the data is presented separately to allow for closer examination of each execution mode.

7.2 Startup and Runtimes

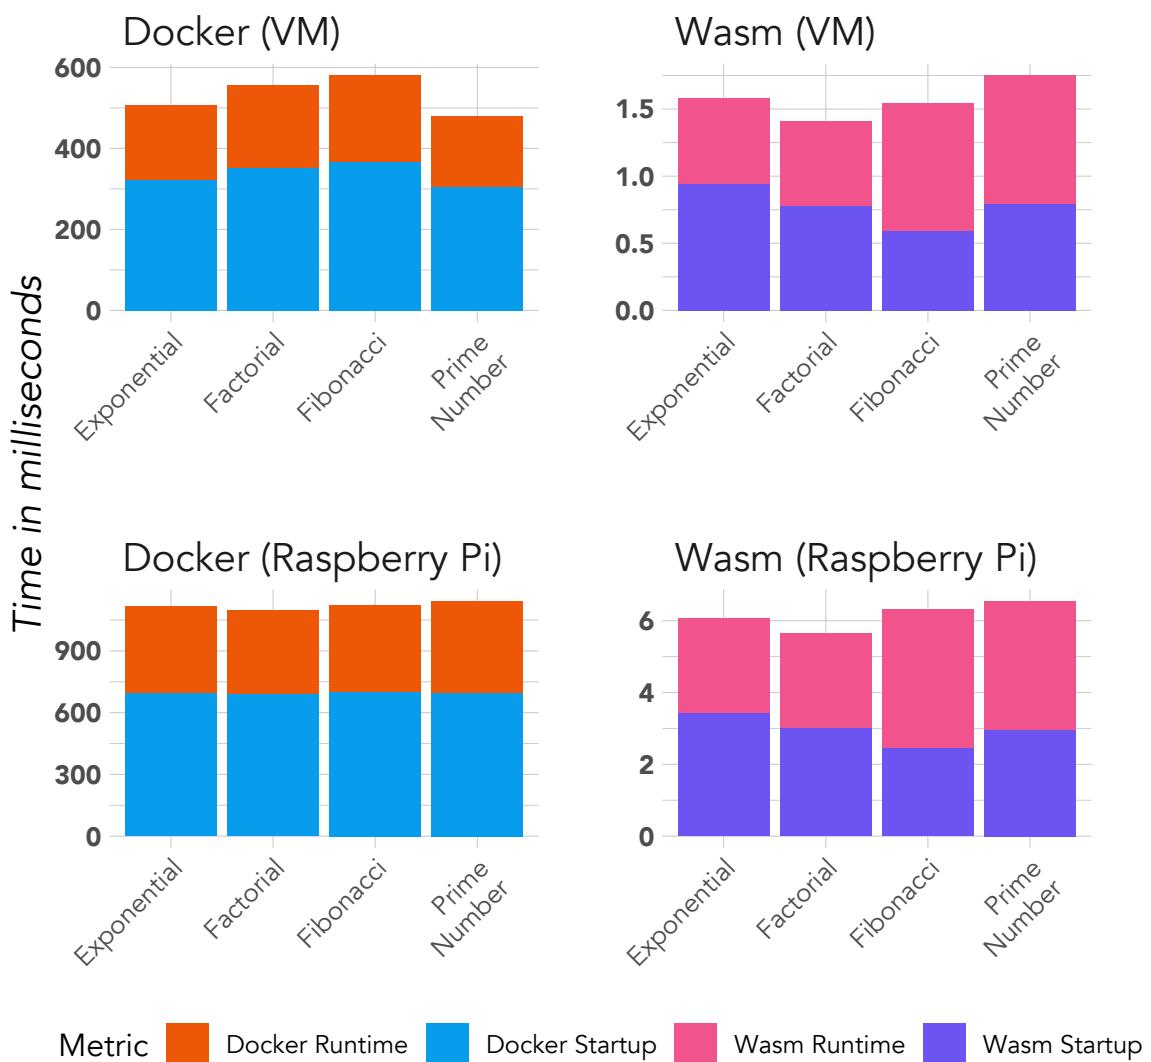


Figure 7.1: Mean startup and runtimes on both our setups.

Figure 7.1 provides an overview of the average startup and runtimes across our benchmarks on both our virtual machine and Raspberry Pi setups. While this offers a general comparison between Wasm and Docker invocations, Figures 7.2 and 7.3 present more detailed scatter plots with an average line drawn across showcasing the performance variance for each input across multiple runs.

At first glance, the startup and runtimes of our Wasm modules and Docker containers may appear similar in Figure 7.1, but a keen eye will notice the differing y-axis scales. While Docker containers startup and return results in the range of hundreds of milliseconds, the Wasm modules startup and shut down in just single-digit milliseconds.

It is evident from these figures that the data follows a similar curve on the Raspberry Pi, albeit with longer runtimes compared to the virtual machine, as expected due to the resource constraints of the former. To account for variance, every permutation of function type, name and input was executed at least five times during the benchmarking process.

For the sake of clarity, measurements exceeding specific input thresholds for the fibonacci (above 30) and prime number (above 1500) functions have been excluded from the graphs, as their steep runtime increases would impede effective visualization.

Based on this, our findings are summarized in Tables 7.1 and 7.2. These tables show the direct comparison of each metric we set out to measure between Wasm modules and Docker containers on our setups.

Table 7.1: Comparison of Wasm vs. Docker cold starts on VM.

	Docker	Wasm	vs Docker
Cold start	324.66 ms	0.89 ms	365×
Runtime	187.03 ms	0.69 ms	271×
Total runtime	511.69 ms	1.58 ms	324×

Table 7.2: Comparison of Wasm vs. Docker cold starts on Raspberry Pi.

	Docker	Wasm	vs Docker
Cold start	692.36 ms	3.19 ms	217×
Runtime	426.82 ms	2.93 ms	146×
Total runtime	1119.18 ms	6.12 ms	183×
Energy consumption	1176.15 µWh	5.95 µWh	198×

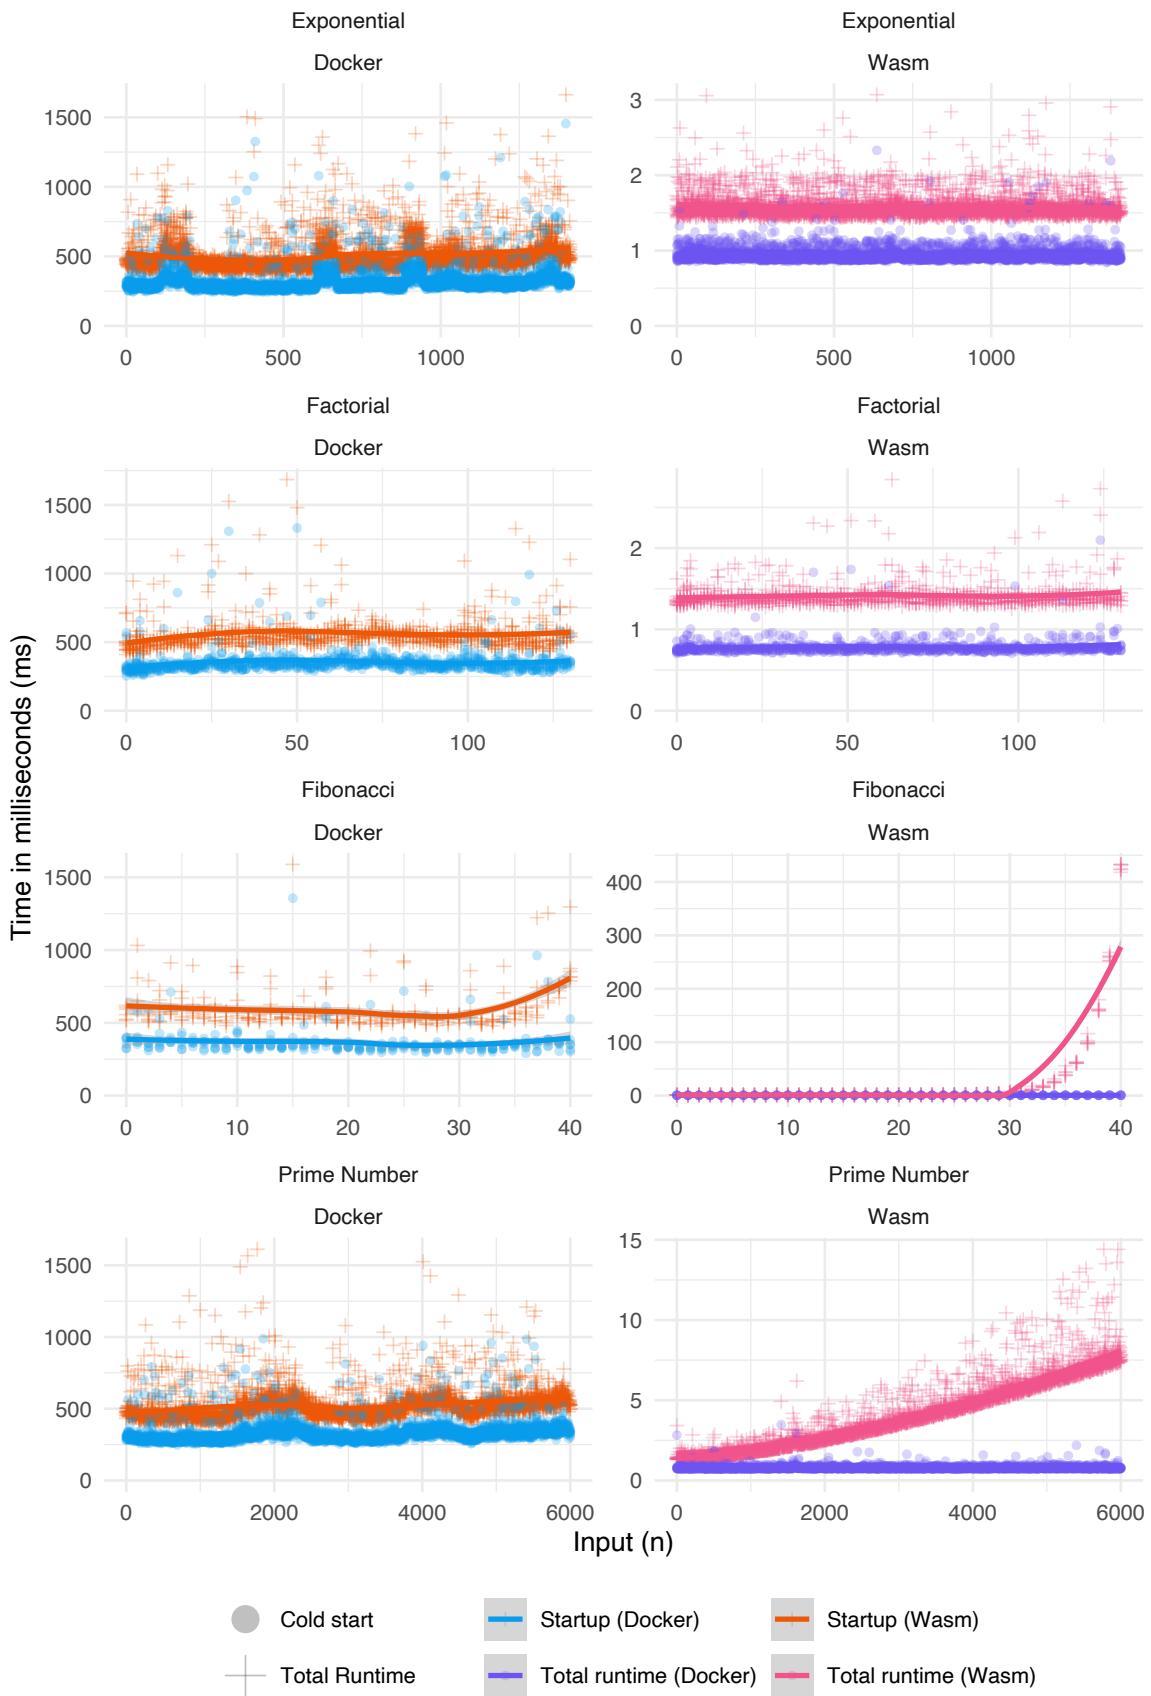


Figure 7.2: Wasm runtime measurements on VM during benchmarks.

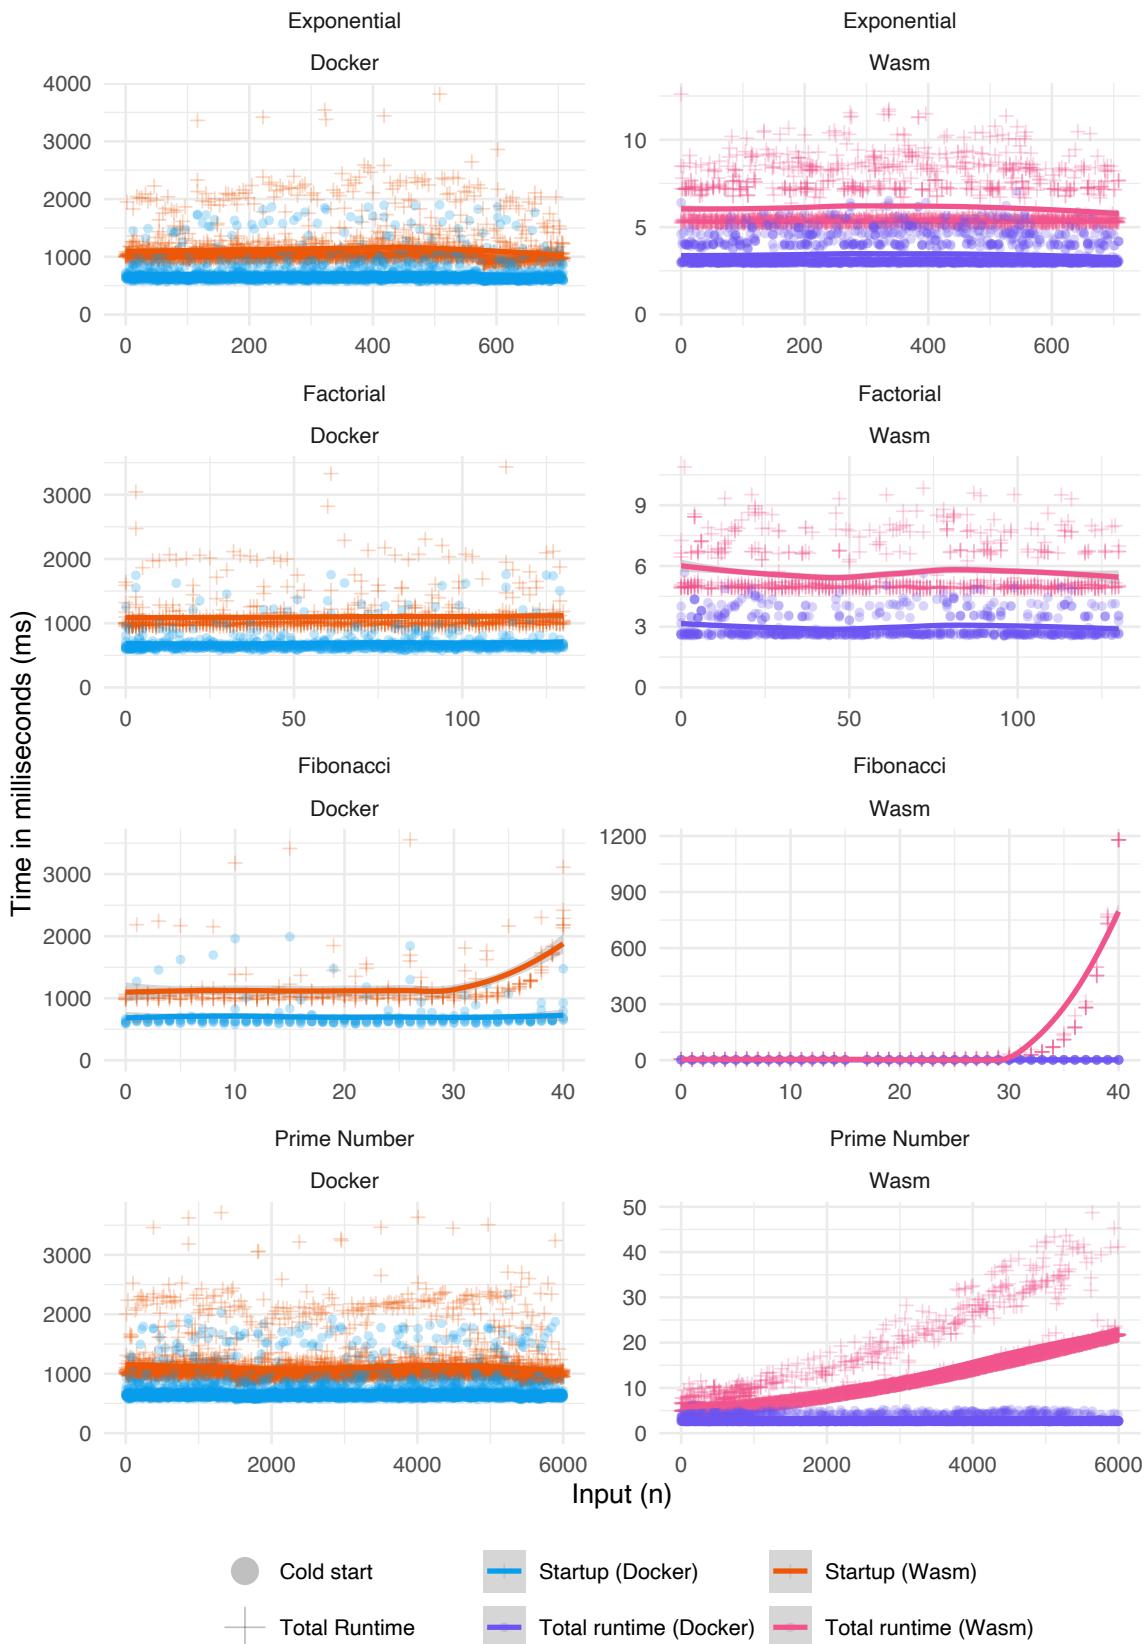


Figure 7.3: Startup and Runtime measurements on Raspberry Pi 4B during benchmarks.

7.3 Power Load and Energy Consumption Results

This section presents the data for power load and energy consumption measured on the Raspberry Pi during our benchmarking experiments. The results for each function are presented in pairs, with the first figure of each pair showing the average power load in watts that the Raspberry Pi drew during function invocations for each input value (n). The second figure displays the estimated energy consumed, in microwatt-hours (μWh), calculated by multiplying the average power load by the total execution time in microseconds for each function. In Figure 7.4, we can see the average energy consumption that our Rasbperry Pi consumes during their execution.

An interesting observation from the fibonacci and prime number functions is that while the average power load converges to a constant value beyond certain input for both Wasm and Docker executions. Despite the consistent power draw during prolonged execution times, the energy consumption curves for each function type maintain their trajectories.

Despite the consistent power draw during prolonged execution times, the energy consumed continues to increase along the same curve of each execution mode. The divergence in energy consumption, even with converging power loads, highlights the importance of considering both power and time factors when optimizing energy efficiency in serverless computing environments.

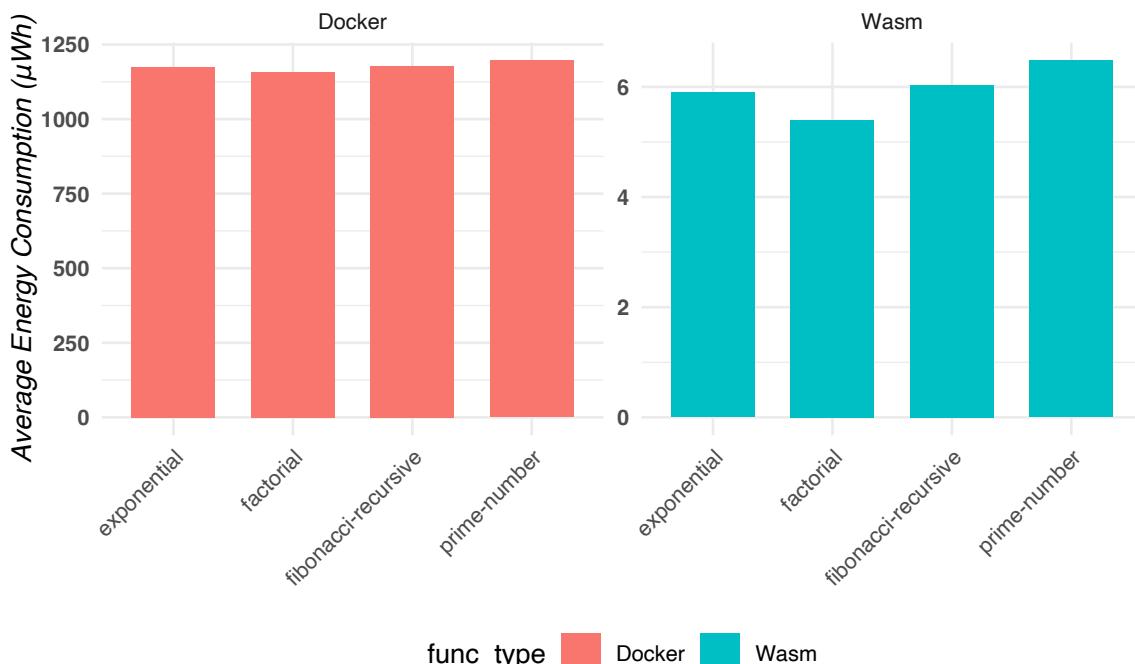


Figure 7.4: Average energy consumption for our functions.

Fibonacci

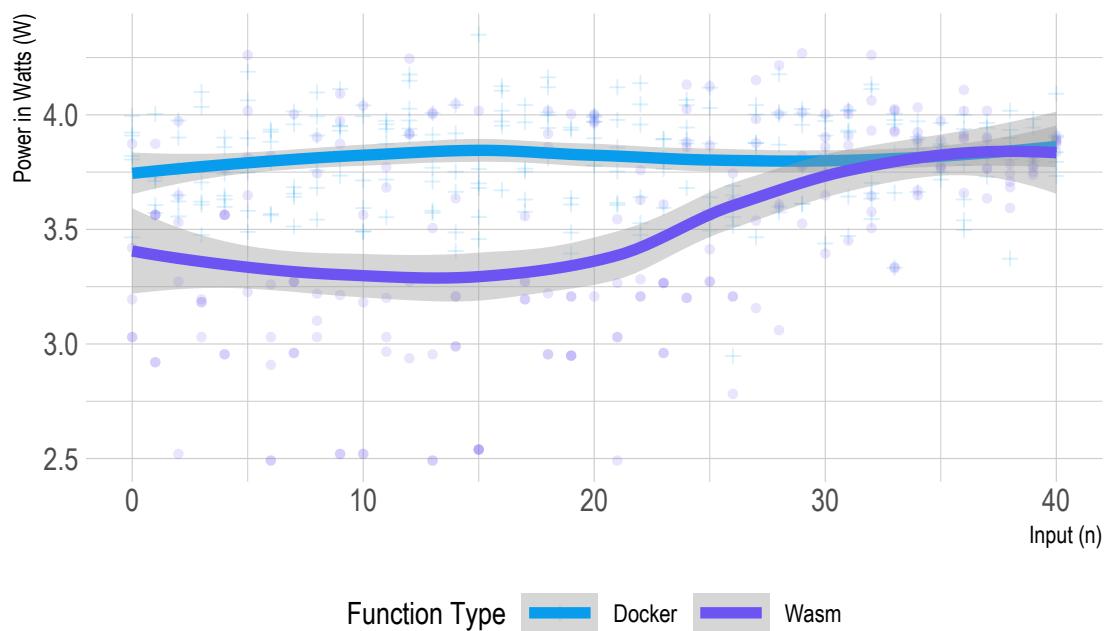


Figure 7.5: Power load on Raspberry Pi during Fibonacci benchmark.

Fibonacci

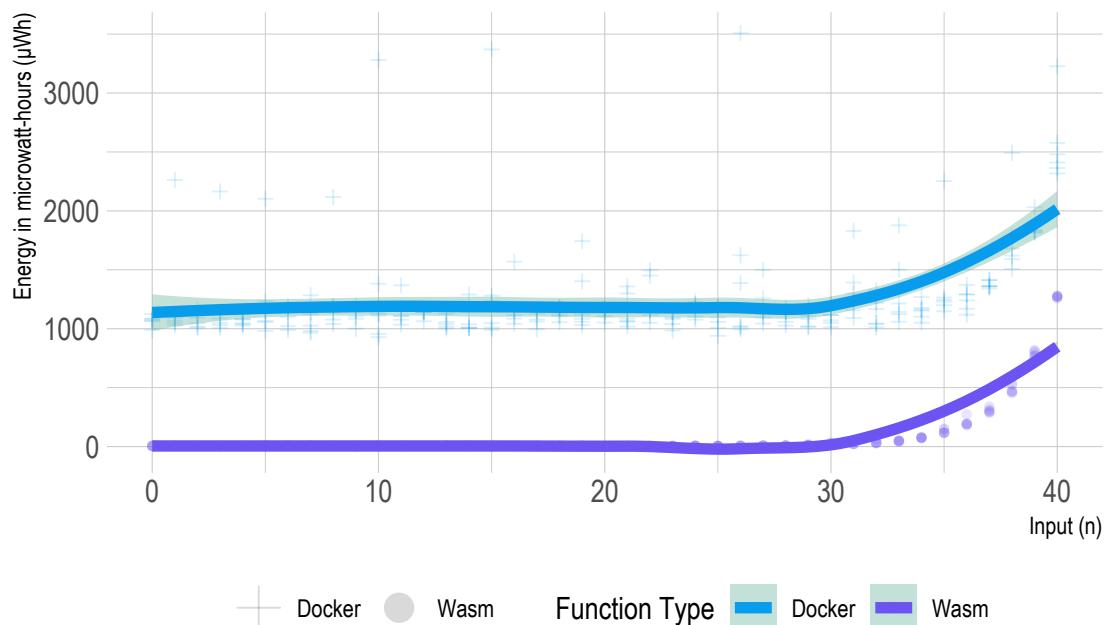


Figure 7.6: Energy consumption on Raspberry Pi during Fibonacci benchmark.

Factorial

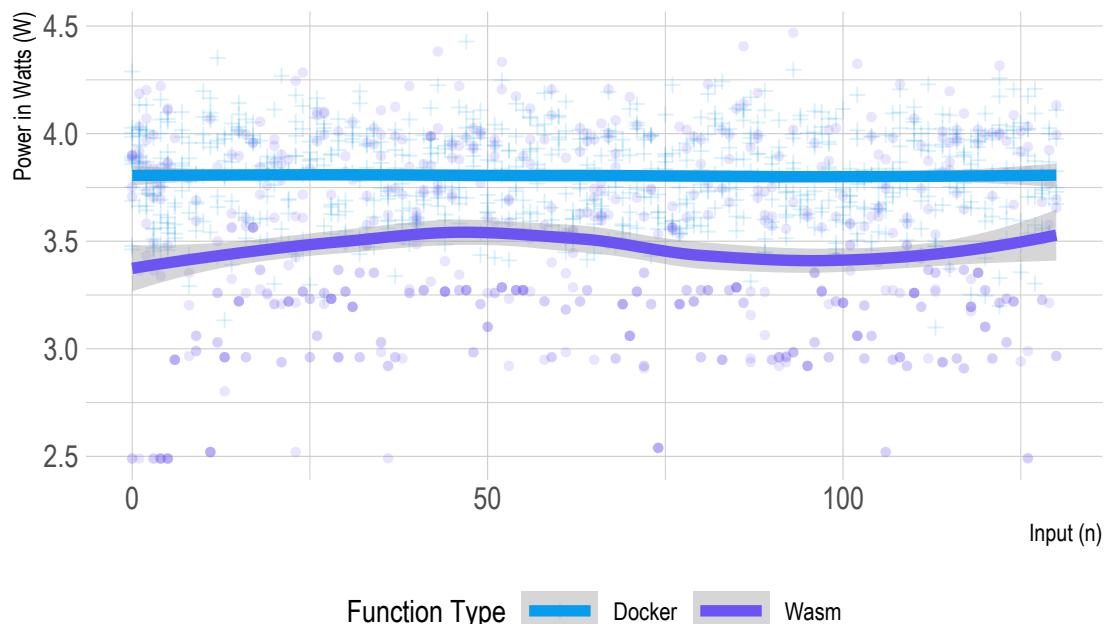


Figure 7.7: Power load on Raspberry Pi during Factorial benchmark.

Factorial

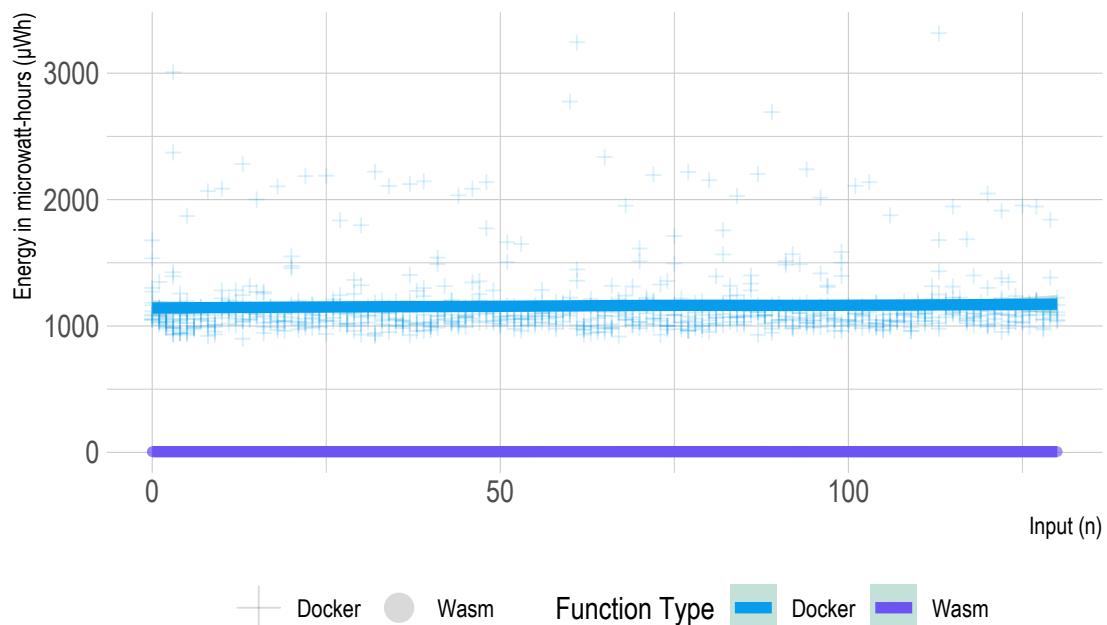


Figure 7.8: Energy consumption on Raspberry Pi during Factorial benchmark.

Exponential

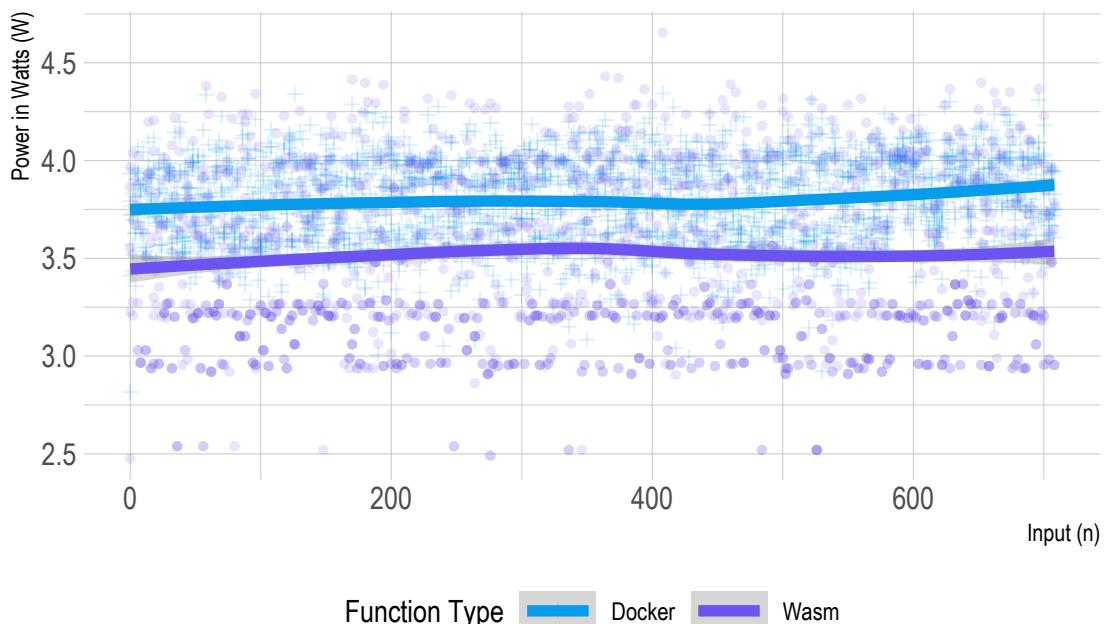


Figure 7.9: Power load on Raspberry Pi during Exponential benchmark.

Exponential

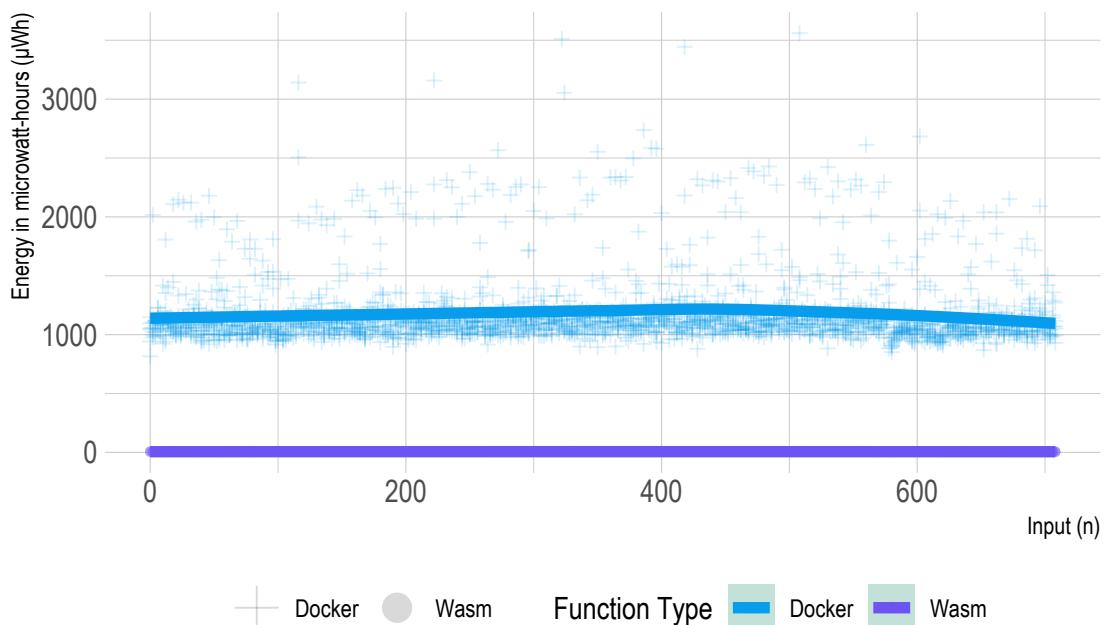


Figure 7.10: Energy consumption on Raspberry Pi during Factorial benchmark.

Prime Number

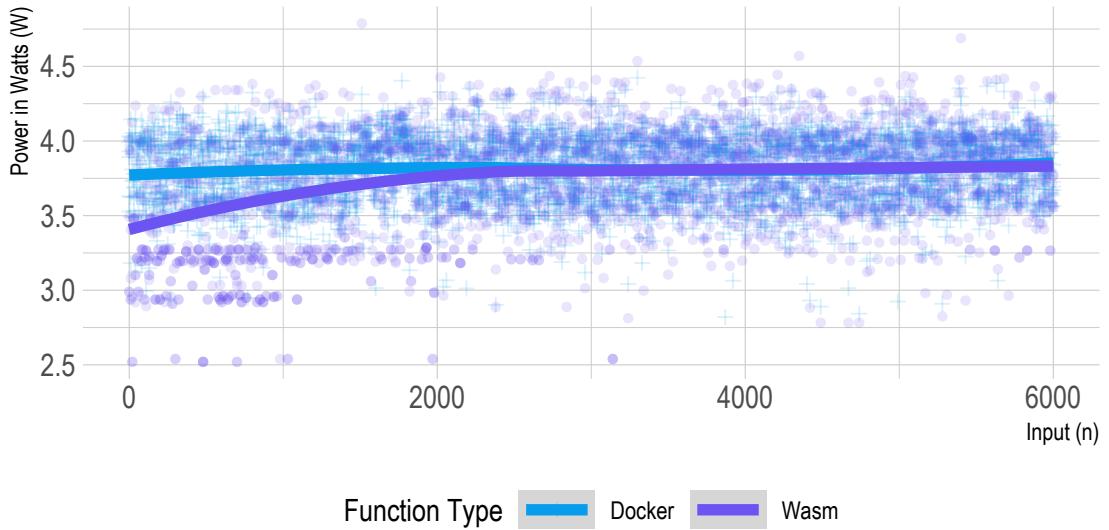


Figure 7.11: Power load on Raspberry Pi during Prime Number benchmark.

Prime Number

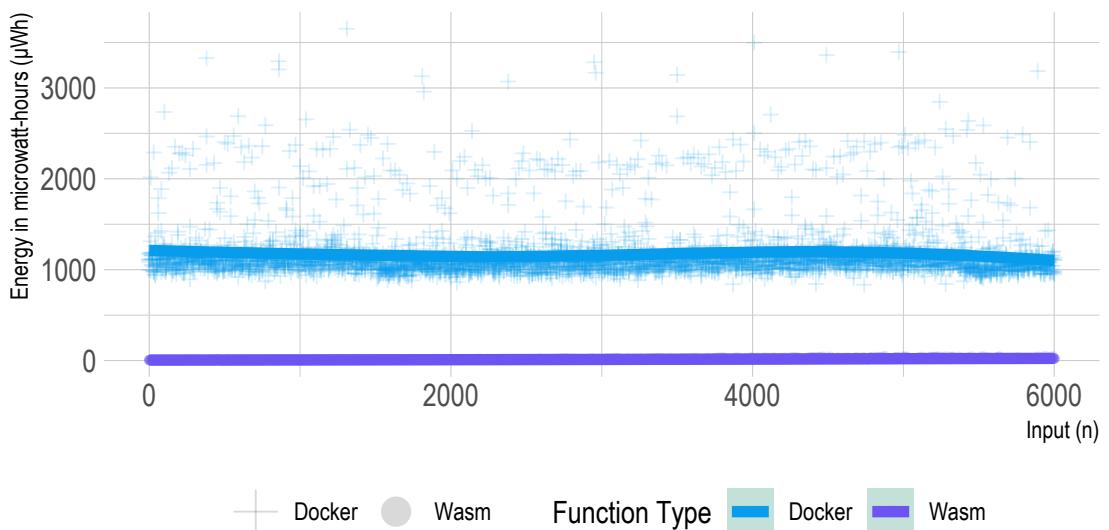


Figure 7.12: Energy consumption on Raspberry Pi during Prime Number benchmark.

These results provide insights into the power and energy implications of executing various functions as Wasm modules and Docker containers on resource-constrained devices like the Raspberry Pi. While this controlled experiment does not necessarily accurately represent larger data centers, the data can still inform decision-making processes related to function deployment, resource allocation, and energy optimization strategies in serverless computing environments.

Chapter 8

Discussion

The aim of argument, or of discussion, should not be victory, but progress.

—Joseph Joubert

The results obtained from our experiments shed light on the performance characteristics and resource implications of serverless function execution using Wasm modules and traditional Docker containers. In this chapter we will discuss the key findings, their significance and compare them to existing research in this domain. As shown in Table 4.2, there is a small degree of inaccuracy in the collected and analyzed data. However, despite this, our results still provide a clear and accurate representation of the performance differences between the two approaches.

8.1 Startup Latency and Runtime Performance

Our experiments revealed a substantial difference in startup latency between Wasm modules and Docker containers. As illustrated in Figure 7.1, Wasm function invocations on average showed startup times approximately 365 times faster than their Docker counterparts on the virtual machine setup. While the gap is smaller on Raspberry Pi, due to its resource constraints, it is still noticeable with Wasm starting up 217 times faster than Docker containers on average. In Figures 7.2 and 7.3 we can see the detailed startup times of our functions across both our setups, and while the values can vary a lot across invocations, we can see that a consistent startup time trend.

In addition to faster startup times, the results in Figure 7.1 indicates that Wasm modules maintained a notable performance advantage over Docker containers in form of overall runtime. On the virtual machine, Wasm runtimes were about 271

times faster than Docker runtimes, and finally 324 times faster for the functions entire lifetime, from start to finish. On the Raspberry Pi, Wasm still showed significant runtime improvements compared to Docker containers, spending 146 times less time calculating the result, and 183 times less time for the functions lifetime. Figures 7.2 and 7.3 shows the total runtime on both setups, including the startup times, and here we can see that most of the time spent during execution of our functions is spent on startup. The exception to this is in the case of our fibonacci and prime number functions, where the total runtime consistently go beyond the startup time at input numbers above 30 and 1500 respectively. This indicates that the improved runtimes of our Wasm-based functions as an alternative to Docker diminishes at heavier workloads.

A very interesting observation that is made apparent from Figure 7.1 is that our Docker containers perform unexpectedly slow in runtime, compared to the runtime of the same code executed as Wasm modules. The ratio of our runtimes across both Wasm and Docker remains largely the same, but with the same source code, one could expect that the Rust binary inside the Docker container would perform close to the same source code compiled to Wasm.

8.2 Power Consumption and Energy Efficiency

Another important aspect explored in our study is the power consumption and energy efficiency of Wasm modules compared to Docker containers. As depicted in Figure 7.5 to fig. 7.12, our results from the Raspberry Pi benchmarks revealed that Wasm function executions consumed approximately 198 times less energy, measured in microwatt-hours (μWh), than their Docker counterparts for the same computational workload.

While the average power load during execution converged to a consistent value for both Wasm and Docker instances at higher input values, the energy consumption curves maintained their trajectories. This divergence highlights the importance of considering both power and time factors when optimizing for energy efficiency in serverless computing environments.

Important to note for our power and energy measurements is that this form of measuring power on a system is inherently not able to paint an accurate representation on how much power and energy our functions consume in isolation. Thus the results presented in Chapter 7 represent a approximation of how much energy our Raspberry Pi as a whole consumes while executing our functions, including its other functions. For our experiments we disabled unnecessary components, such as wifi, bluetooth, HDMI and other connectors that draw constant power. We

could measure a “baseline” power draw that our system draws while idle and not executing any functions and subtract this from our measurements, but this would also only result in a close approximation with mostly the same graphs with lower numbers.

8.3 Function storage

Interestingly, and not something we set out to explore, but is still worth a mention, is that our Wasm modules were relatively small in size. Especially compared to the same source code packaged as a Docker image and gzipped. Table 8.1 shows the storage size, which were mostly identical, for each of our benchmark functions. Using the strip-tool of the wasm-tools program reduces the size of our Wasm by ten times, while still being able to run on Wasmtime without issue.

Table 8.1: Comparison of storage size of functions.

Storage Size	
Docker	32MB
Wasm	2.1MB
Wasm (stripped)	236KB
Wasm vs Docker	14×
Wasm (stripped) vs Docker	138×

8.4 Related Work

Our findings on the performance and efficiency advantages of Wasm align with and complement the existing body of research in this field.

Shillaker and Pietzuch (2020) introduced a novel serverless runtime, which employs Wasm based “Faaslets” for software-fault isolation. Like our observations, they reported that Faaslets achieved cold start times 538 times faster than traditional container-based operators and occupied between 6.5 to 25 times less memory during operation (Table 8.2).

Table 8.2: Comparison of Faaslets vs. container cold starts. Reprinted with permission.

	Docker	Faaslets	vs.Docker
Initialization	2.8 s	5.2 ms	538×
CPU cycles	251M	1.4K	179K×
PSS memory	1.3 MB	200 KB	6.5×
RSS memory	5.0 MB	200 KB	25×
Capacity	~8 K	~70 K	8×

Additionally, they demonstrated a 2x speed-up in machine learning model training and doubled throughput for inference tasks while reducing memory usage by 90% compared to traditional containers with Knative (Figure 8.1).

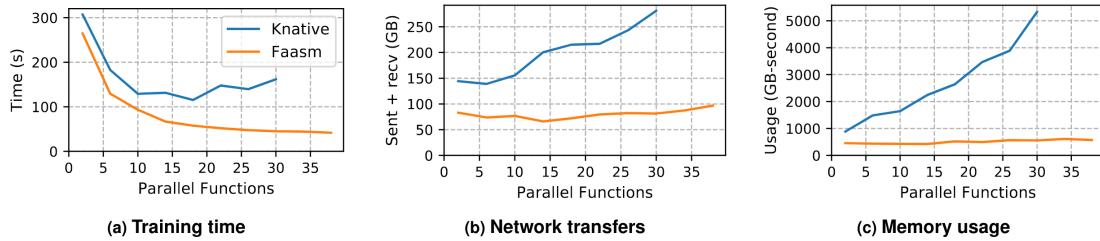


Figure 8.1: Machine learning training with SGD with Faaslets (FAASM) and containers (Knative). Reprinted with permission.

Our results indicating the drastic reduction in startup latency and runtime improvements with Wasm modules align closely with the performance gains reported by Shillaker and Pietzuch (2020), further validating the potential of Wasm as a lightweight and efficient efficient execution environment.

Sebrechts et al. (2022) developed a Wasm-based framework for running lightweight controllers on Kubernetes. They found that their Wasm-based operators significantly reduced memory footprint compared to traditional container-based solutions. In their research they saw reductions from 1405MiB to 227MiB for 100 synthetic operators, and for idle operators they saw a reduction from 1131MiB to 86MiB. (Figure 8.2)

We didn't measure memory usage during our experiments, so the findings of Sebrechts et al. (2022) are not directly applicable to our results, but their results

corroborate with our findings and show that Wasm is a viable option on resource-constrained setups.

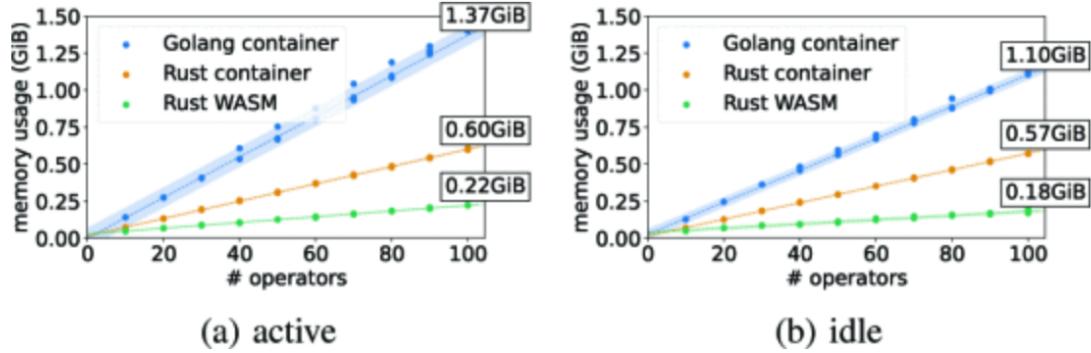


Figure 8.2: Different Kubernetes operators memory usage. Comparing Rust to WASM vs Rust container vs Golang container. © 2022 IEEE

Chapter 9

Conclusion

The more we got into WebAssembly, the more we thought, oh, this is actually a very, very powerful technology that can solve a wide variety of problems.

—Matt Butcher, CEO of Fermyon

With the first goal of this thesis achieved, we move our attention to the topic of whether our hypothesis that Wasm offers a more efficient alternative to Docker in regards to startup, runtime, and energy consumption.

The findings from our experiments conclusively demonstrate the superiority of Wasm over Docker in terms of our scoped domain. As evident from Tables 7.1 and 7.2, Wasm outperformed Docker by significant margins across all aspects, including cold start times, runtime, total runtime, and our estimation for energy consumption. It is important to note that for certain computationally intensive functions with very large input ranges (fibonacci-recursive with input > 30 and prime-number with input > 1500). We excluded those data points to avoid skewing the averages due to excessive runtimes, like we did with figure Figure 7.1.

On the VM environment Table 7.1, Wasm showed a remarkable $365\times$ faster cold start, $271\times$ faster runtime, and $324\times$ faster total runtime compared to the same source code packaged into Docker images. Similarly, on the Raspberry Pi Table 7.2, Wasm showcased impressive improvements, with $217\times$ faster cold start, $146\times$ faster runtime, $183\times$ faster total runtime, and $198\times$ less energy consumed than Docker.

These findings underscore the potential of Wasm as a lightweight and efficient execution environment for serverless functions, offering significant performance benefits over traditional container-based approaches. The drastic reduction in startup latency and runtime times can be a serious advantage in scenarios where

minimizing latency and maximizing throughput are critical, such as in latency-sensitive applications or high-throughput workloads. Furthermore, the significant energy savings, the reduced storage required for storing functions, and the low memory usage (Shillaker & Pietzuch, 2020) achieved by Wasm make it an attractive choice for resource-constrained environments, such as edge computing devices or embedded systems.

While our experiments focused on specific workloads and environments, the observed trends suggest that Wasm could be a viable and compelling alternative to Docker for a wide range of serverless applications. However, it is essential to note that the choice between Wasm and Docker may also depend on other factors, such as the specific requirements of the application, the available resources, and the trade-offs between performance and other considerations, such as security and portability.

In summary, our research has demonstrated the potential of Wasm as a high-performance and energy-efficient alternative to Docker for serverless computing, particularly in scenarios involving small, repeated workloads. It is important to note that our testing focused exclusively on this specific type of workload. Therefore, while Wasm shows superiority in such situations, further research is needed to evaluate its performance in other types of services.

Nevertheless, the findings presented in this thesis provide a solid foundation for further exploration and adoption of Wasm in the serverless ecosystem, particularly in scenarios where performance, latency, and energy efficiency are critical concerns

Chapter 10

Future work

“So will wasm replace Docker?” No, but imagine a future where Docker runs linux containers, windows containers and wasm containers side by side. Over time wasm might become the most popular container type. Docker will love them all equally, and run it all :)

—Solomon Hykes, *Founder of Docker*

While our findings sure are impressive, with Wasm clearly coming ahead in our controlled experiments in regards to the measurements we set out to research, there is still a lot of room for exploring this domain further. In essence, Nebula as a FaaS platform is flawed. For this project it was scoped down to focusing on pure functions, where for every input we provide, we will always get the same output back. This is fine for our controlled experiments, but the real world is not that black or white.

10.1 Investigating Docker Performance Discrepancies

One intriguing observation from our results was the unexpectedly slow runtime performance of the Docker containers compared to the Wasm modules, despite executing the same source code. This discrepancy warrants further investigation to understand the underlying causes and potential optimizations. Future work could involve a comprehensive analysis of the Docker runtime environment inside the container. Factors such as container initialization overhead, resource allocation strategies, and potential bottlenecks in the execution pipeline would prove very useful to improve the body of research.

10.2 Other Programming Languages as Source

As described in Section 6.1.1, Rust has a very close relationship with Wasm. They were both developed by the same company, and the two have a natural alignment for working together. We chose Rust as the source language for this project, because it allowed us to tap into this strength, and achieve synergy with the development of Nebula itself. However, in theory any function written in a language that can compile to the `wasm_wasi` target and expose `stdin` and `stdout` as the means of reading input and writing output, should be able to run on our prototype. Therefore it would be very interesting to see research that explore the idea presented in Figure 3.7, where any code written in a language that supports compiling to Wasm will be able to run on any platform that a Wasm runtime can run on.

10.3 Distributed Deployment and Power Measurement

While our experiments focused on evaluating Wasm modules and Docker containers on individual devices, the serverless computing paradigm often involves distributed deployments across multiple nodes. Future research could explore the deployment of Nebula across a cluster of devices, enabling the measurement of power consumption and energy efficiency in a distributed setting. This would provide valuable insights into the scalability and resource utilization of each deployment approach in a more realistic serverless environment.

Due to the small size of each individual Wasm module binary that we observed in our implementation, there could be an interesting case of exploring the idea of shipping the function code to the execution environment. If the Wasm binaries remain below 300Kb in size when using tools like `wasm-tools`' strip feature, there could be cases where it might be beneficial to move the Wasm binary closer to where the request came from, instead of passing large bodies of JSON data across the globe.

10.4 Accurate Power Consumption Measurements

The power consumption measurements obtained in this study, while insightful, represent approximations of the overall system power draw during function execution. To gain a more precise understanding of the energy efficiency of Wasm modules and Docker containers, future work could leverage advanced power measurement

techniques. For example, the recently released Raspberry Pi 5 incorporates internal systems for measuring power consumption at a granular level, allowing for the isolation of CPU and RAM power consumption from other components. Additionally, Intel CPUs offer the RAPL (Running Average Power Limit) feature, which could provide detailed insights into the power consumption of specific hardware components during function execution.

By addressing these areas, future research can build upon the foundations laid by this thesis, further advancing our understanding of the performance and energy implications of Wasm modules in serverless computing environments.

Acronyms

IaaS Infrastructure-as-a-Service

FaaS Functions-as-a-Service

Wasm WebAssembly

WASI WebAssembly System Interface

VMs virtual machines

VM virtual machine

VM Java Virtual Machine

NIST National Institute of Standards and Technology

AWS Amazon Web Services

GCP Google Cloud Platform

ACI Azure Container Instances

SBC Single board computer

AOT ahead-of-time

MQTT Message Queuing Telemetry Transport

NREC Norwegian Reasearch and Education Cloud

Bibliography

- Adzic, G., & Chatley, R. (2017). Serverless computing: Economic and architectural impact. *Proceedings of the 2017 11th Joint Meeting on Foundations of Software Engineering*, 884–889. <https://doi.org/10.1145/3106237.311776>
- Aeotec. (n.d.). *Aeotec Smart Switch 6 user guide*. Aeotec Help Desk. Retrieved April 29, 2024, from <https://aeotec.freshdesk.com/support/solutions/articles/6000056905-smart-switch-6-user-guide->
- Akkus, I. E., Chen, R., Rimac, I., Stein, M., Satzke, K., Beck, A., Aditya, P., & Hilt, V. (2018). SAND: Towards High-Performance Serverless Computing. <https://www.usenix.org/conference/atc18/presentation/akkus>
- Al-Fuqaha, A., Guizani, M., Mohammadi, M., Aledhari, M., & Ayyash, M. (2015, Winter). Internet of Things: A Survey on Enabling Technologies, Protocols, and Applications. *IEEE Commun. Surv. Tutorials*, 17(4), 2347–2376. <https://doi.org/10.1109/COMST.2015.2444095>
- Amazon. (2019, September 25). *The Climate Pledge*. Retrieved April 14, 2024, from <https://www.aboutamazon.eu/news/sustainability/the-climate-pledge>
- Armbrust, M., Fox, A., Griffith, R., Joseph, A., Katz, R., Konwinski, A., Lee, G., Patterson, D., Rabkin, A., Stoica, I., & Zaharia, M. (2010). A View of Cloud Computing. *Communications of the ACM*, 53, 50–58. <https://doi.org/10.1145/1721654.1721672>
- Baldini, I., Castro, P., Chang, K., Cheng, P., Fink, S., Ishakian, V., Mitchell, N., Muthusamy, V., Rabbah, R., Slominski, A., & Suter, P. (2017, June 10). *Serverless Computing: Current Trends and Open Problems*. Retrieved April 9, 2024, from <http://arxiv.org/abs/1706.03178>
- Bao, W., Hong, C., Sudheer Chunduri, Chunduri, S., Krishnamoorthy, S., Sriram Krishnamoorthy, Sriram Krishnamoorthy, Pouchet, L.-N., Rastello, F., & Sadayappan, P. (2016). Static and Dynamic Frequency Scaling on Multicore CPUs. *ACM Transactions on Architecture and Code Optimization*, 13(4), 51. <https://doi.org/10.1145/3011017>

- Barham, P., Dragovic, B., Fraser, K., Hand, S., Harris, T., Ho, A., Neugebauer, R., Pratt, I., & Warfield, A. (2003). Xen and the Art of Virtualization. *Association for Computing Machinery, Proceedings of the 19th ACM Symposium on Operating Systems Principles*. <https://doi.org/10.1145/945445.945462>
- Barr, J. (2006, August 25). *Amazon EC2 Beta*. AWS News Blog. Retrieved April 1, 2024, from https://aws.amazon.com/blogs/aws/amazon_ec2_beta/
- Bekaroo, G., & Santokhee, A. (2016). Power consumption of the Raspberry Pi: A comparative analysis, 361–366. <https://doi.org/10.1109/EmergiTech.2016.7737367>
- Beloglazov, A., Abawajy, J., & Buyya, R. (2012). Energy-aware resource allocation heuristics for efficient management of data centers for Cloud computing. *Future Generation Computer Systems*, 28(5), 755–768. <https://doi.org/10.1016/j.future.2011.04.017>
- Bernstein, D. (2014). Containers and Cloud: From LXC to Docker to Kubernetes. *IEEE Cloud Computing*, 1(3), 81–84. <https://doi.org/10.1109/MCC.2014.51>
- Bologna, M., & Aquino, G. (2020). Deforestation and world population sustainability: A quantitative analysis. *Scientific Reports*, 10(1), 7631. <https://doi.org/10.1038/s41598-020-63657-6>
- Burns, B., Grant, B., Oppenheimer, D., Brewer, E., & Wilkes, J. (2016). Borg, Omega, and Kubernetes: Lessons learned from three container-management systems over a decade. *Queue*, 14(1), 70–93. <https://doi.org/10.1145/2898442.2898444>
- Butcher, M., & Dodds, E. (2024, January 10). *How WebAssembly is enabling the third wave of cloud compute with Matt Butcher of Fermyon Technologies* (Audio podcast episode No. 172). <https://datastackshow.com/podcast/how-webassembly-is-enabling-the-third-wave-of-cloud-compute-with-matt-butcher-of-fermyon-technologies/>
- Cappelletti, F. (2022, June 21). *The greenest and most ethical energy is the energy we do not consume*. European Commission. Retrieved May 11, 2024, from https://sustainable-energy-week.ec.europa.eu/news/greenest-and-most-ethical-energy-energy-we-do-not-consume-2022-06-21_en
- Castro, P., Ishakian, V., Muthusamy, V., & Slominski, A. (2019). The rise of serverless computing. *Communications of the ACM*, 62(12), 44–54. <https://doi.org/10.1145/3368454>
- Chiueh, S., & Nanda, T.-c. (2005). A Survey on Virtualization Technologies. <https://courses.cs.vt.edu/~cs5204/fallo7-kafura/Papers/Virtualization/Survey-Virtualization-Technologies-TR.pdf>

- Crisman, P. A. (1963). *Computer Time-Sharing System*. MIT. https://people.csail.mit.edu/saltzer/Multics/CTSS-Documents/CTSS_ProgrammersGuide_1963.pdf
- Cuadrado-Cordero, I., Orgerie, A.-C., & Menaud, J.-M. (2018). Comparative Experimental Analysis of the Quality-of-Service and Energy-Efficiency of VMs and Containers' Consolidation for Cloud Applications, 1–6. <https://doi.org/10.23919/SOFTCOM.2017.8115516>
- Durieux, T. (2024, March 12). *Empirical Study of the Docker Smells Impact on the Image Size*. <https://doi.org/10.1145/3597503.3639143>
- Eismann, S., Scheuner, J., van Eyk, E., Schwinger, M., Grohmann, J., Herbst, N., Abad, C. L., & Iosup, A. (2021, January 28). *A Review of Serverless Use Cases and their Characteristics*. arXiv: 2008.11110. <http://arxiv.org/abs/2008.11110>
- Etro, F. (2009). The Economic Impact of Cloud Computing on Business Creation, Employment and Output in Europe. An application of the Endogenous Market Structures Approach to a GPT innovation. *Review of Business and Economics*, 54, 179–208. <https://ideas.repec.org/a/ete/revbec/20090204.html>
- Favaloro, G., & O'Sullivan, S. (1996, November 14). *Internet Solutions Division Strategy for Cloud Computing*. https://s3.amazonaws.com/files.technologyreview.com/p/pub/legacy/compaq_cst_1996_0.pdf
- Freitag, C., Berners-Lee, M., Widdicks, K., Knowles, B., Blair, G. S., & Friday, A. (2021). The real climate and transformative impact of ICT: A critique of estimates, trends, and regulations. *Patterns*, 2(9), 100340. <https://doi.org/10.1016/j.patter.2021.100340>
- GmbH, G. S. (2023). Expert Power Control 1104 / 1105 Manual Guide. *Gude Systems*. https://files.gude-systems.com/pdf/gude/manual-epc1105-series_v1.5.1.pdf
- Golec, M., Walia, G. K., Kumar, M., Cuadrado, F., Gill, S. S., & Uhlig, S. (2023, October 12). *Cold Start Latency in Serverless Computing: A Systematic Review, Taxonomy, and Future Directions* [Comment: 34 Pages, 16 Figures]. arXiv: 2310.08437 [cs]. Retrieved April 15, 2024, from <http://arxiv.org/abs/2310.08437>
- Google. (n.d.). *Tracking Our Carbon-Free Energy Progress*. Google Sustainability. Retrieved April 14, 2024, from <https://sustainability.google/progress/energy/>
- Google 2023 Environmental Report. (2023). <https://www.gstatic.com/gumdrop/sustainability/google-2023-environmental-report.pdf>

- Gottlieb, N. (2018, June 20). *On serverless, data lock-in and vendor choice*. Retrieved April 15, 2024, from <https://serverless.com/blog/data-lockin-vendor-choice-portability>
- Gunnerød, J. L. (2022). Statnett - Kortsiktig Markedsanalyse November 2022.
- Haas, A., Rossberg, A., Schuff, D. L., Titzer, B. L., Holman, M., Gohman, D., Wagner, L., Zakai, A., & Bastien, J. (2017). Bringing the web up to speed with WebAssembly. *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation*, 185–200. <https://doi.org/10.1145/3062341.3062363>
- Kaplan, J. M., Forrest, W., & Kindler, N. (2008). Revolutionizing Data Center Energy Efficiency. *McKinsey & Company*. <https://docplayer.net/5045147-Revolutionizing-data-center-energy-efficiency.html>
- Kjorveziroski, V., Filiposka, S., & Mishev, A. (2022). Evaluating WebAssembly for Orchestrated Deployment of Serverless Functions. *2022 30th Telecommunications Forum (TELFOR)*, 1–4. <https://doi.org/10.1109/TELFOR56187.2022.9983733>
- Leonard, J. (2024, March 26). *WebAssembly heralds 'third wave of cloud computing'*. Retrieved April 19, 2024, from <https://wwwcomputing.co.uk/news/4189373/webassembly-heralds-wave-cloud-computing>
- Malmgren, P. (2022, April 17). *Getting data in and out of WASI modules - Peter Malmgren*. Retrieved May 7, 2024, from <https://petermalmgren.com/serverside-wasm-data/>
- McDonald, P. (2008, April 7). *Introducing Google App Engine + our new blog*. Google App Engine Blog. Retrieved April 8, 2024, from <https://googleappengine.blogspot.com/2008/04/introducing-google-app-engine-our-new.html>
- McGrath, G., & Brenner, P. R. (2017). Serverless Computing: Design, Implementation, and Performance. *2017 IEEE 37th International Conference on Distributed Computing Systems Workshops (ICDCSW)*, 405–410. <https://doi.org/10.1109/ICDCSW.2017.36>
- Mell, P., & Grance, T. (2011, September 1). The NIST Definition of Cloud Computing. *Special Publication (NIST SP), National Institute of StandardsTechnology, Gaithersburg, MD*. <https://doi.org/10.6028/NIST.SP.800-145>
- Merkel, D. (2014, May 9). *Docker: Lightweight Linux Containers for Consistent Development and Deployment*. Linux Journal. Retrieved April 16, 2024, from <https://www.linuxjournal.com/content/docker-lightweight-linux-containers-consistent-development-and-deployment>
- Mytton, D. (2020). Hiding greenhouse gas emissions in the cloud. *Nat. Clim. Chang.*, 10(8), 701–701. <https://doi.org/10.1038/s41558-020-0837-6>

- Ni, J., & Bai, X. (2017). A review of air conditioning energy performance in data centers. *Renewable and Sustainable Energy Reviews*, 67, 625–640. <https://doi.org/10.1016/j.rser.2016.09.050>
- Norwegian-Energy. (2023, November 8). *Electricity production*. Norwegian Energy. Retrieved April 14, 2024, from <https://energifaktanorge.no/en/norsk-energiforsyning/kraftproduksjon/>
- Qian, Z. (2022, January 7). *Energy footprinting of WordPress websites on the user side* (Master Thesis). University of Bristol.
- Rasmussen, N. (2018). Impact of Leading Power Factor on Data Center Generator Systems. https://3dfs.com/wp-content/uploads/2015/11/VAVR-9EJ262_Ro_EN.pdf
- Rivrud, K. (2024, February 7). *Investeringen av et Google-senter i Skien er gigantisk – det blir også strømforbruket*. NRK. Retrieved April 12, 2024, from <https://www.nrk.no/vestfoldogtelemark/investeringen-av-et-google-senter-i-skien-er-gigantisk--det-blir-ogsa-stromforbruket-1.16753588>
- Roberts, M. (2018, May 22). *Serverless Architectures*. martinfowler.com. Retrieved April 15, 2024, from <https://martinfowler.com/articles/serverless.html>
- Sebrechts, M., Ramlot, T., Borny, S., Goethals, T., Volckaert, B., & De Turck, F. (2022). Adapting Kubernetes controllers to the edge: On-demand control planes using Wasm and WASI. *2022 IEEE 11th International Conference on Cloud Networking (CloudNet)*, 195–202. <https://doi.org/10.1109/CloudNet55617.2022.9978884>
- Sergeev, A., Rezedinova, E., & Khakhina, A. (2022). Docker Container Performance Comparison on Windows and Linux Operating Systems, 1–4. <https://doi.org/10.1109/CIEES55704.2022.9990683>
- Sewak, M., & Singh, S. (2018). Winning in the Era of Serverless Computing and Function as a Service. *2018 3rd International Conference for Convergence in Technology (I2CT)*, 1–5. <https://doi.org/10.1109/I2CT.2018.8529465>
- Shehabi, A., Smith, S., Sartor, D., Brown, R., Herrlin, M., Koomey, J., Masanet, E., Horner, N., Azevedo, I., & Lintner, W. (2016, June 1). *United States Data Center Energy Usage Report* (LBNL-1005775, 1372902). <https://doi.org/10.2172/1372902>
- Shillaker, S., & Pietzuch, P. (2020). Faasm: Lightweight Isolation for Efficient Stateful Serverless Computing, 419–433. Retrieved April 22, 2024, from <https://www.usenix.org/conference/atc20/presentation/shillaker>
- Shirinbab, S., Lundberg, L., & Casalicchio, E. (2020). Performance evaluation of containers and virtual machines when running Cassandra workload concurrently. *Concurrency and Computation: Practice and Experience*, 32. <https://doi.org/10.1002/cpe.5693>

- Smith, B. (2020, January 16). *Microsoft will be carbon negative by 2030*. The Official Microsoft Blog. Retrieved April 14, 2024, from <https://blogs.microsoft.com/blog/2020/01/16/microsoft-will-be-carbon-negative-by-2030/>
- Solomon Hykes [@solomonstre]. (2019, March 27). *If WASM+WASI existed in 2008, we wouldn't have needed to created Docker. That's how important it is. Webassembly on the server is the future of computing. A standardized system interface was the missing link. Let's hope WASI is up to the task!* Twitter. Retrieved April 15, 2024, from <https://twitter.com/solomonstre/status/1111004913222324225>
- Sundt, S., & Rehdanz, K. (2015). Consumers' willingness to pay for green electricity: A meta-analysis of the literature. *Energy Economics*, 51, 1–8. <https://doi.org/10.1016/j.eneco.2015.06.005>
- Tong, W., Gao, J., Jin, X., & Li, Z. (2015). Study on Ethernet Communication Strategy for Electric Energy Data Acquisition System. *2015 Fifth International Conference on Instrumentation and Measurement, Computer, Communication and Control (IMCCC)*, 1240–1244. <https://doi.org/10.1109/IMCCC.2015.266>
- Van Eyk, E., Toader, L., Talluri, S., Versluis, L., Uta, A., & Iosup, A. (2018). Serverless is More: From PaaS to Present Cloud Computing. *IEEE Internet Comput.*, 22(5), 8–17. <https://doi.org/10.1109/MIC.2018.053681358>
- Wang, L., Li, M., Zhang, Y., Ristenpart, T., & Swift, M. (2018). Peeking Behind the Curtains of Serverless Platforms.
- WASI/preview2/README.md at main · WebAssembly/WASI*. (n.d.). GitHub. Retrieved April 24, 2024, from <https://github.com/WebAssembly/WASI/blob/main/preview2/README.md>
- WebAssembly.org. (n.d.). *FAQ - WebAssembly*. Retrieved April 16, 2024, from <https://webassembly.org/docs/faq/>
- Zhang, Y., Liu, M., Wang, H., Ma, Y., Huang, G., & Liu, X. (2024, April 19). *Research on WebAssembly Runtimes: A Survey*. Retrieved April 23, 2024, from <http://arxiv.org/abs/2404.12621>